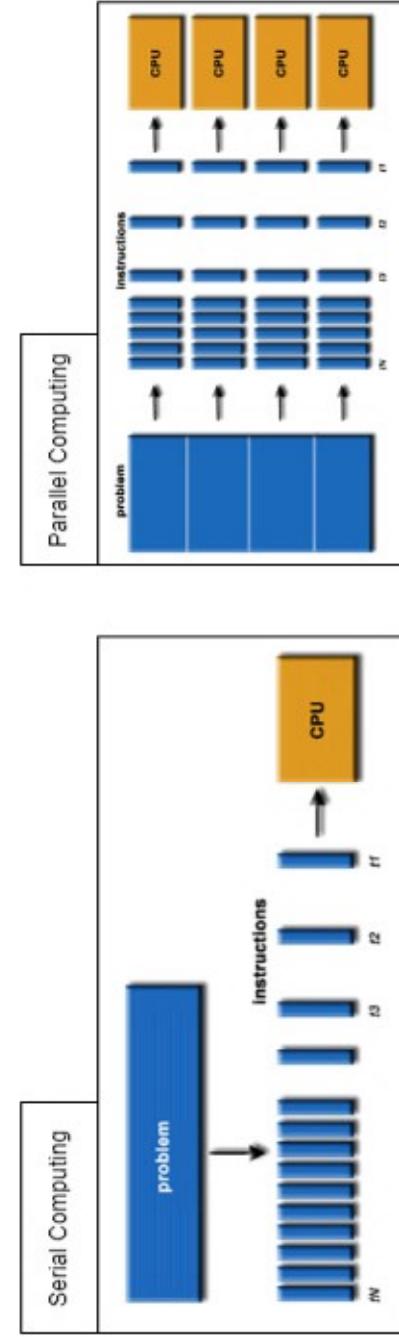
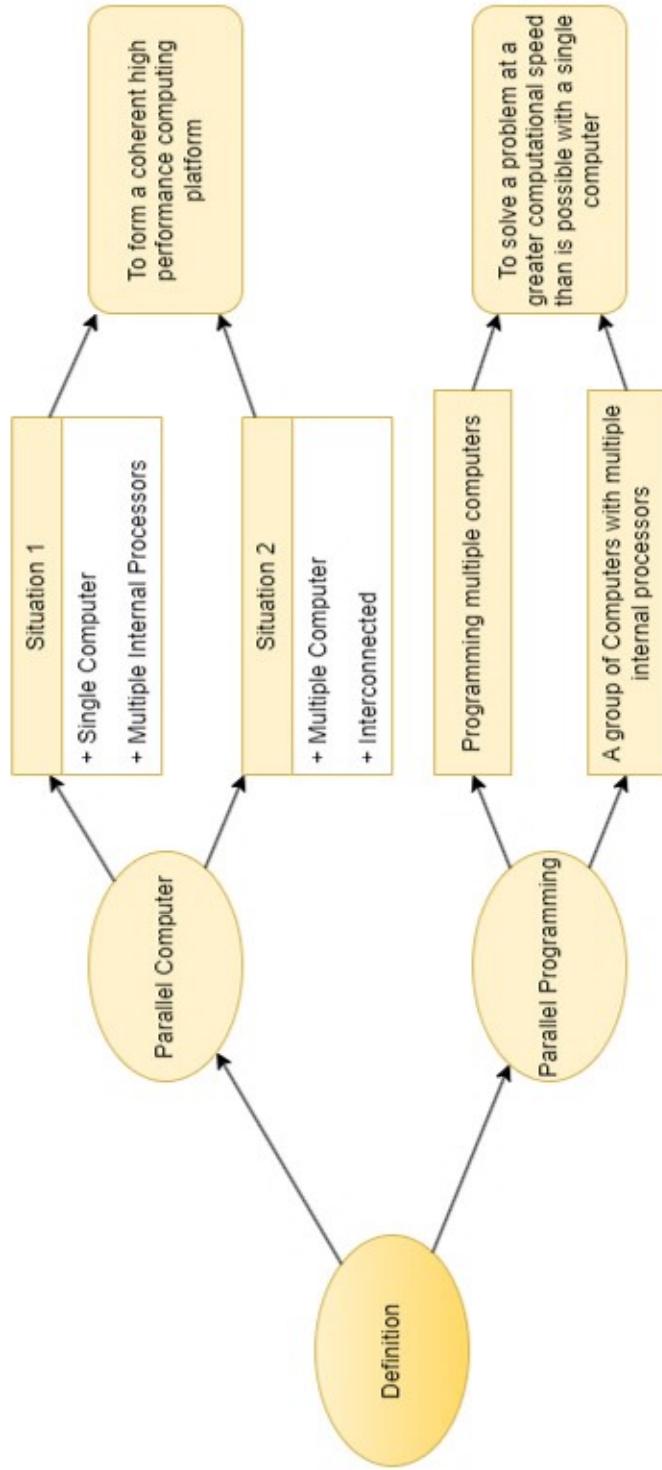


Lecture 1: Introduction

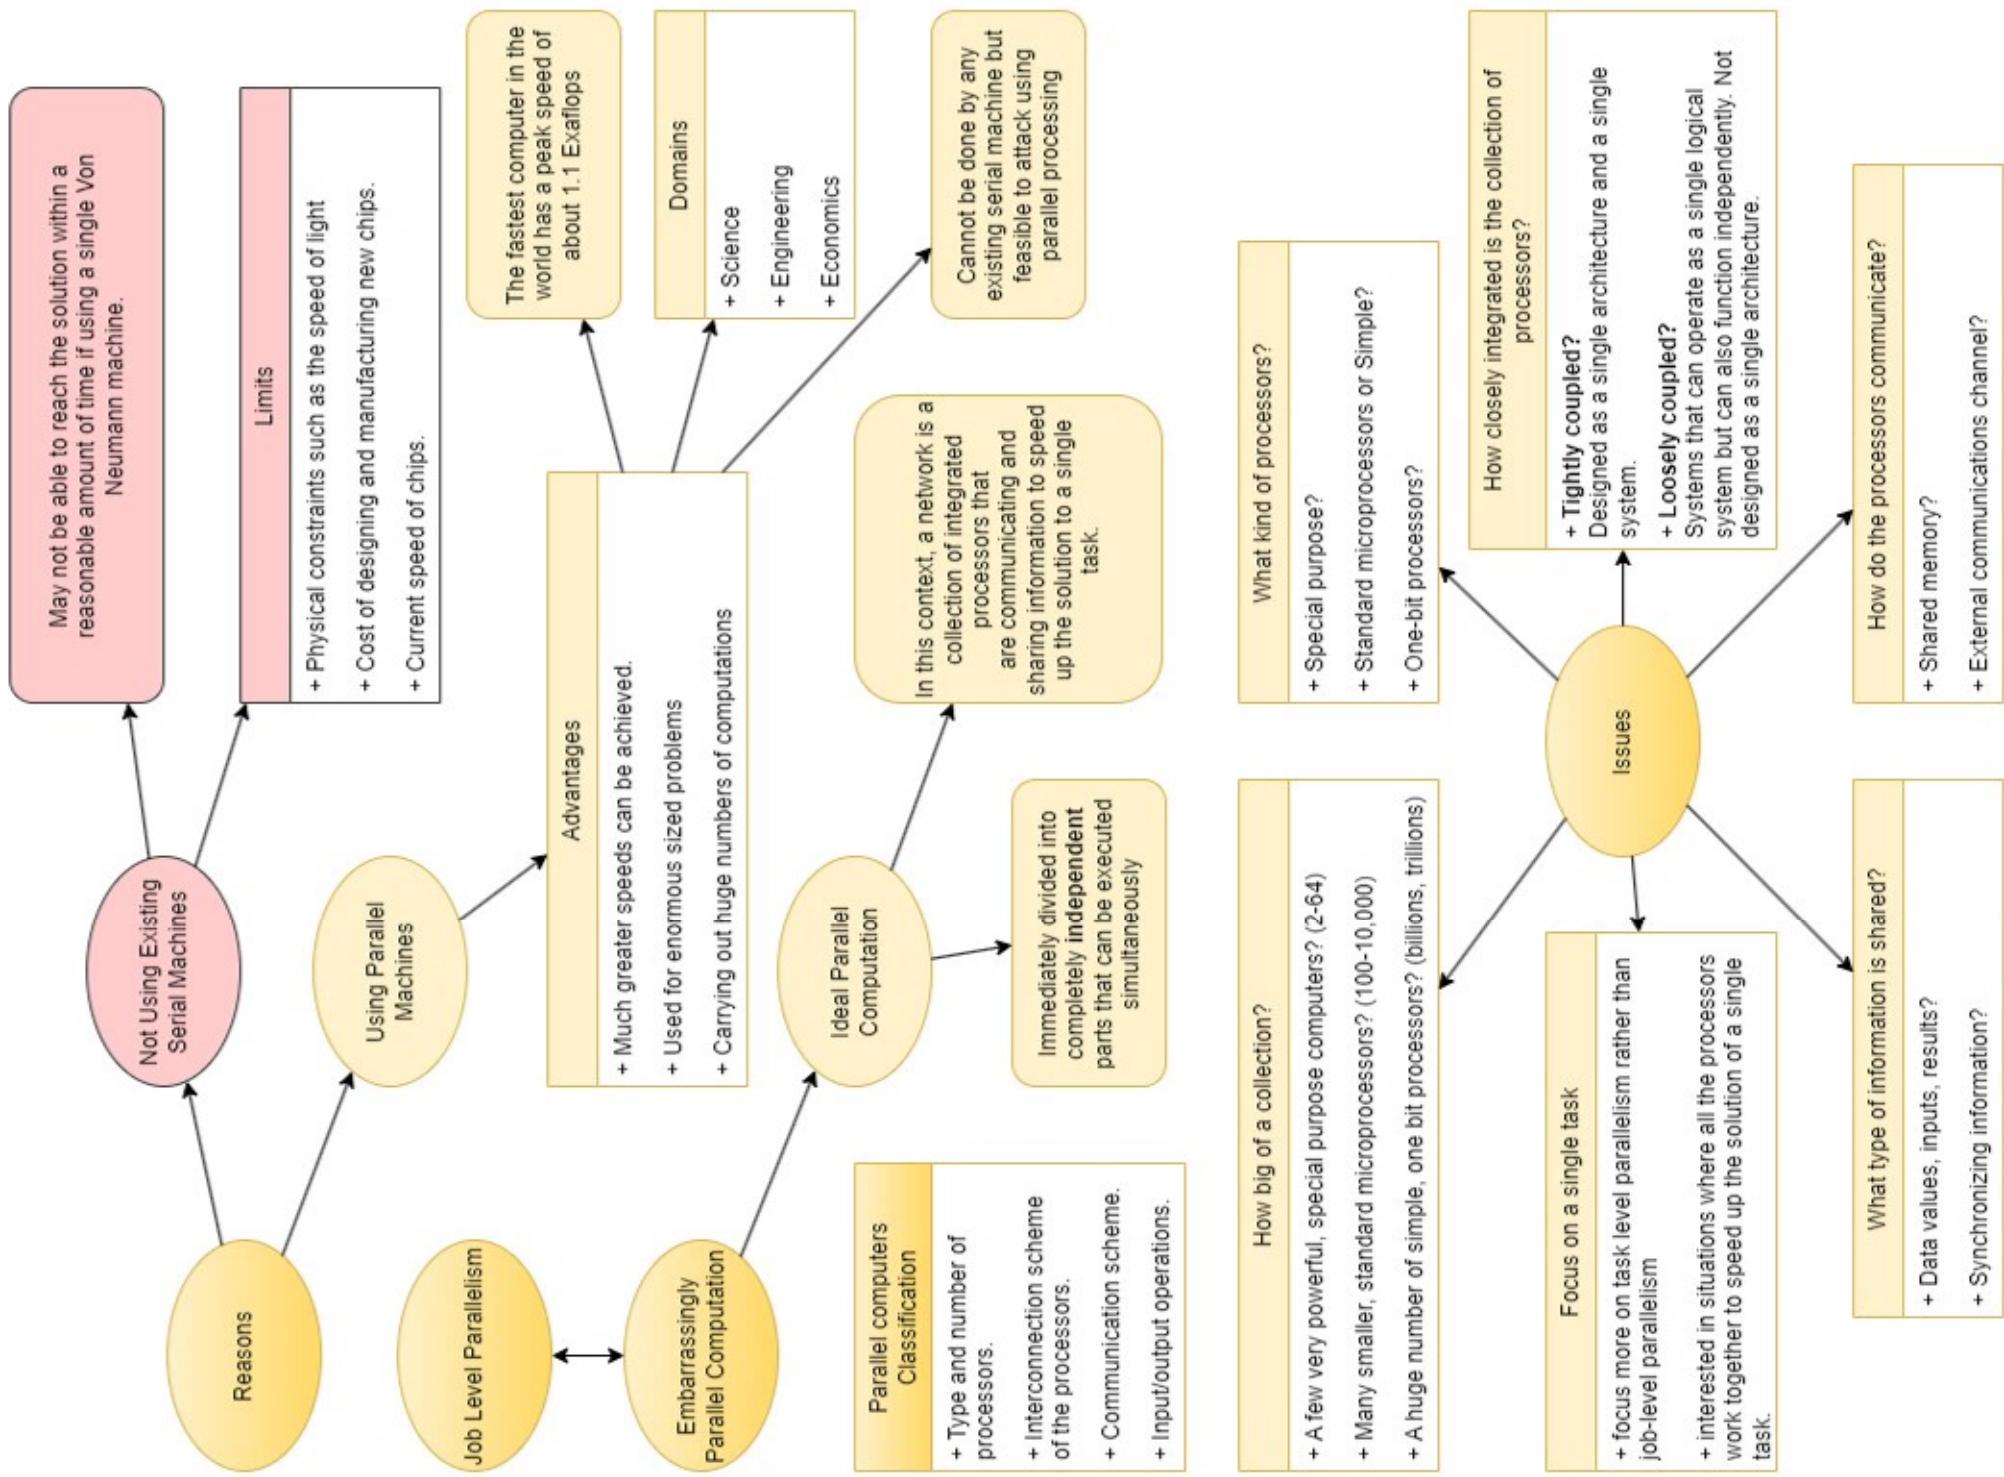


Demand for Computational Speed
Computations must be completed within a "reasonable" time period.
Continual Demand for greater computational speed from a computer system than is currently possible.

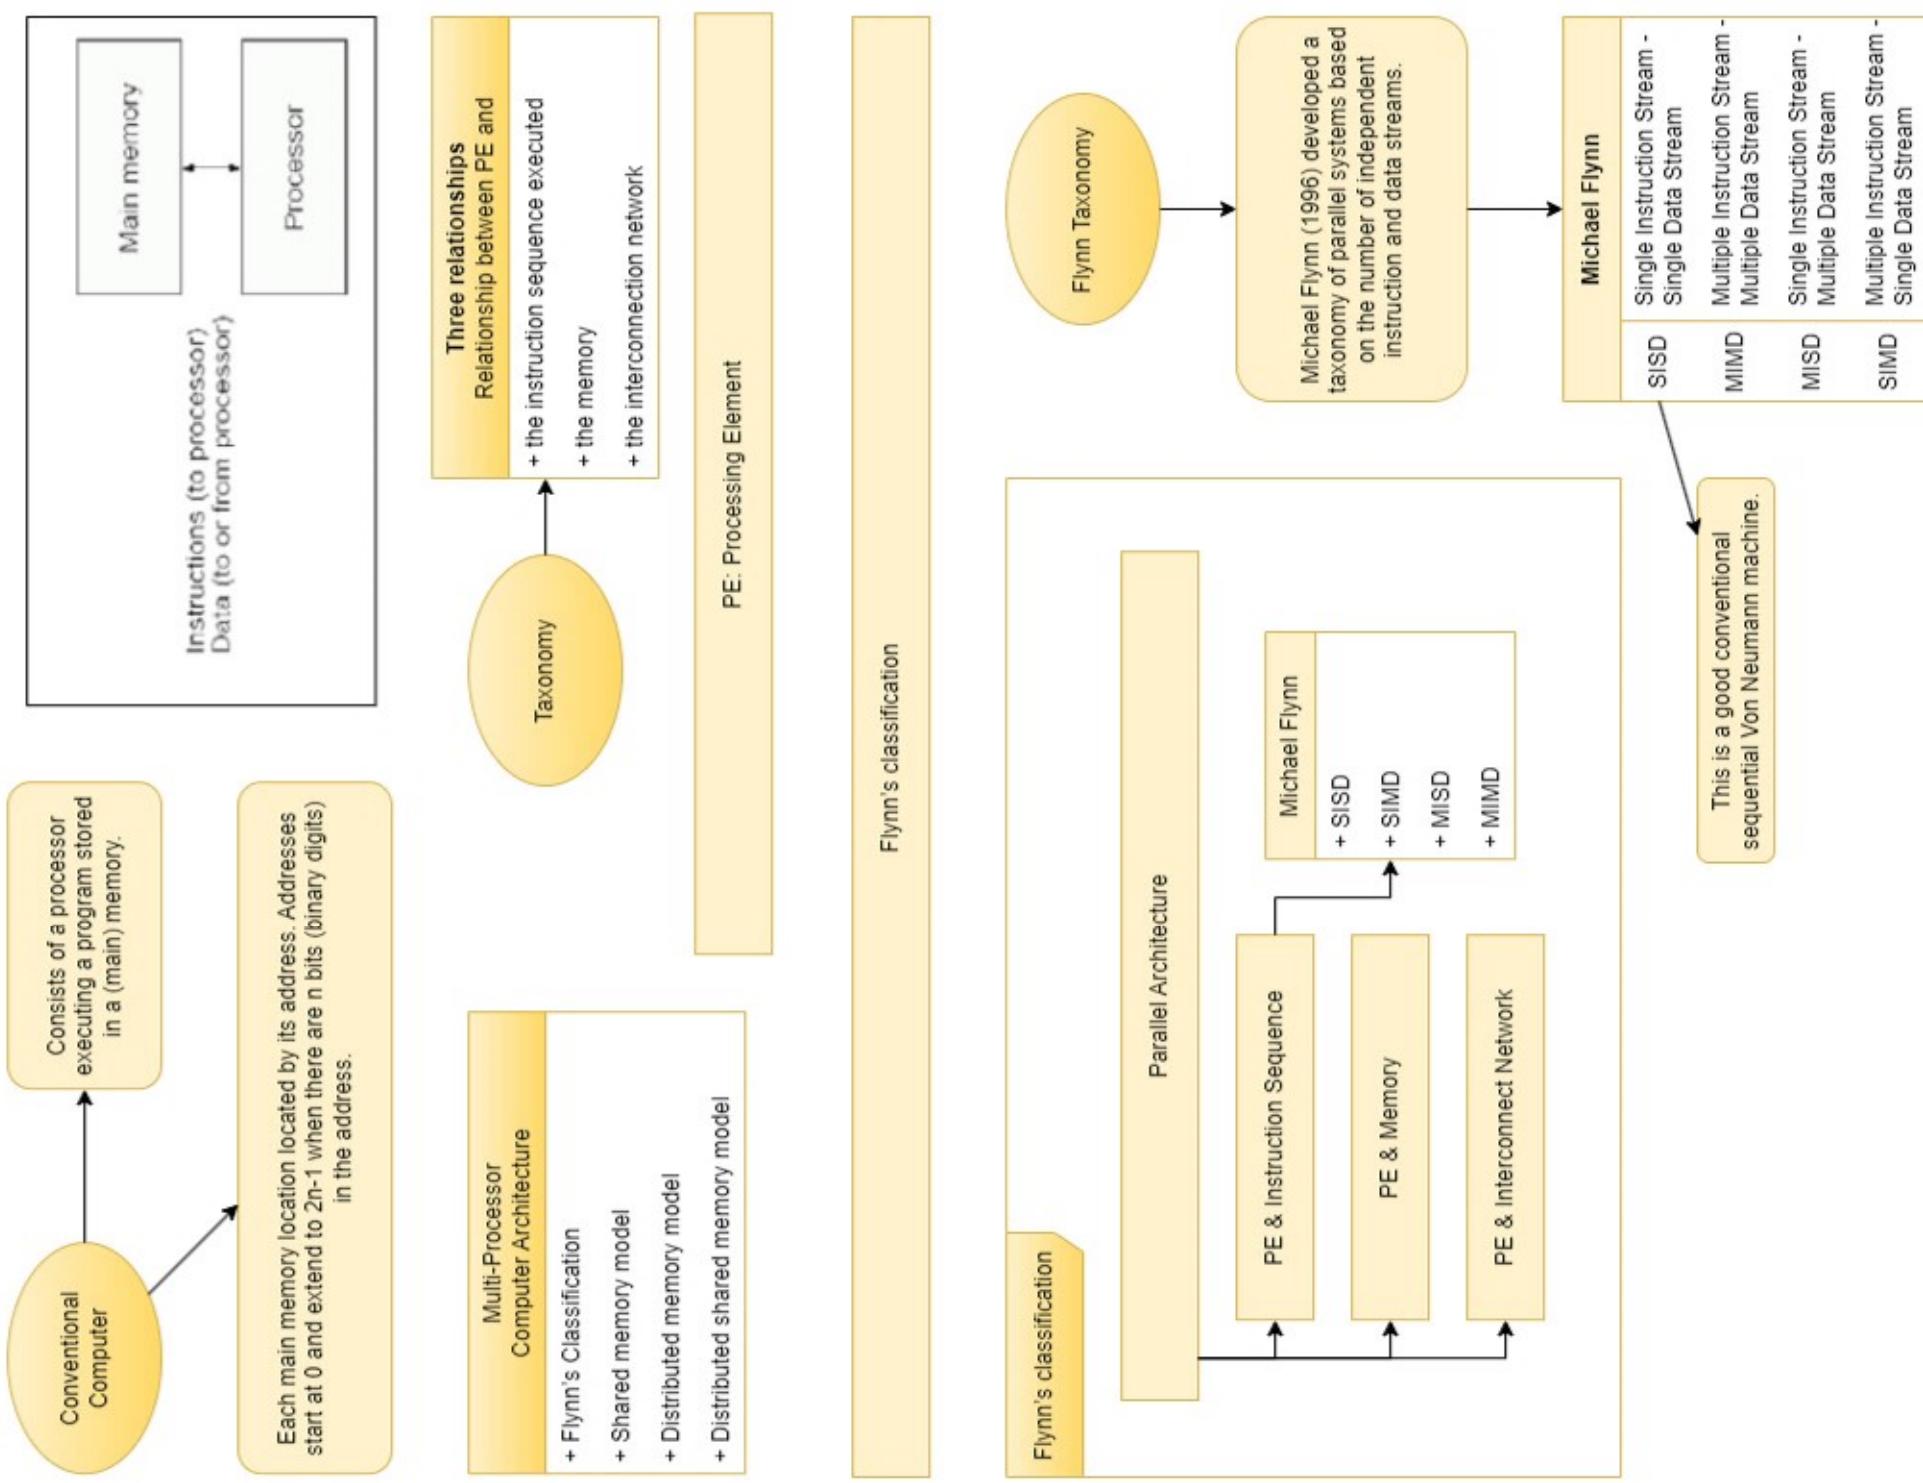
Areas requiring great computational speed
+ Numeric Modelling
+ Simulation of scientific and engineering problems

Grand Challenge Problems
one that cannot be solved in a reasonable amount of time with today's computers.
Example
+ Modelling large DNA Structures
+ Global weather forecasting
+ Modelling motion of astronomical bodies

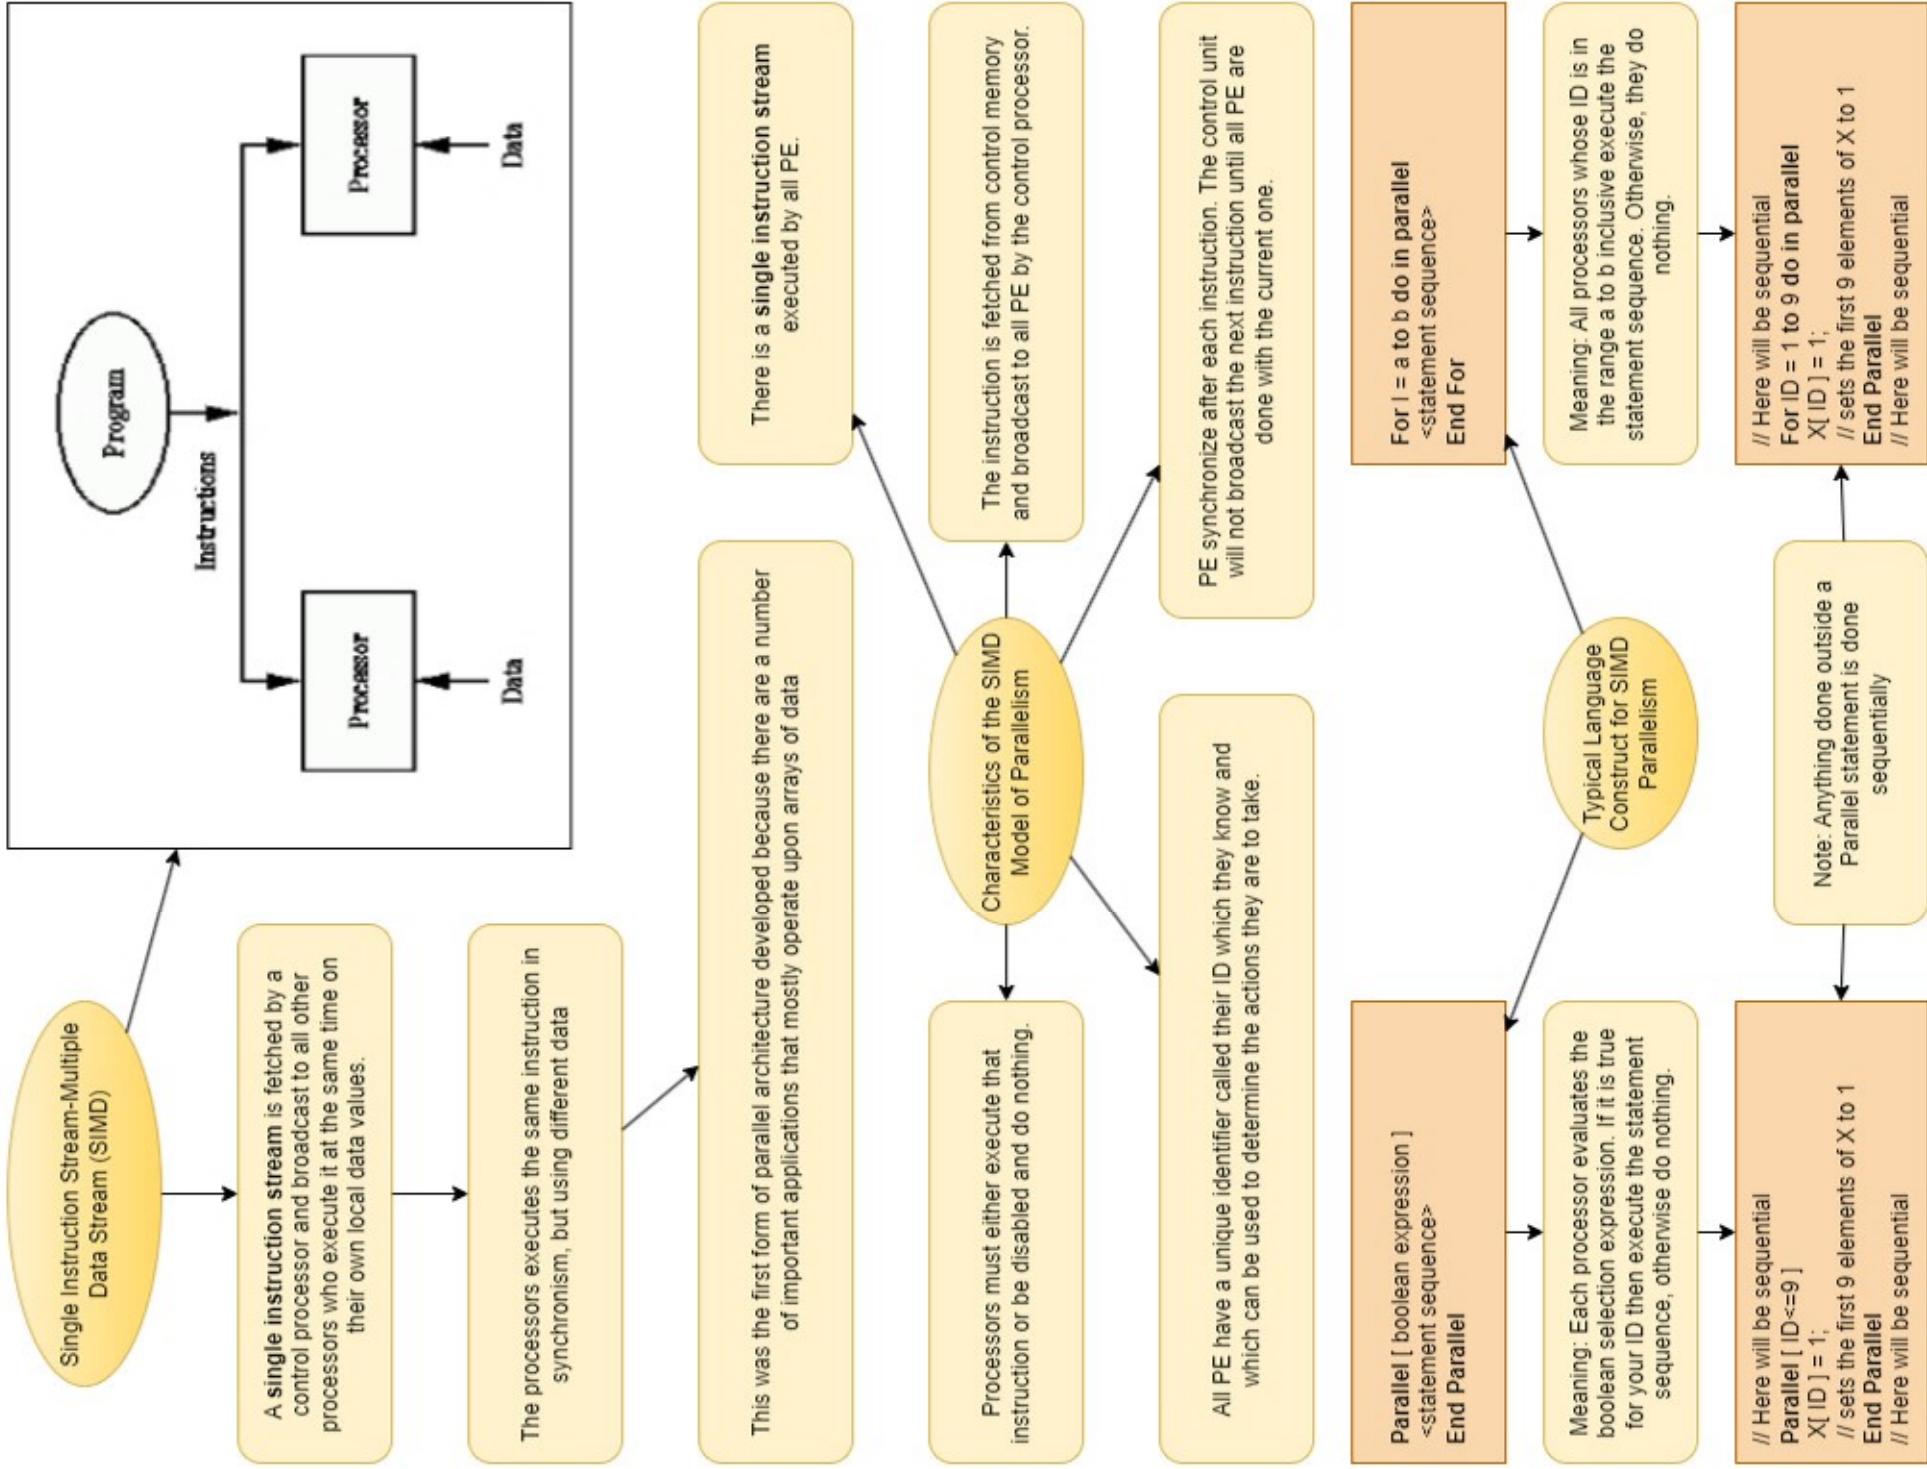
Lecture 1: Introduction



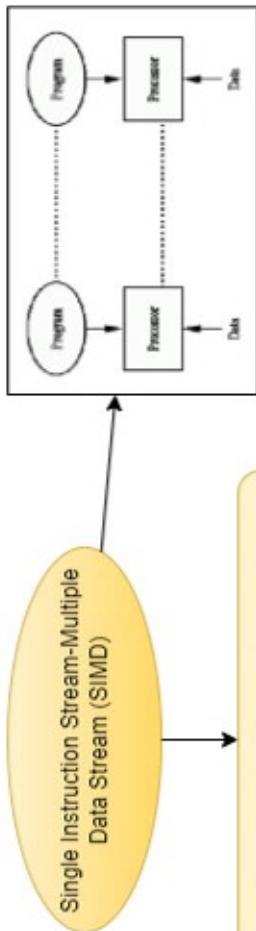
Lecture 1: Multi-Processor Computer Architecture



Lecture 1: Multi-Processor Computer Architecture



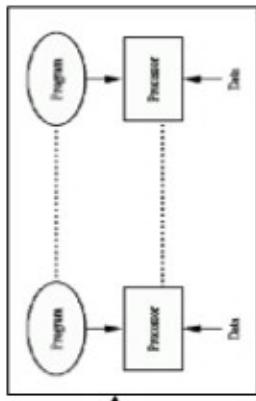
Lecture 1: Multi-Processor Computer Architecture



General-purpose multiprocessor system - each processor has a separate program and one instruction stream is generated from each program for each processor.

Each instruction operates upon different data.

Both the shared memory and the message-passing multiprocessors so far described are in the MIMD classification.



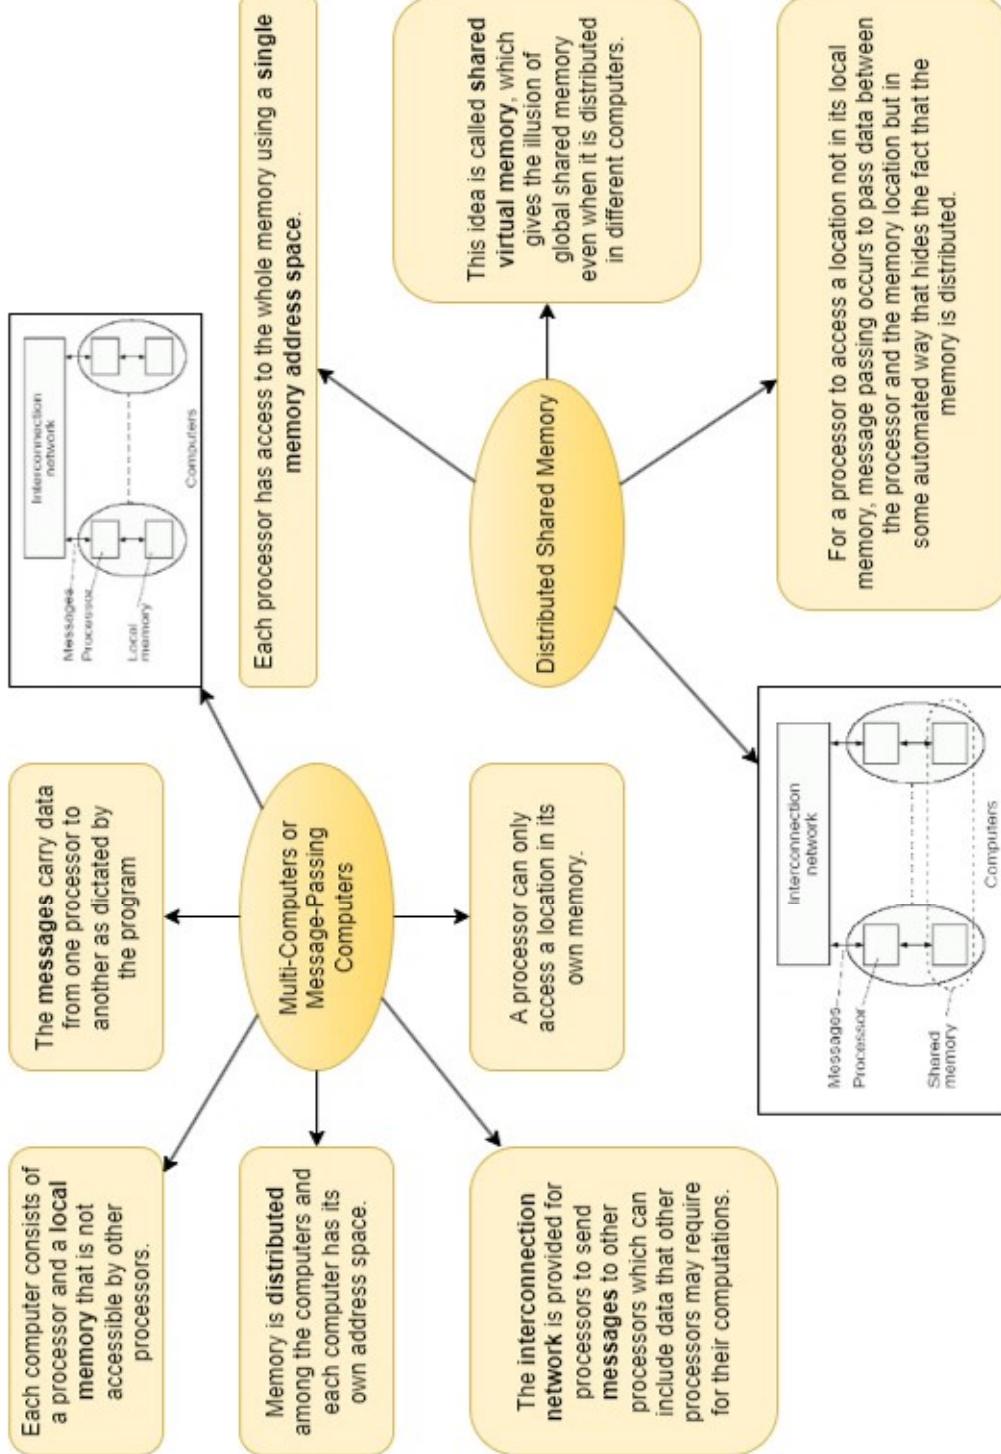
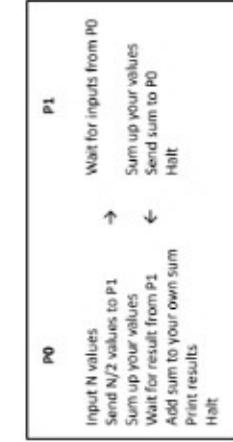
Each processing element executes its own program stream.

Characteristics of the MIMD Model of Parallelism

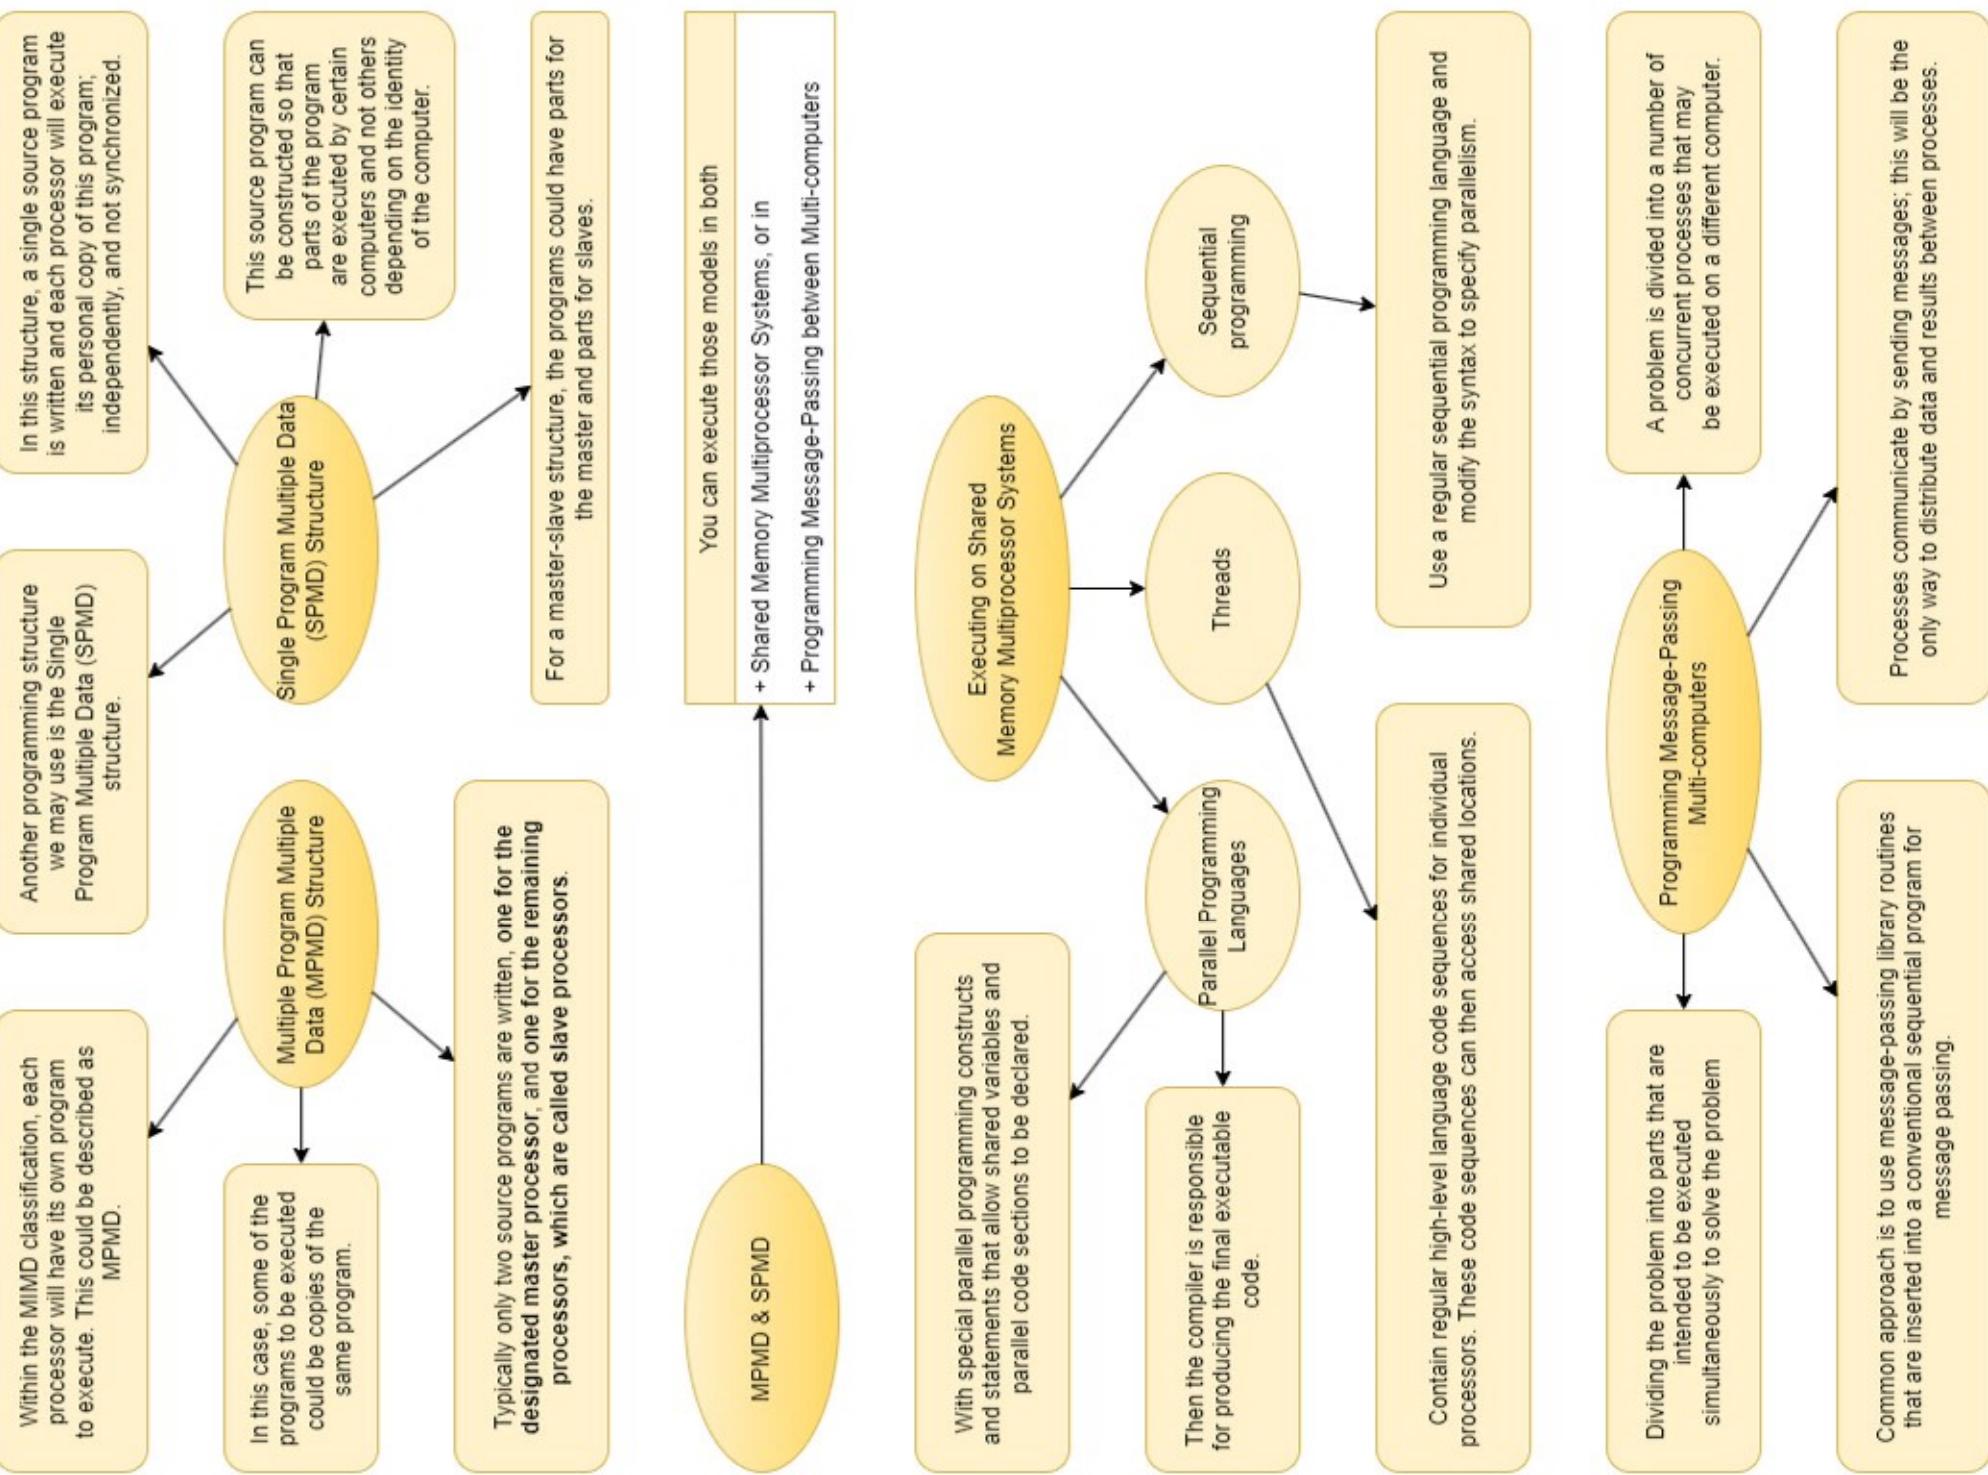
There is no central control. Each processor executes independently of all other processors.

There is no central clock. Each processor executes at its own rate of speed.

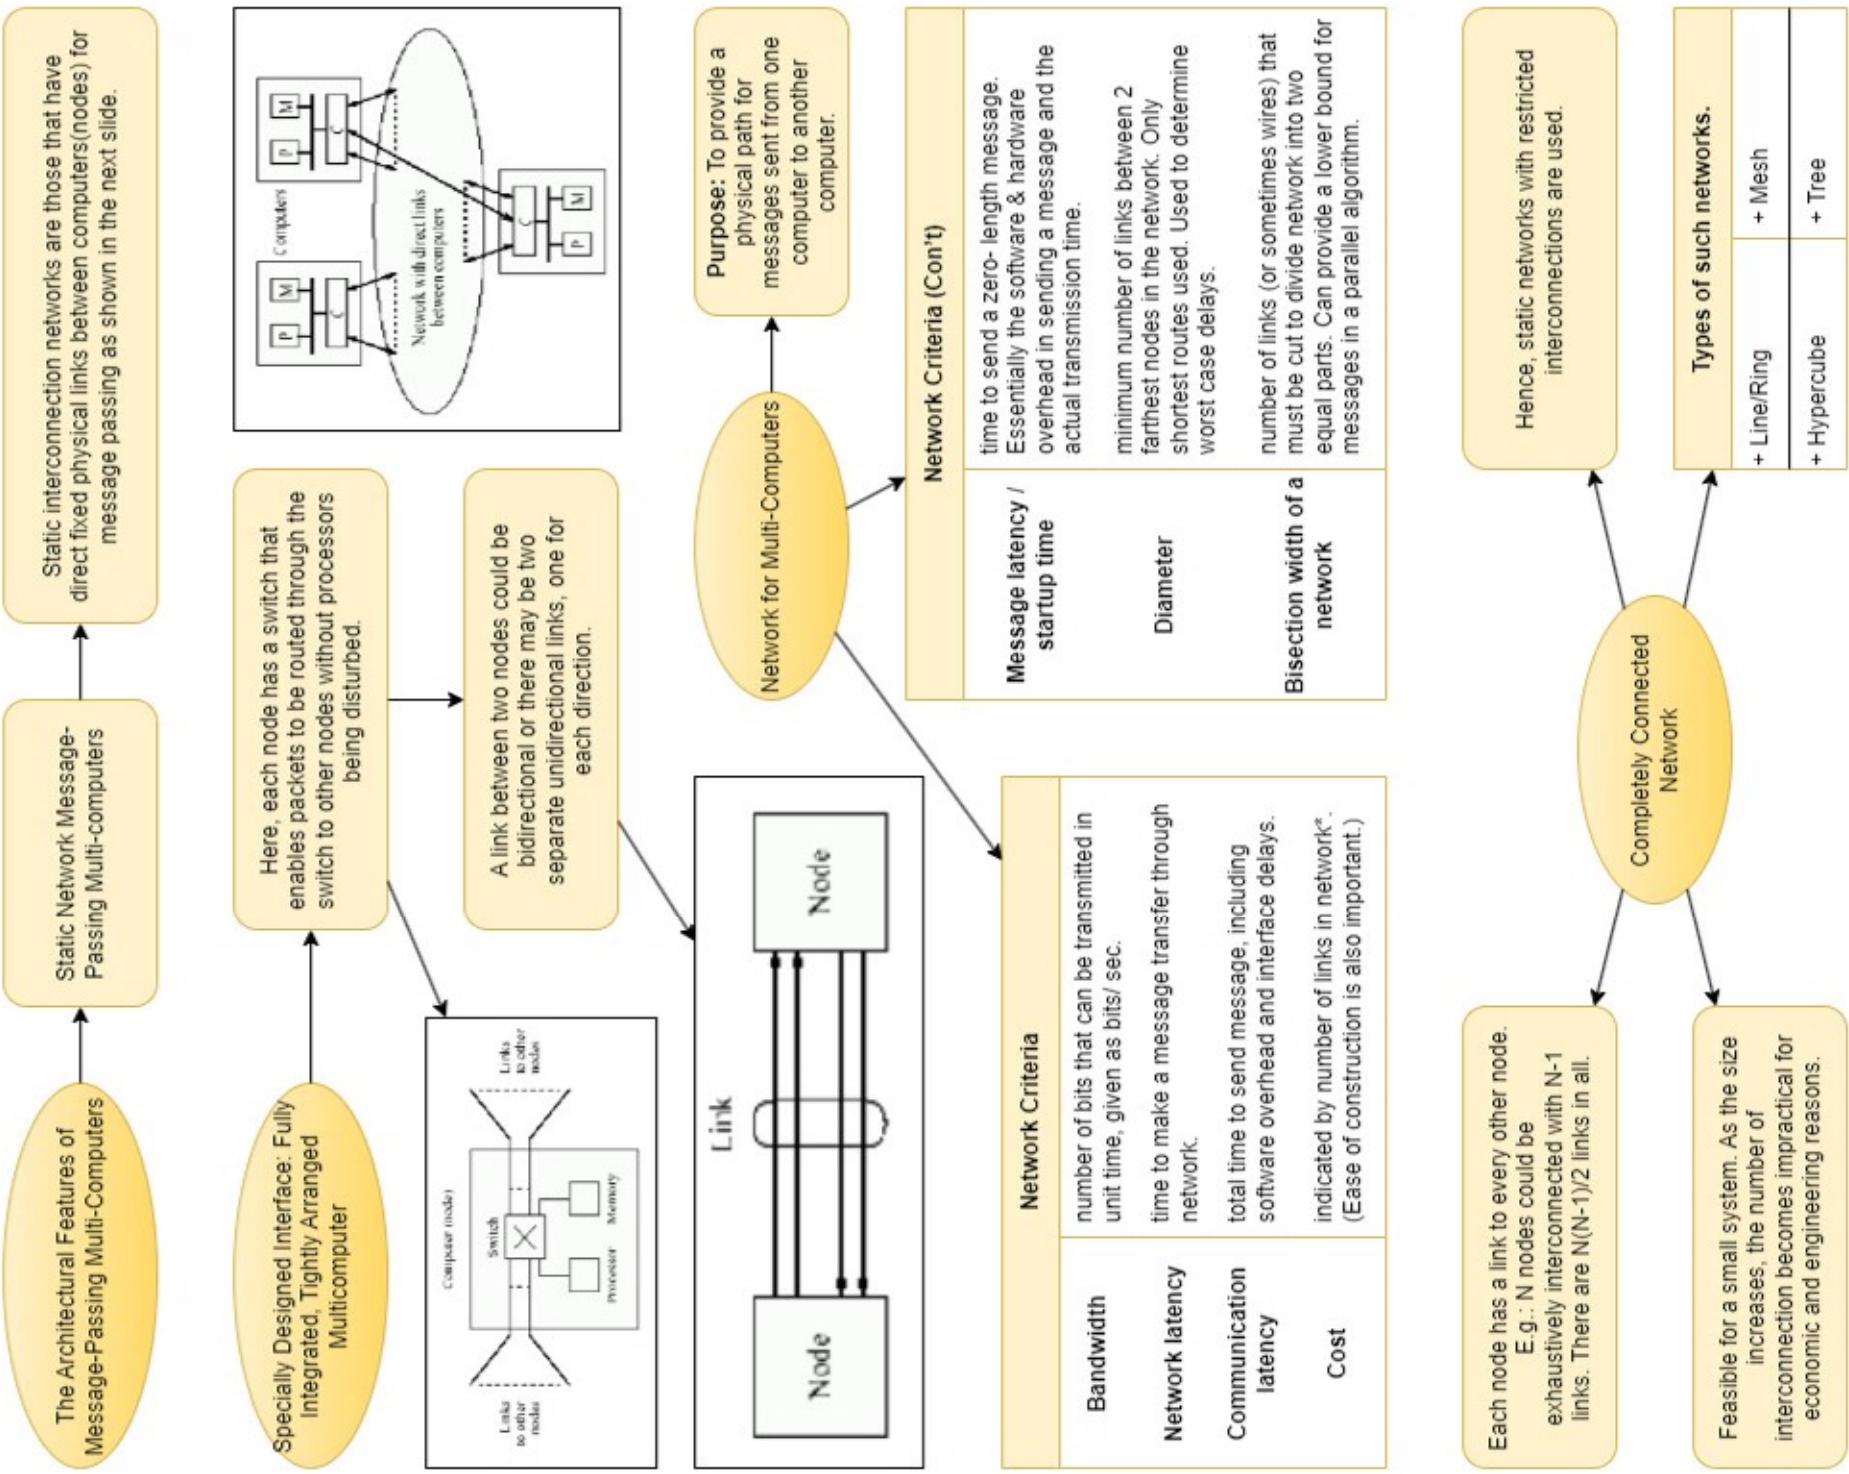
There is no automatic hardware synchronization. Synchronization must be done in software by the programs themselves.



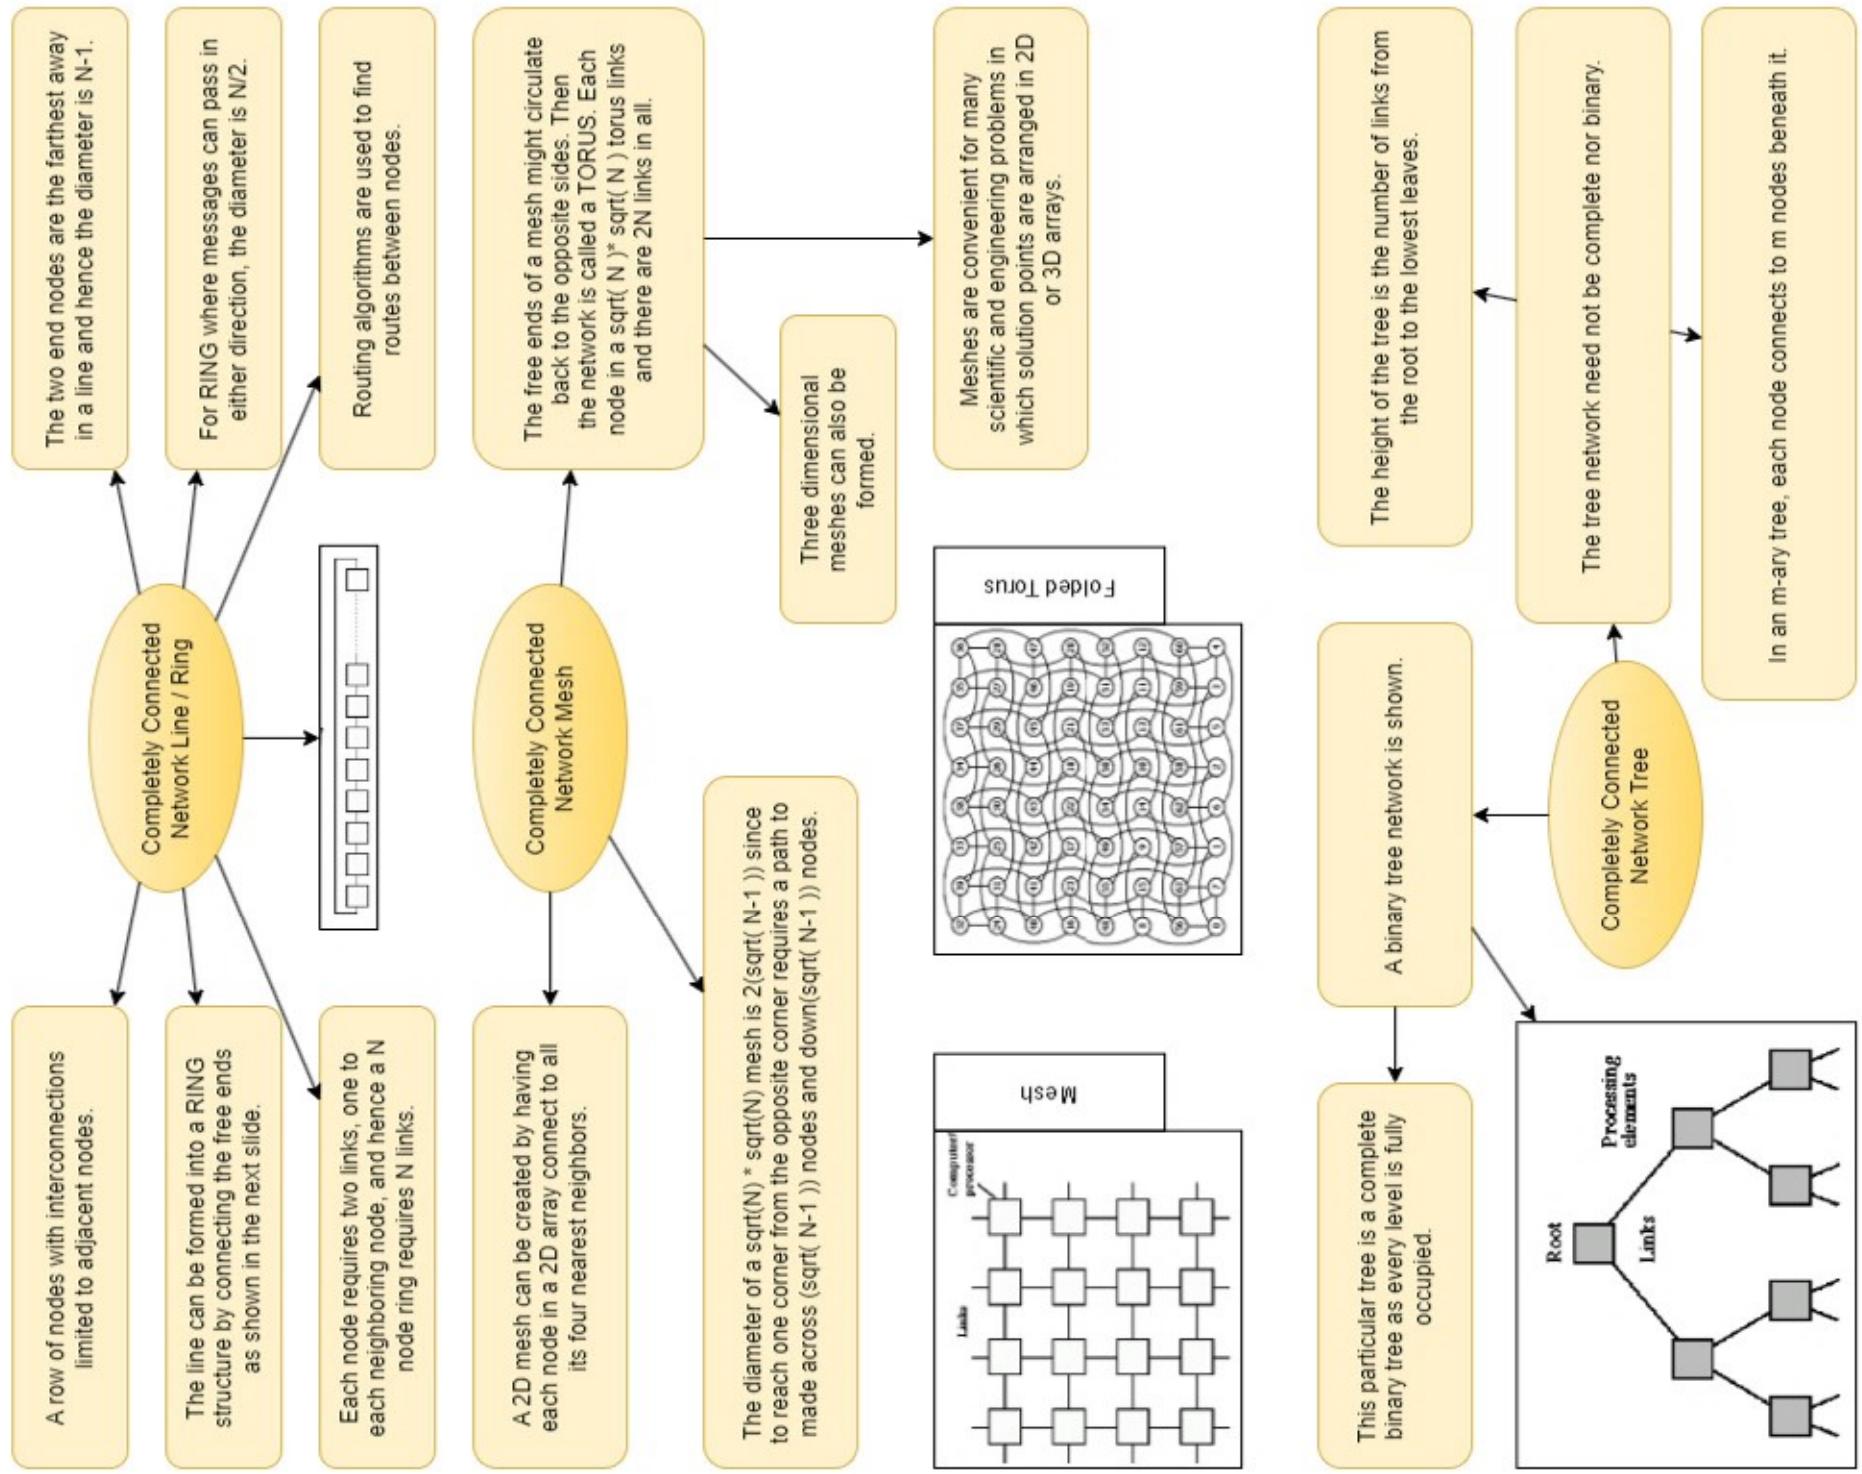
Lecture 2: Parallel Programming Models



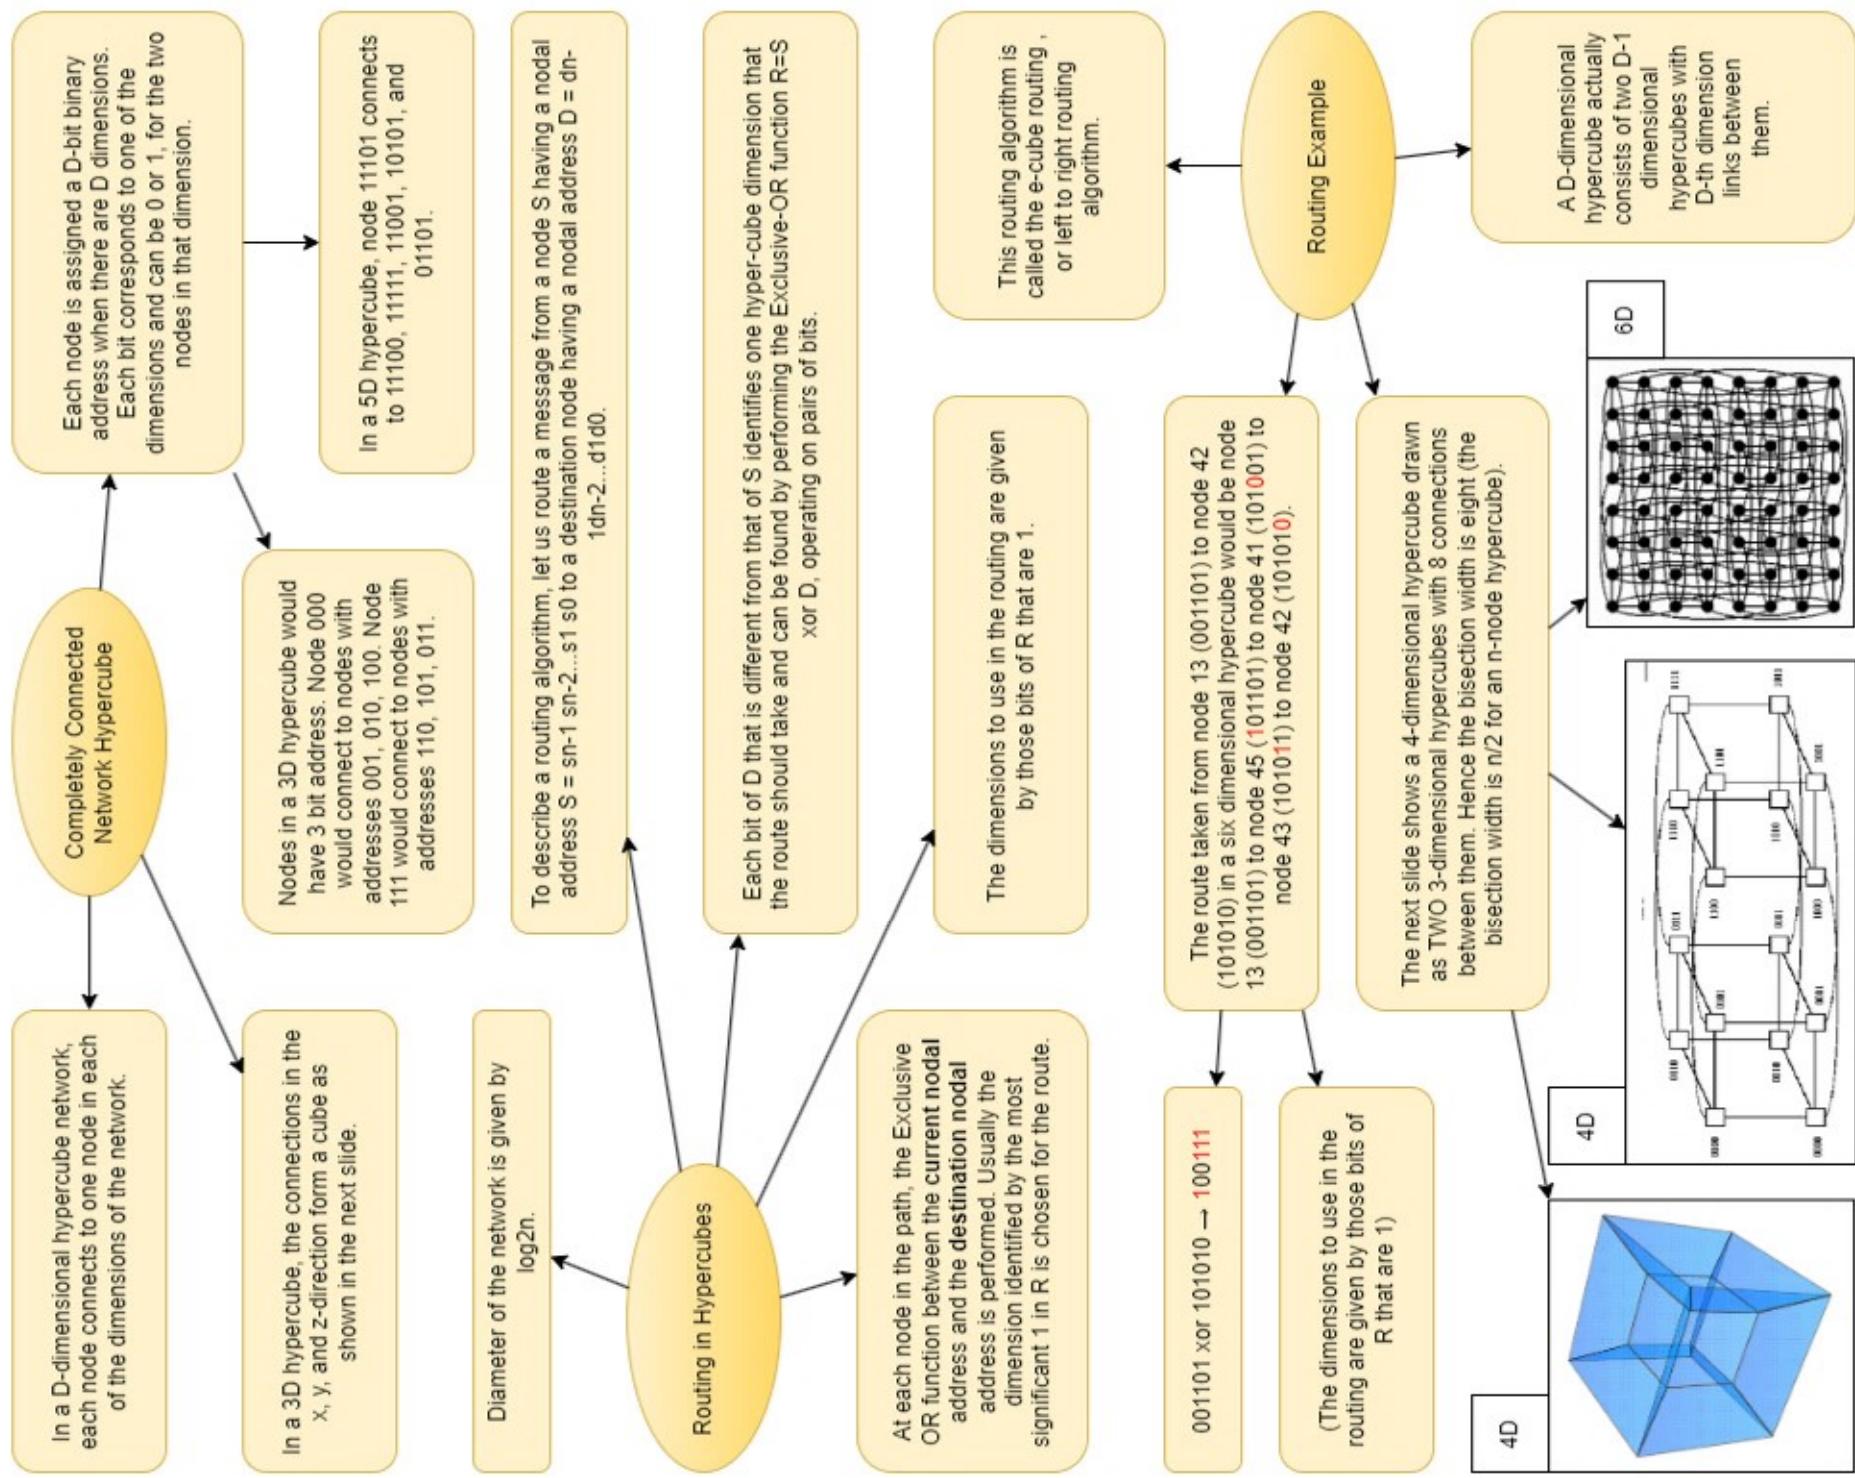
Lecture 2: Interconnection network for multi-computers or message-passing computers



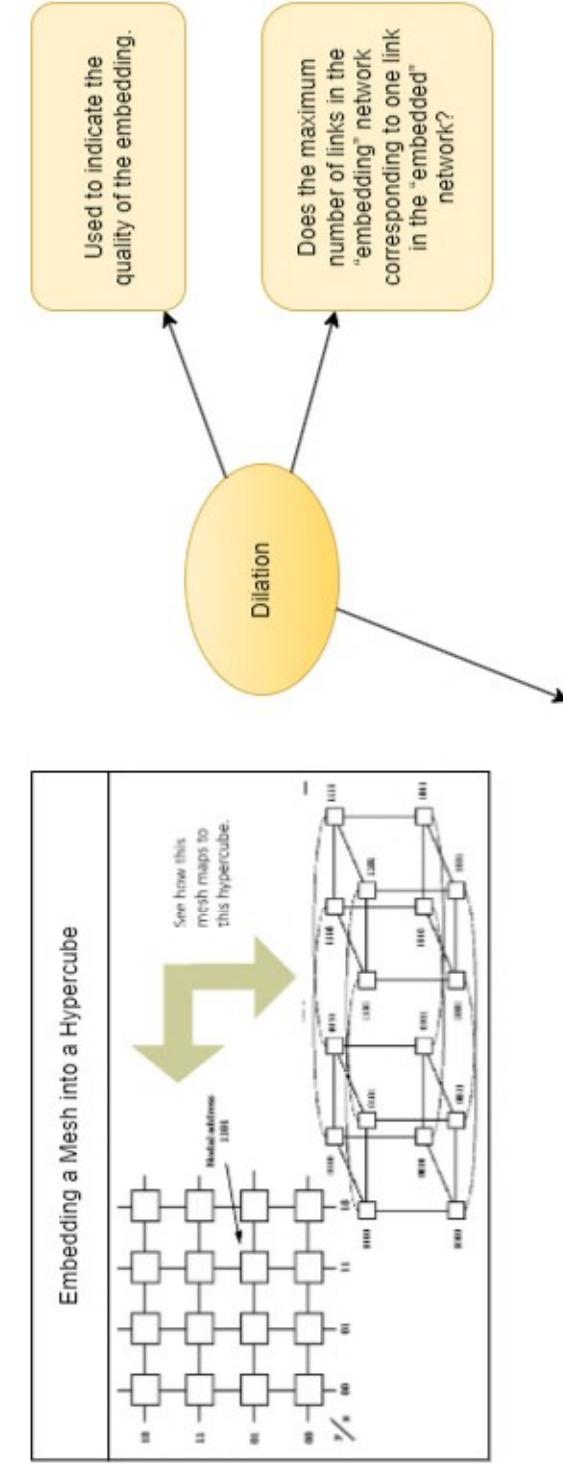
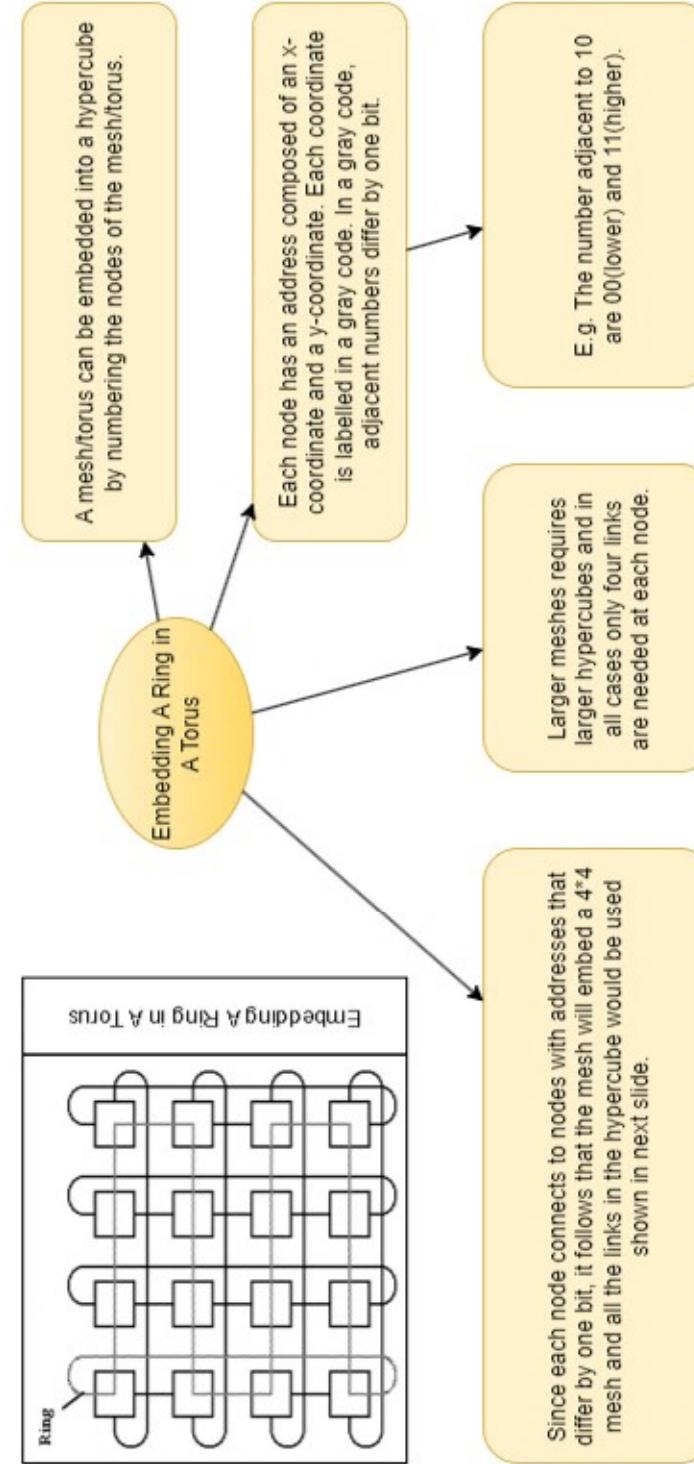
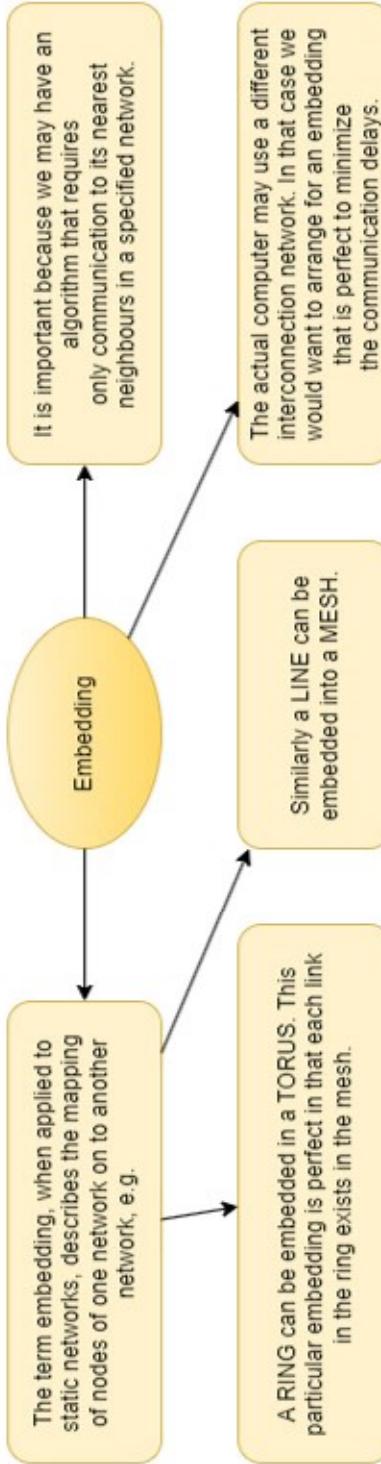
Lecture 2: Interconnection network for multi-computers or message-passing computers



Lecture 2: Interconnection network for multi-computers or message-passing computers

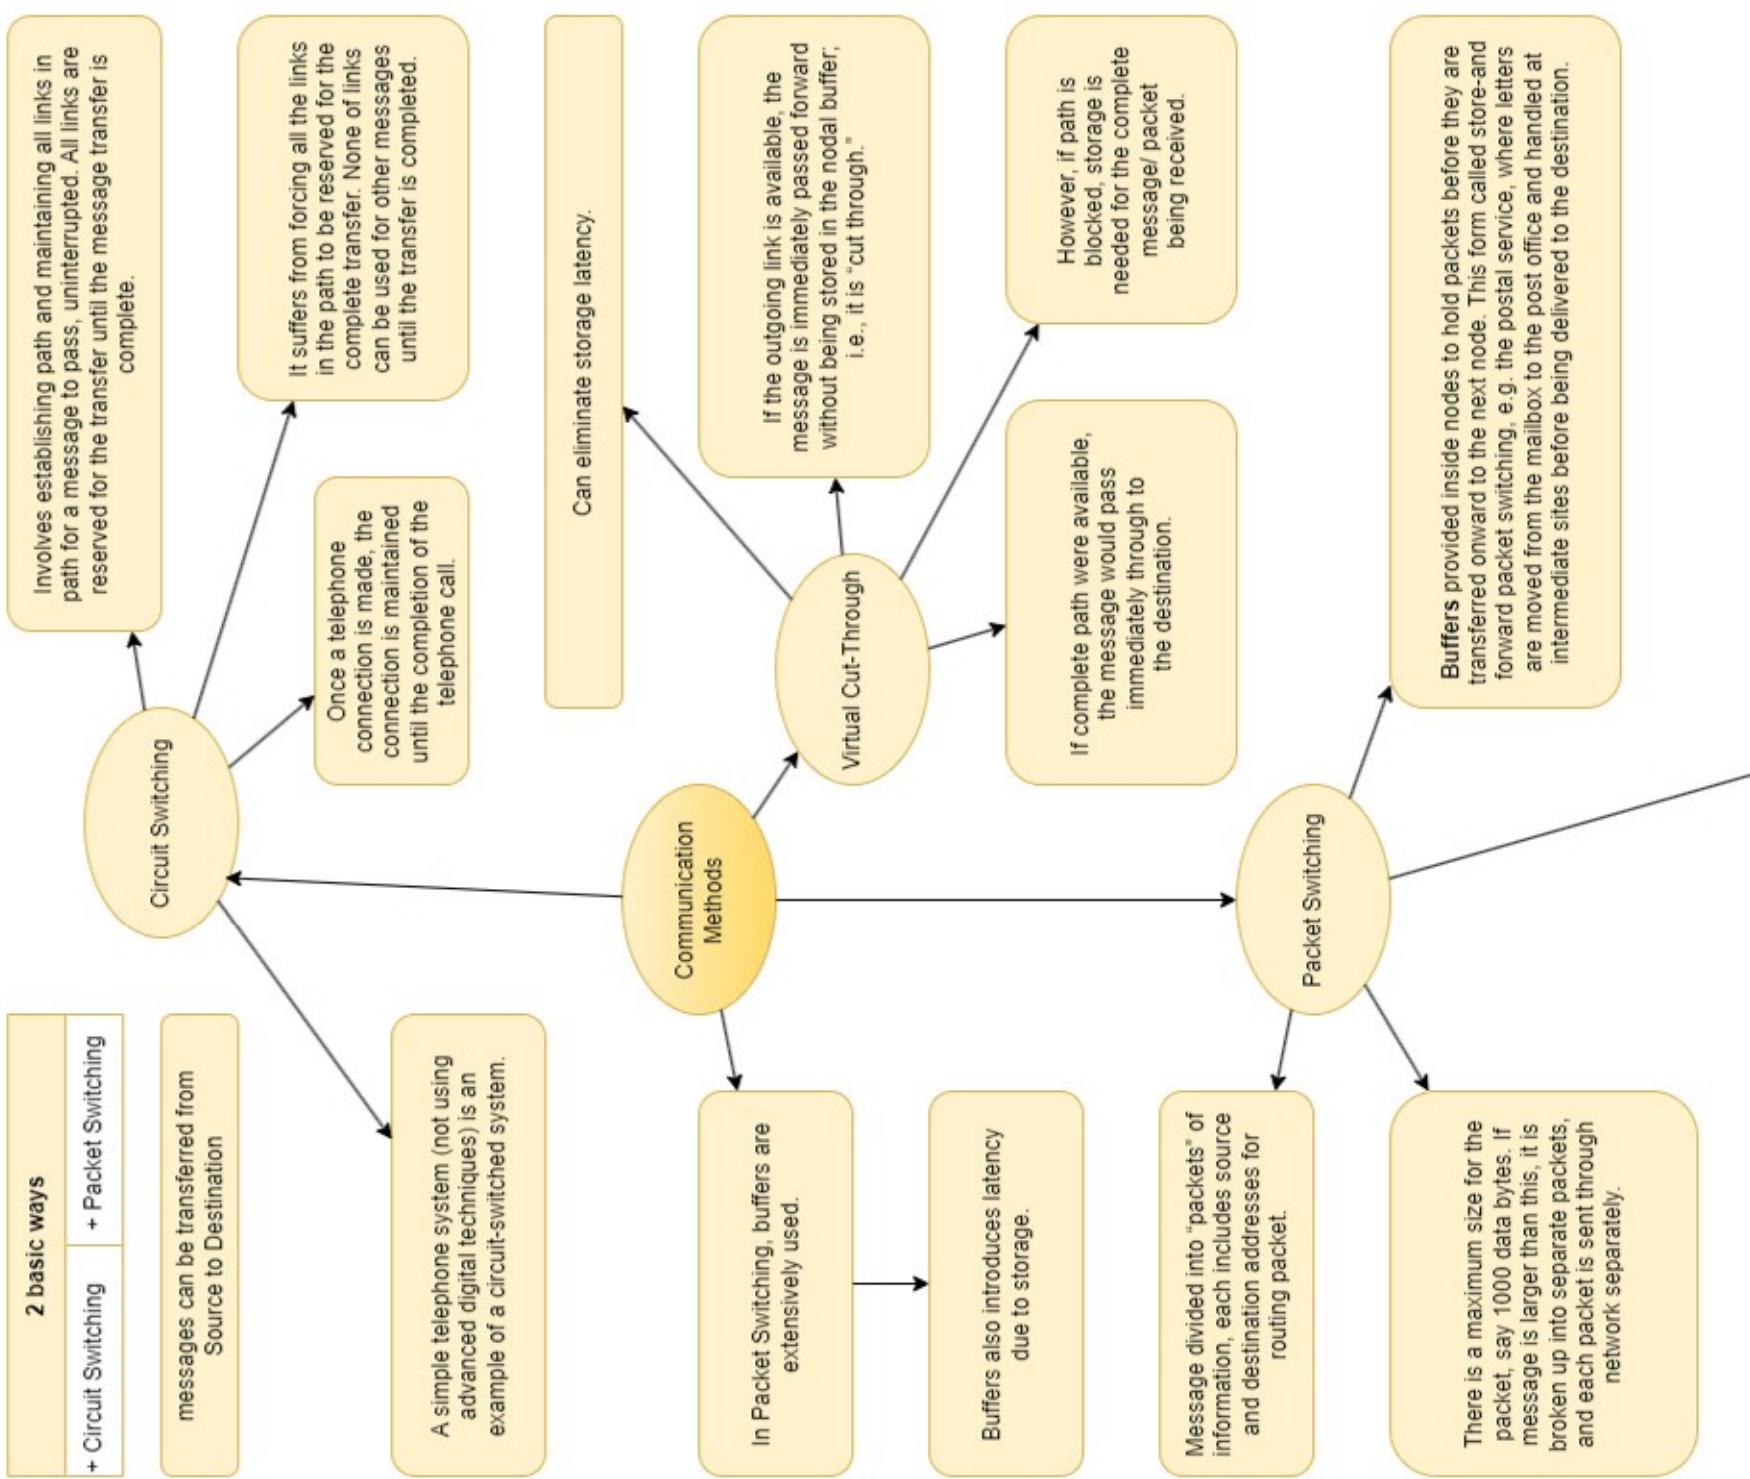


Lecture 2: Interconnection network for multi-computers

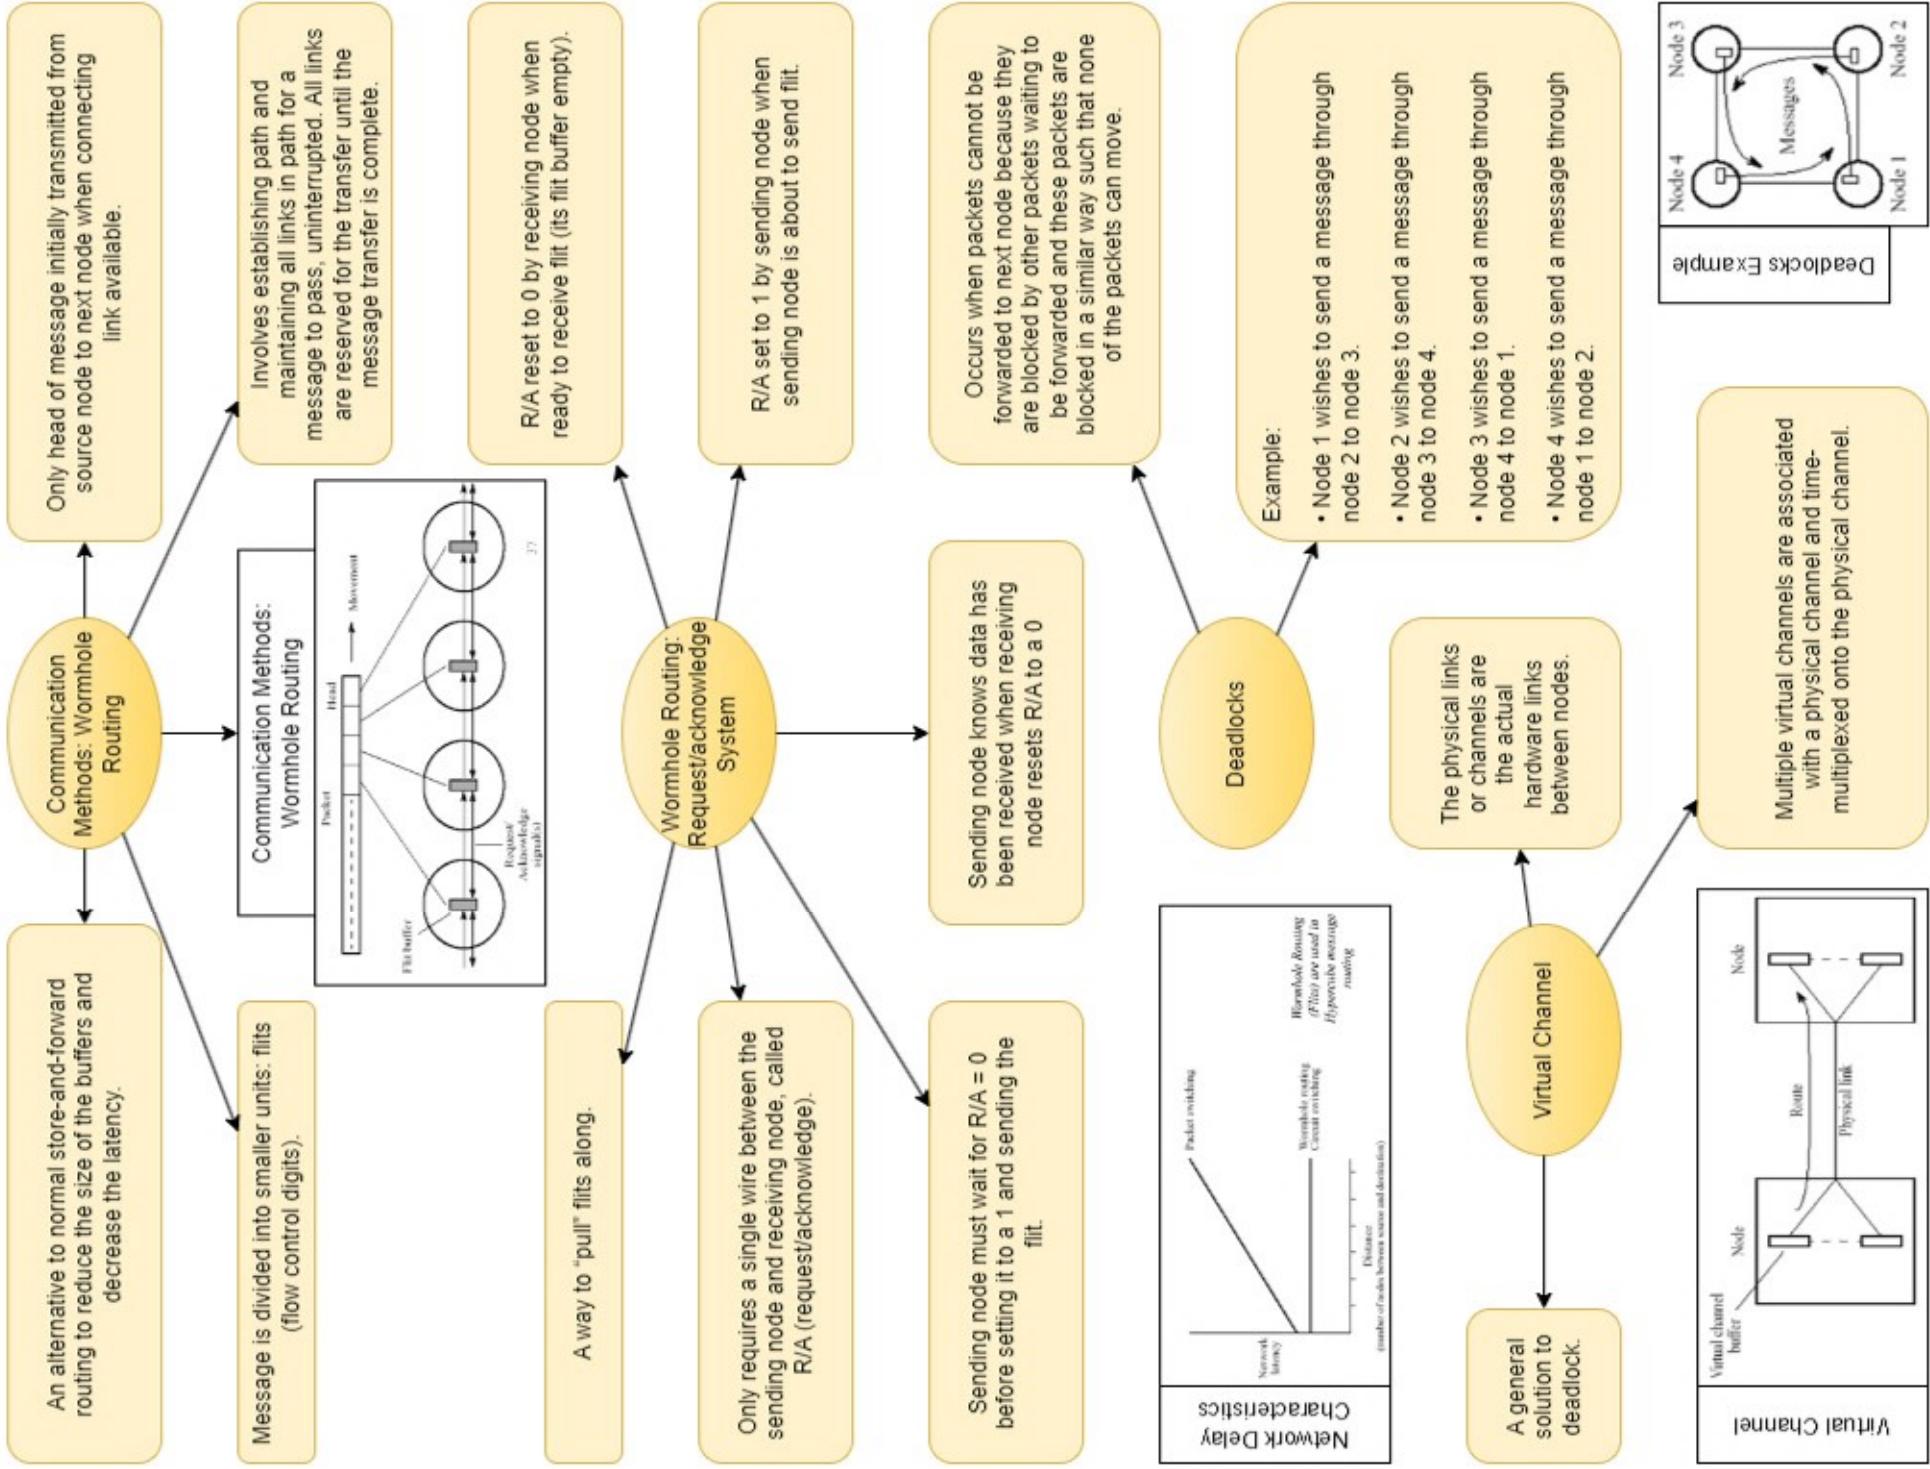


Perfect embeddings, such as a line/ ring into mesh/ torus or a mesh onto a hypercube, have a dilation of 1.

Lecture 2: Communication methods and deadlocks

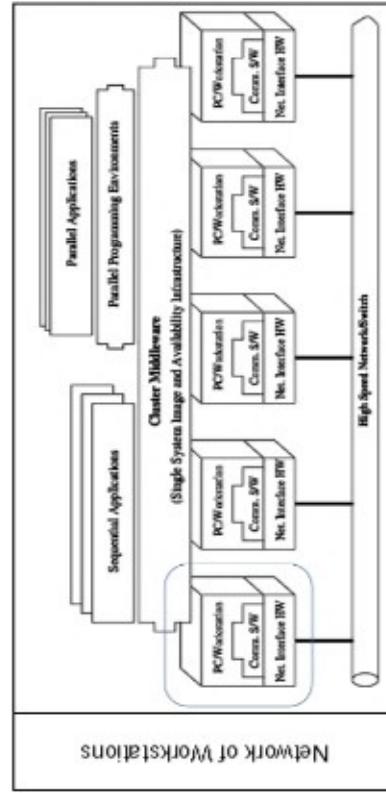
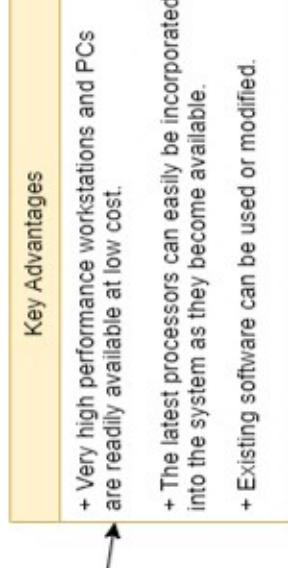


Lecture 2: Communication methods and deadlocks



Lecture 2: Cluster Computing

A cluster of workstations (COWs) or network of workstations (NOWs) offers a very attractive alternative to expensive supercomputers and parallel computer systems for high-performance computing.



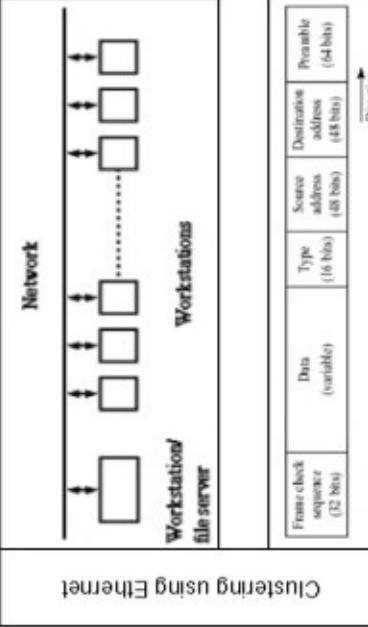
Clustering using Ethernet

All transfers between a source and a destination are carried in packets (or frames) along the single wire, one bit at a time. The packet carries the source address, the destination address, and the data.

Common communication network for workstations.

Message-Passing Interface (MPI)

+ Standard was defined in 1990s.



Ring Structures

Examples

- + Token rings/FDDI networks
- + Star-connected networks

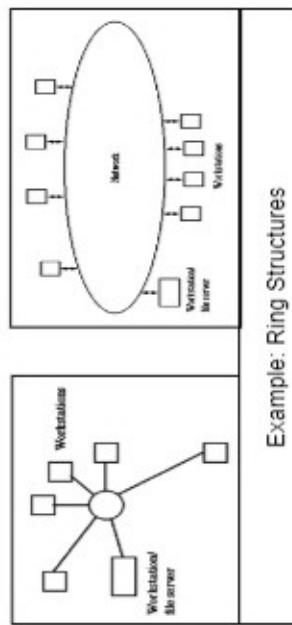
Other forms of networks are also in wide use for work stations.

Provides the highest interconnection bandwidth.

Point-to-Point Communication

Each cluster could be Ethernet, ring or star-connected.

Individual workstations make request to the fileserver (or to each other) on the network.

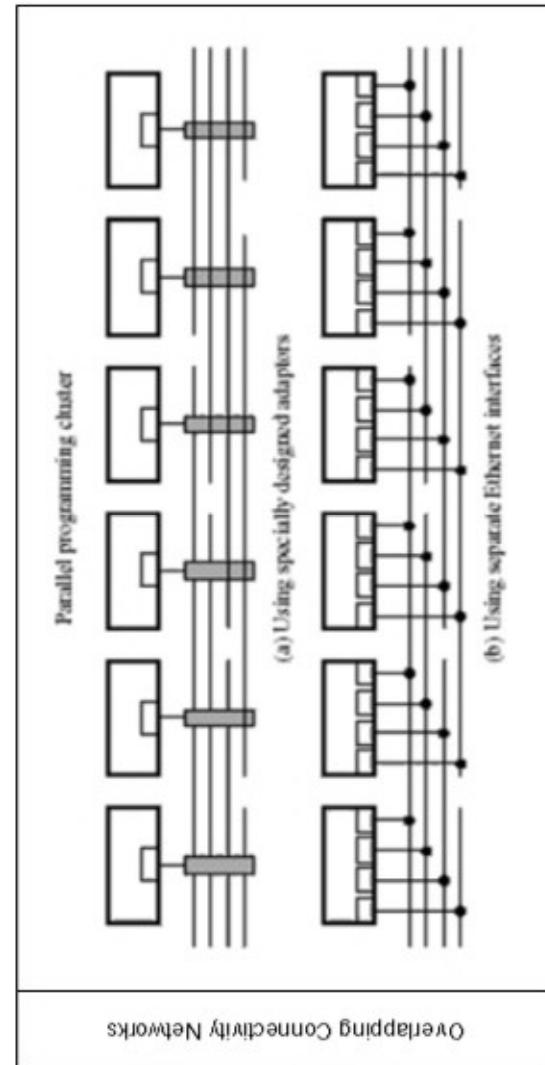
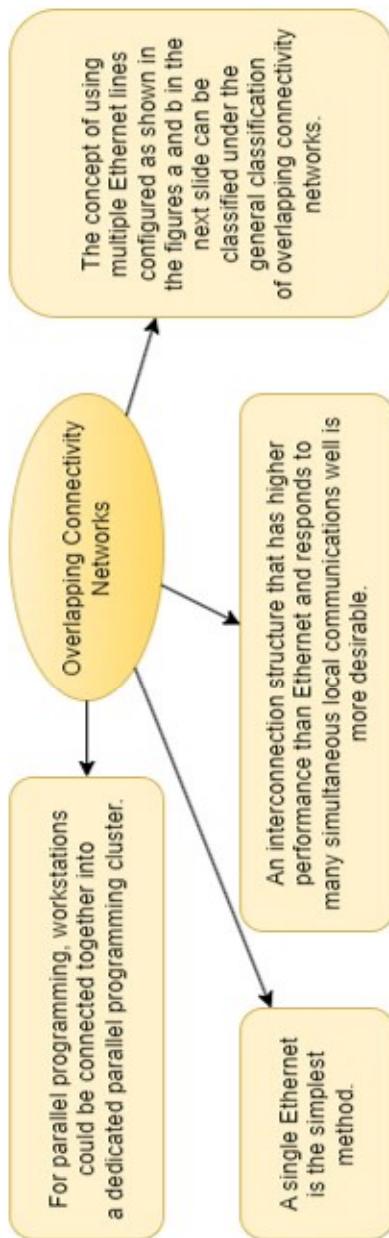


Example: Ring Structures

Various point-to-point configurations can be created using hubs and switches. Examples:

High Performance Parallel Interface (HPI), Fast (100 MHz) and Gigabit Ethernet, and fiber optics.

Lecture 2: Cluster Computing



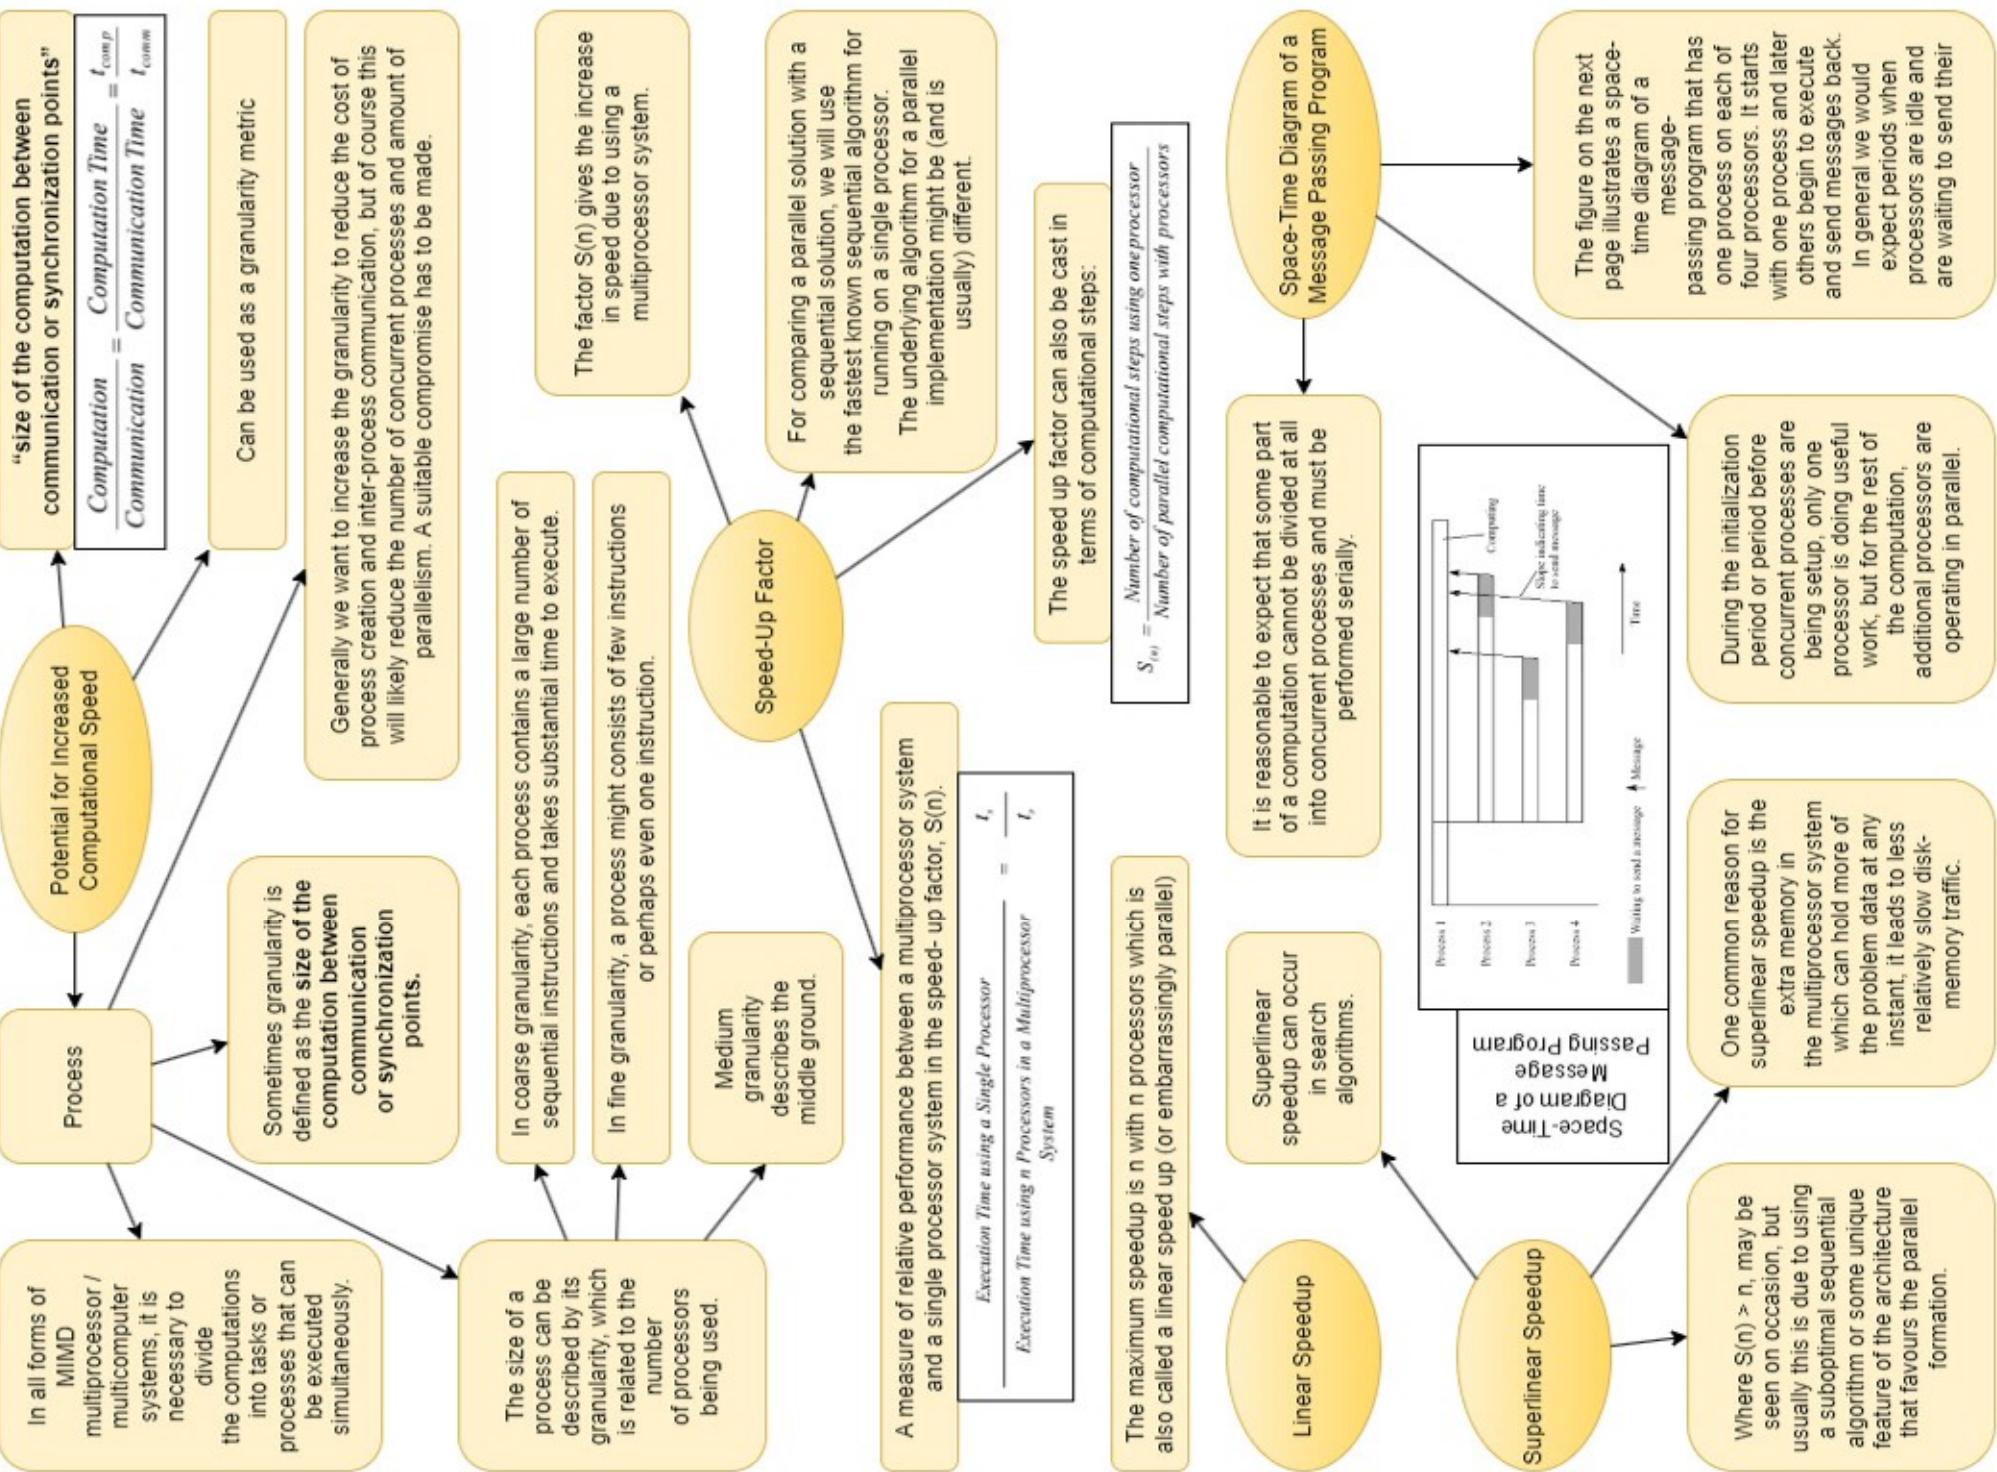
Overlapping connectivity networks have the characteristic that regions of connectivity are provided and the regions overlap.

Overlapping Connectivity Networks

There are several ways that overlapping connectivity can be achieved. An example using multi-tiered Ethernet switches.

Parallel Computing	Distributed Processing
PEs are clustered together (tightly coupled)	Processing nodes are geographically distributed
Requires a fast communication network	Can do without a fast communication network
Nodes are homogeneous in general	Heterogeneous nodes are very common
Motivation: Fast processing	Information/Resource sharing and fault tolerance
May use shared memory	Memory is always distributed

Lecture 3: Computational Speed



Lecture 3: Computational Speed

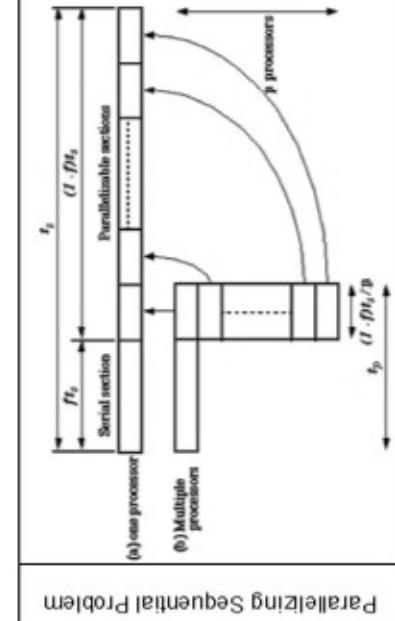
Illustrated is the case with a single serial part at the beginning of the computation which cannot be parallelized, but the remaining parts could be distributed throughout the computation.

Maximum Speedup

The speed up factor $S(p)$ is therefore given by

$$S(p) = \frac{t_s + (1-f)t_s/p}{f t_s + (1-f)t_s/p} = \frac{p}{1+(p-1)f}$$

If the fraction of the computation that cannot be divided into concurrent tasks is f , the time to perform the computation with p processors is given by $t_s + (1-f)t_s/p$ as shown in the next slide.



$S(p) = \frac{t_s + (1-f)t_s/p}{f t_s + (1-f)t_s/p} = \frac{p}{1+(p-1)f}$	Amdahl's Law
$S(p) = \frac{p}{1+(p-1)f}$	Amdahl's Law

It shows $S(n)$ plotted against number of processors and against f .

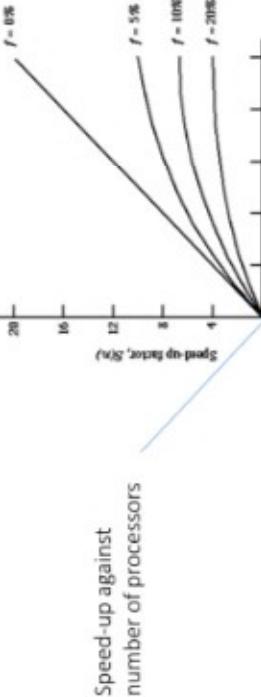
If a significant increase in speed of computation is to be achieved, the fraction of computation that is executed by concurrent processes needs to be substantial fraction of the overall computation.

Even with infinite number of processors, maximum speedup limited to $1/f$.

i.e., $S(n) = 1/f$
 $n \rightarrow \infty$

Amdahl's Law

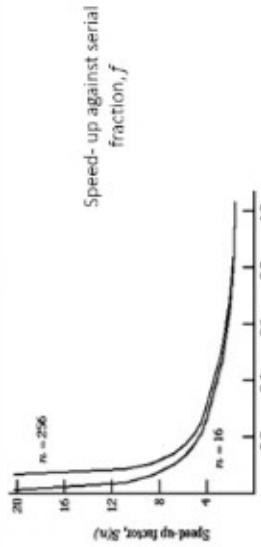
Amdahl's Law



$$E = \frac{S(n)}{n} \times 100\%$$

Efficiency gives fraction of time that processors are being used on computation.

Efficiency



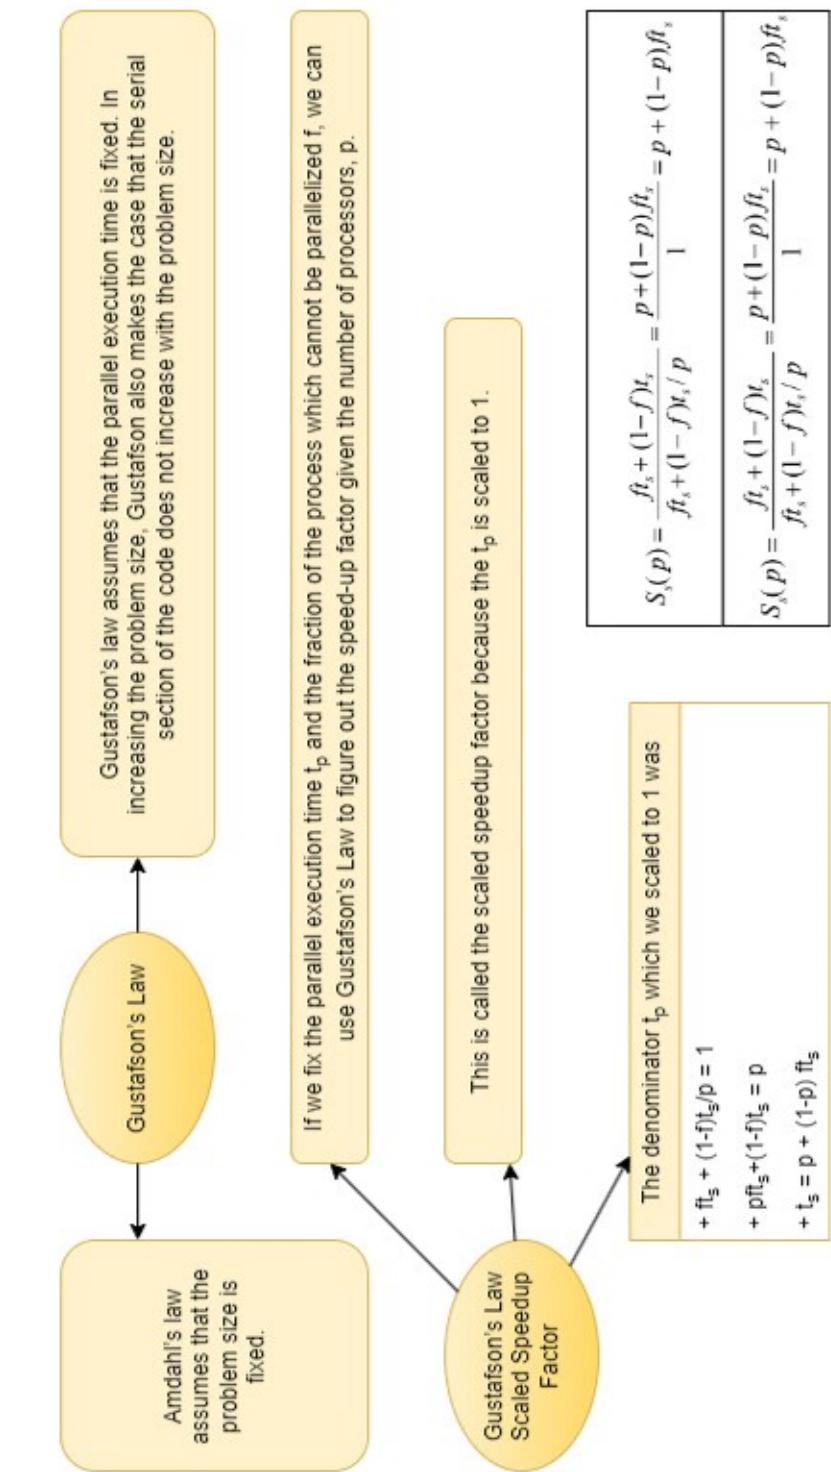
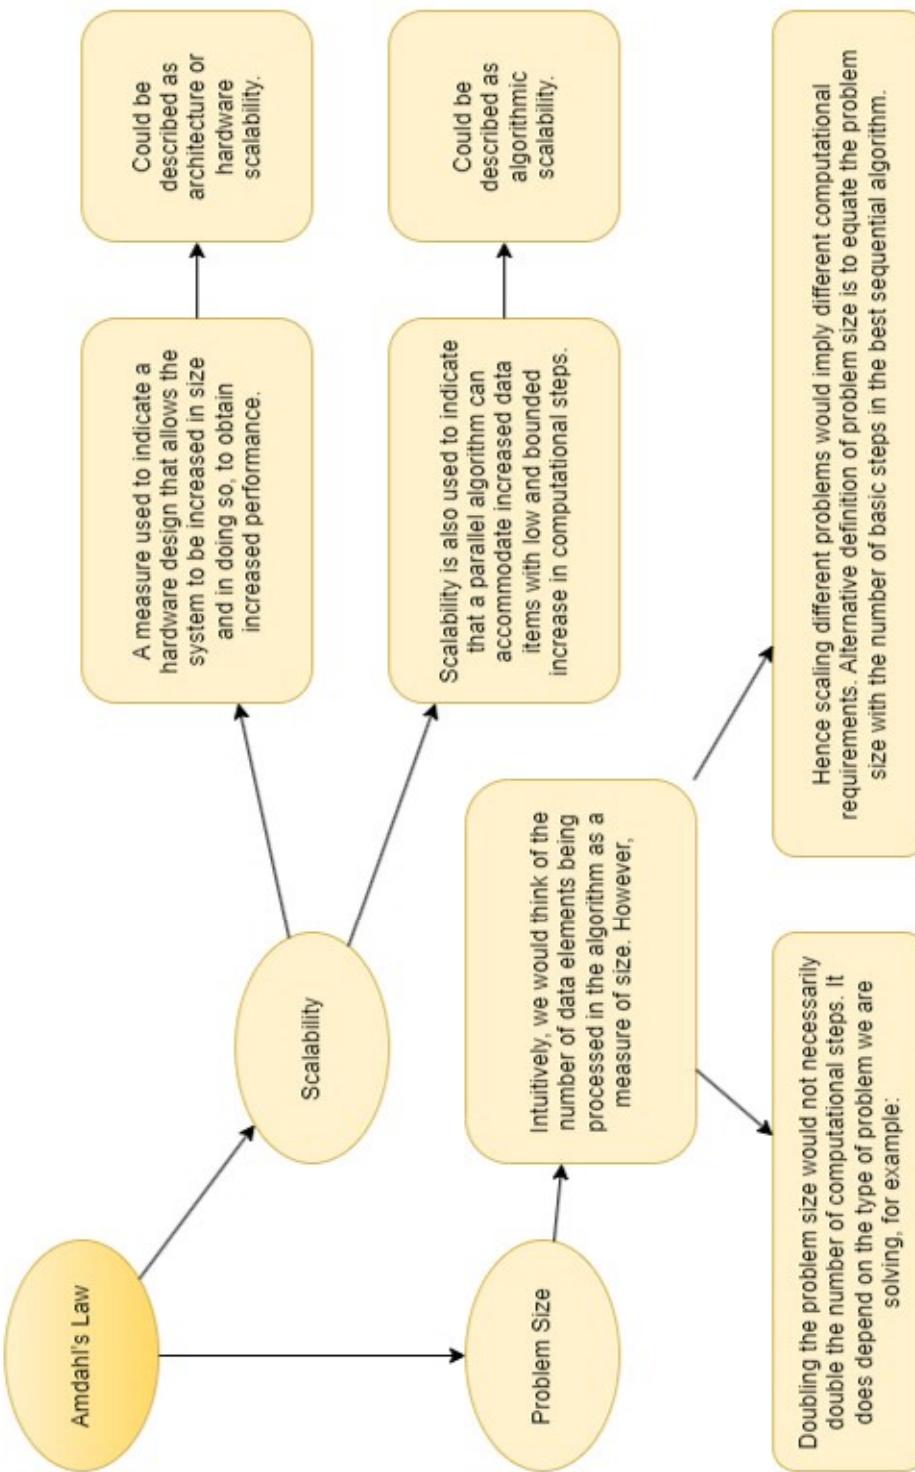
The processor-time product or cost (or work) of a computation defined as
 $\text{Cost} = (\text{execution time}) \times (\text{total number of processors used})$
The cost of a sequential computation is simply its execution time, t_s . The cost of a parallel computation is $t_p \times n$. The parallel execution time, t_p is given by $t_p/S(n)$.

Cost-Optimal Parallel Algorithm

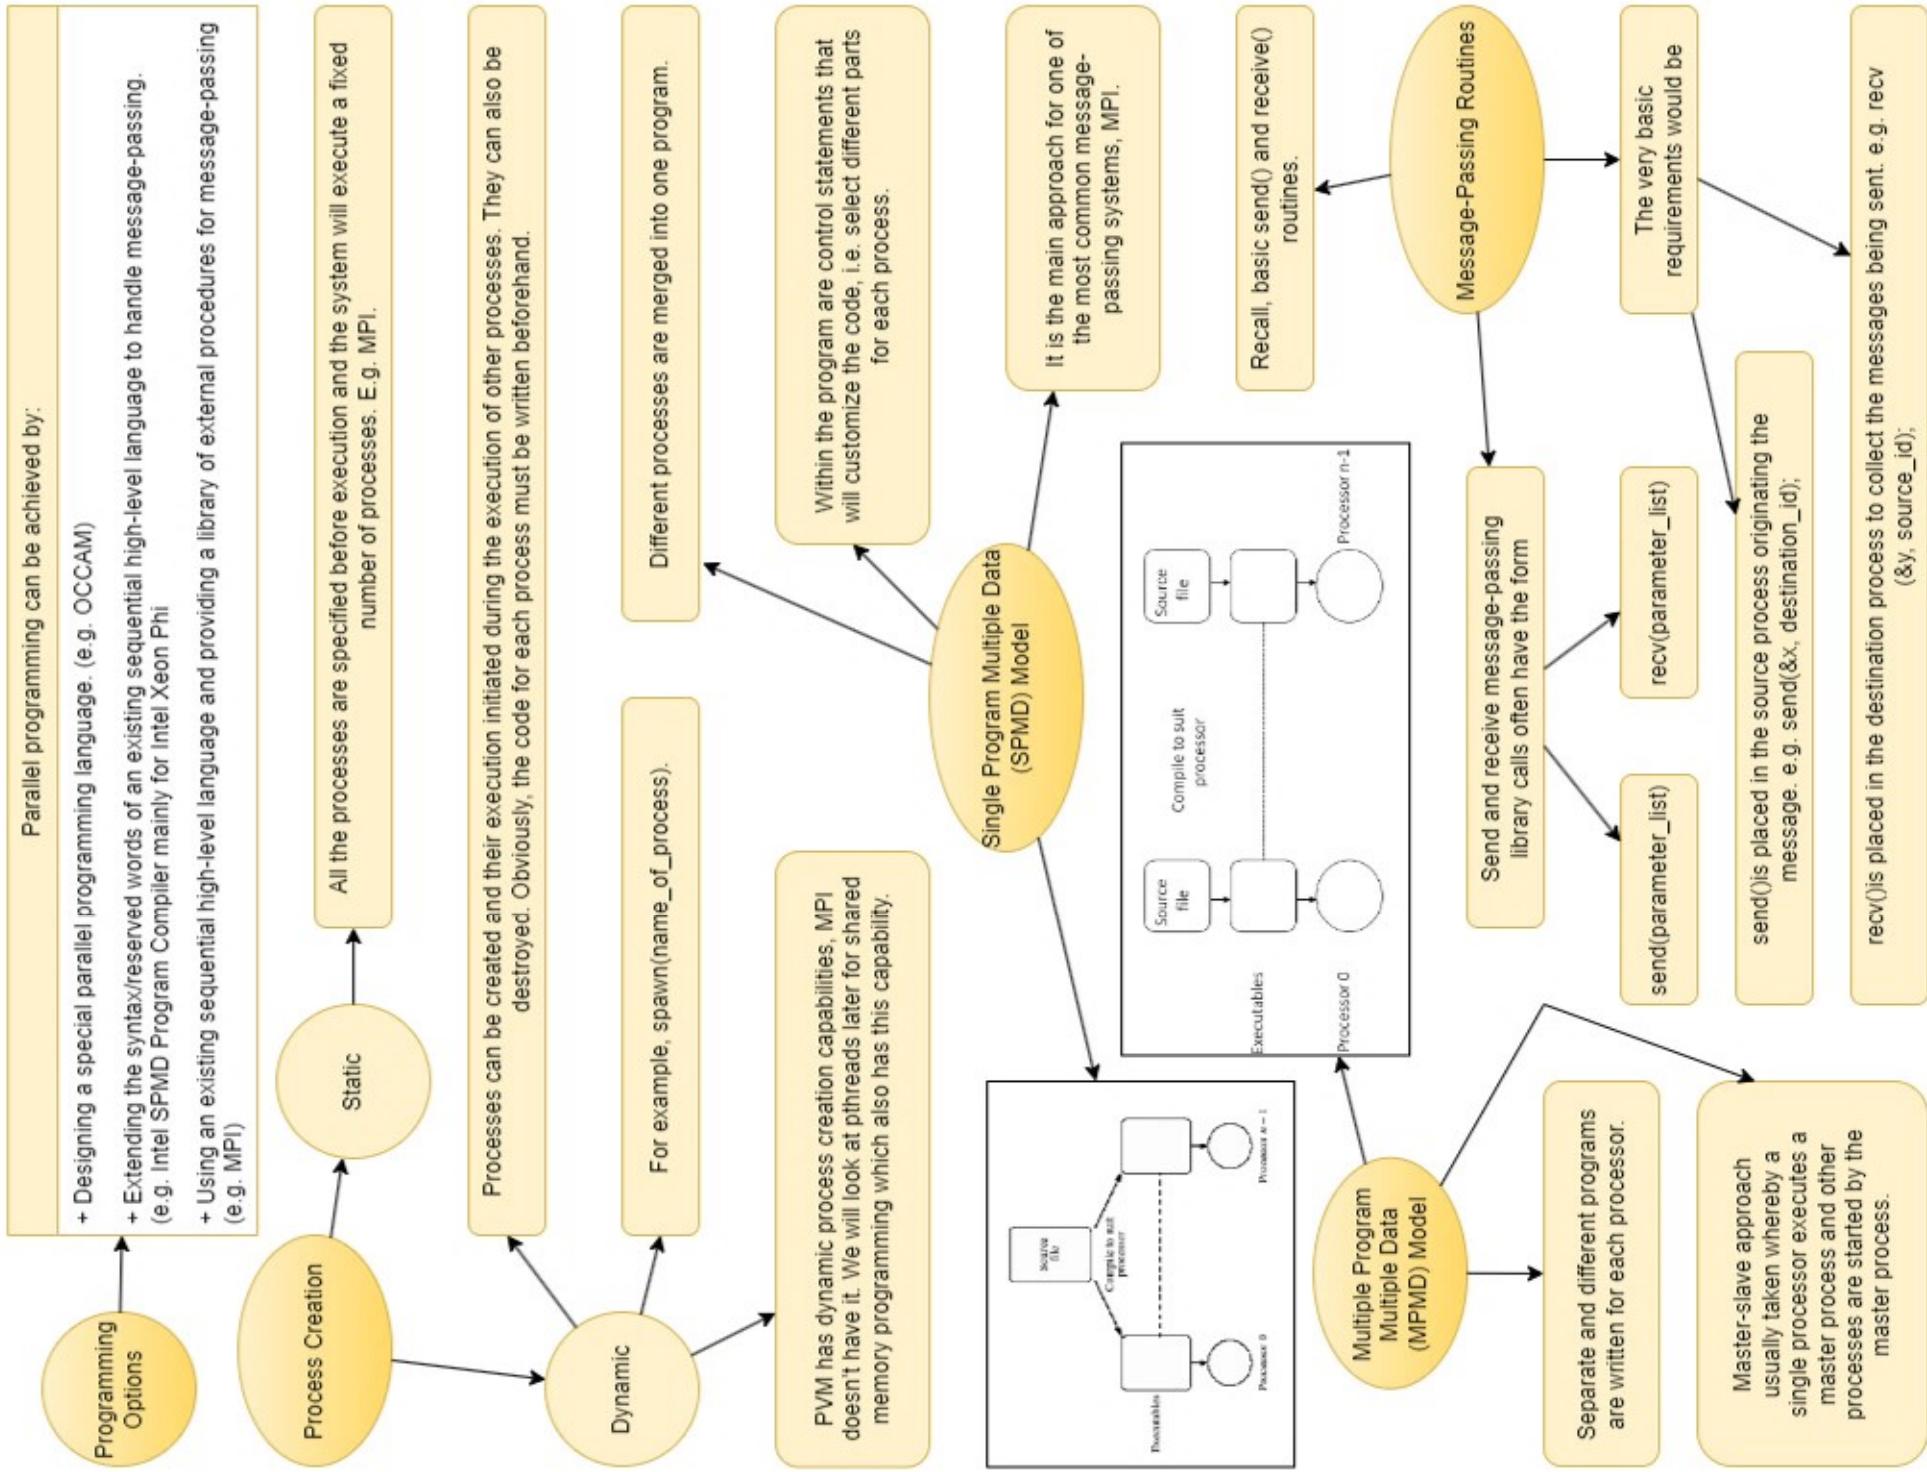
One in which the cost to solve a problem on a multiprocessor is proportional to the cost (i.e., execution time) on a single processor system.

Cost

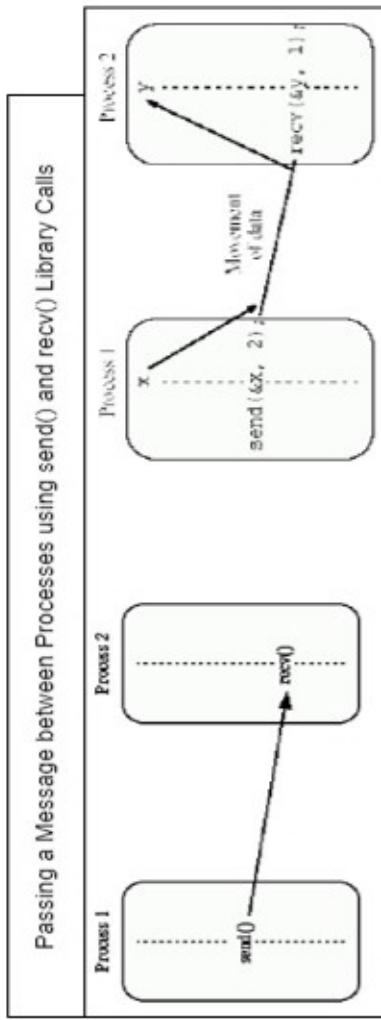
Lecture 3: Computational Speed



Lecture 3: Analytical Models for Parallel Processing

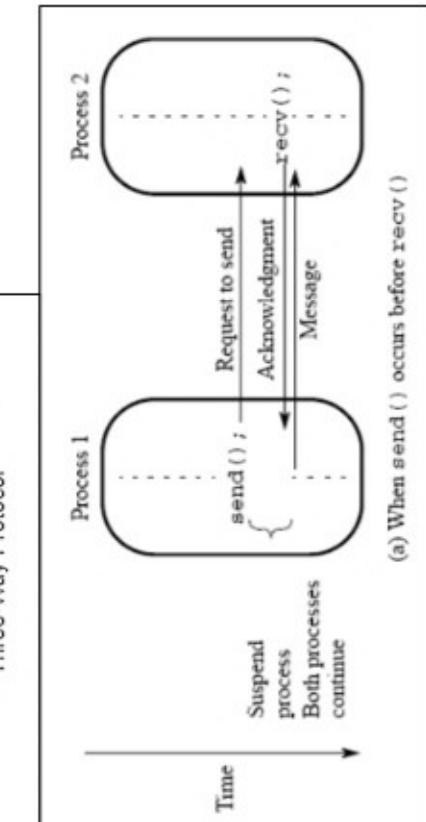


Lecture 3: Analytical Models for Parallel Processing



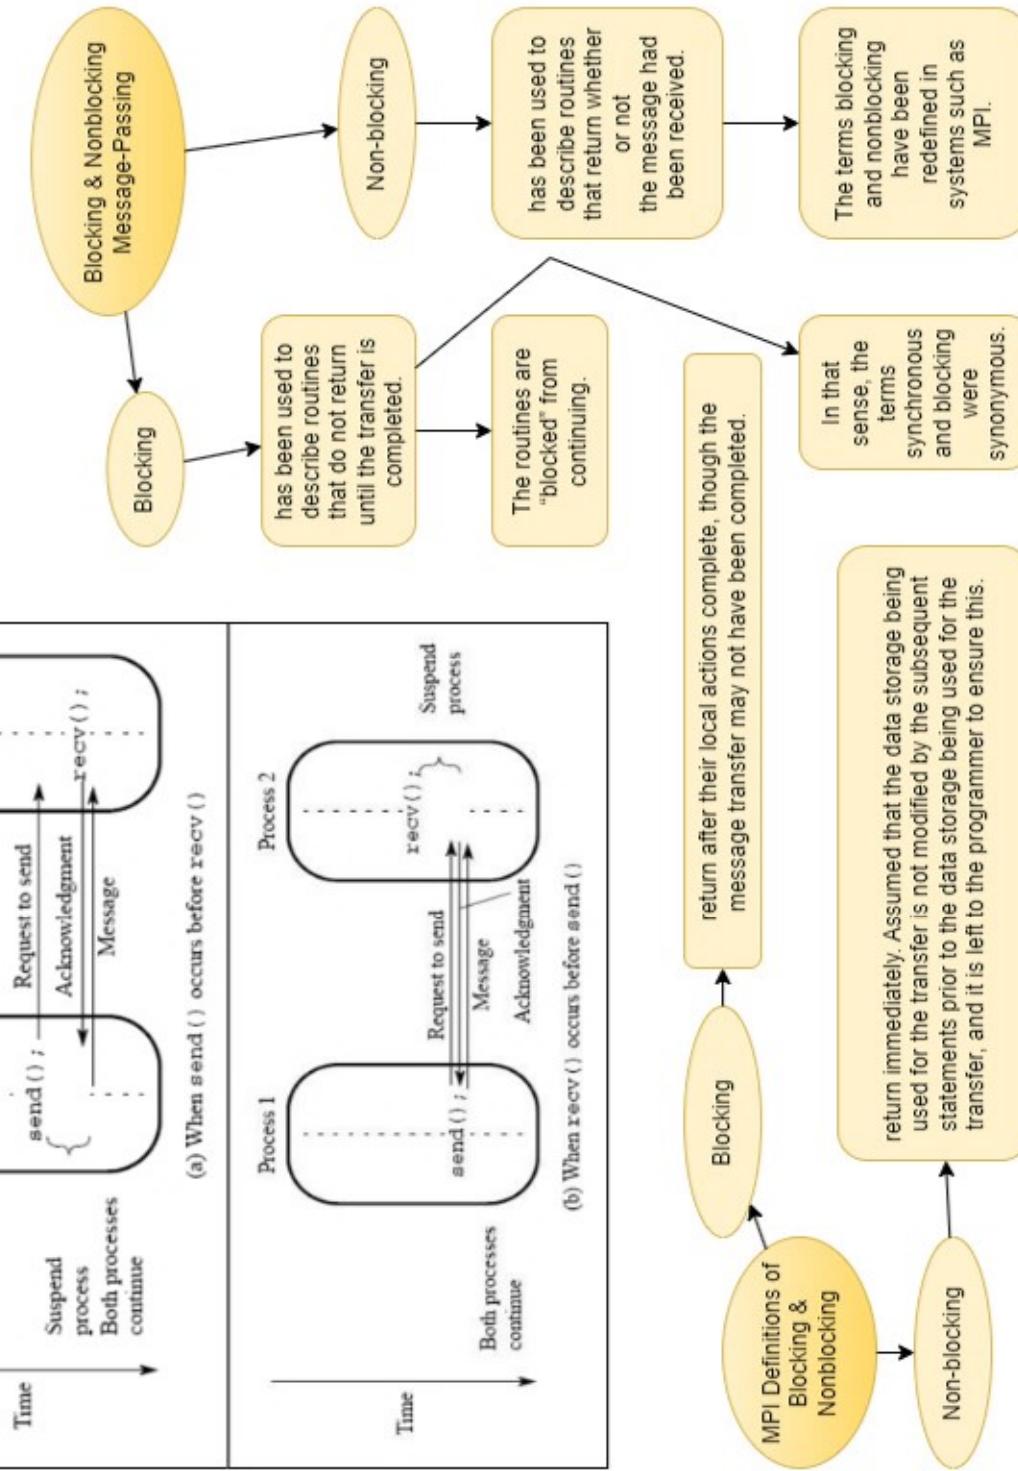
Do not need message buffer storage. A synchronous send routine could wait until the complete message can be accepted by the receiving process before sending the message.

Synchronous send() and recv() Library Calls using a Three-Way Protocol

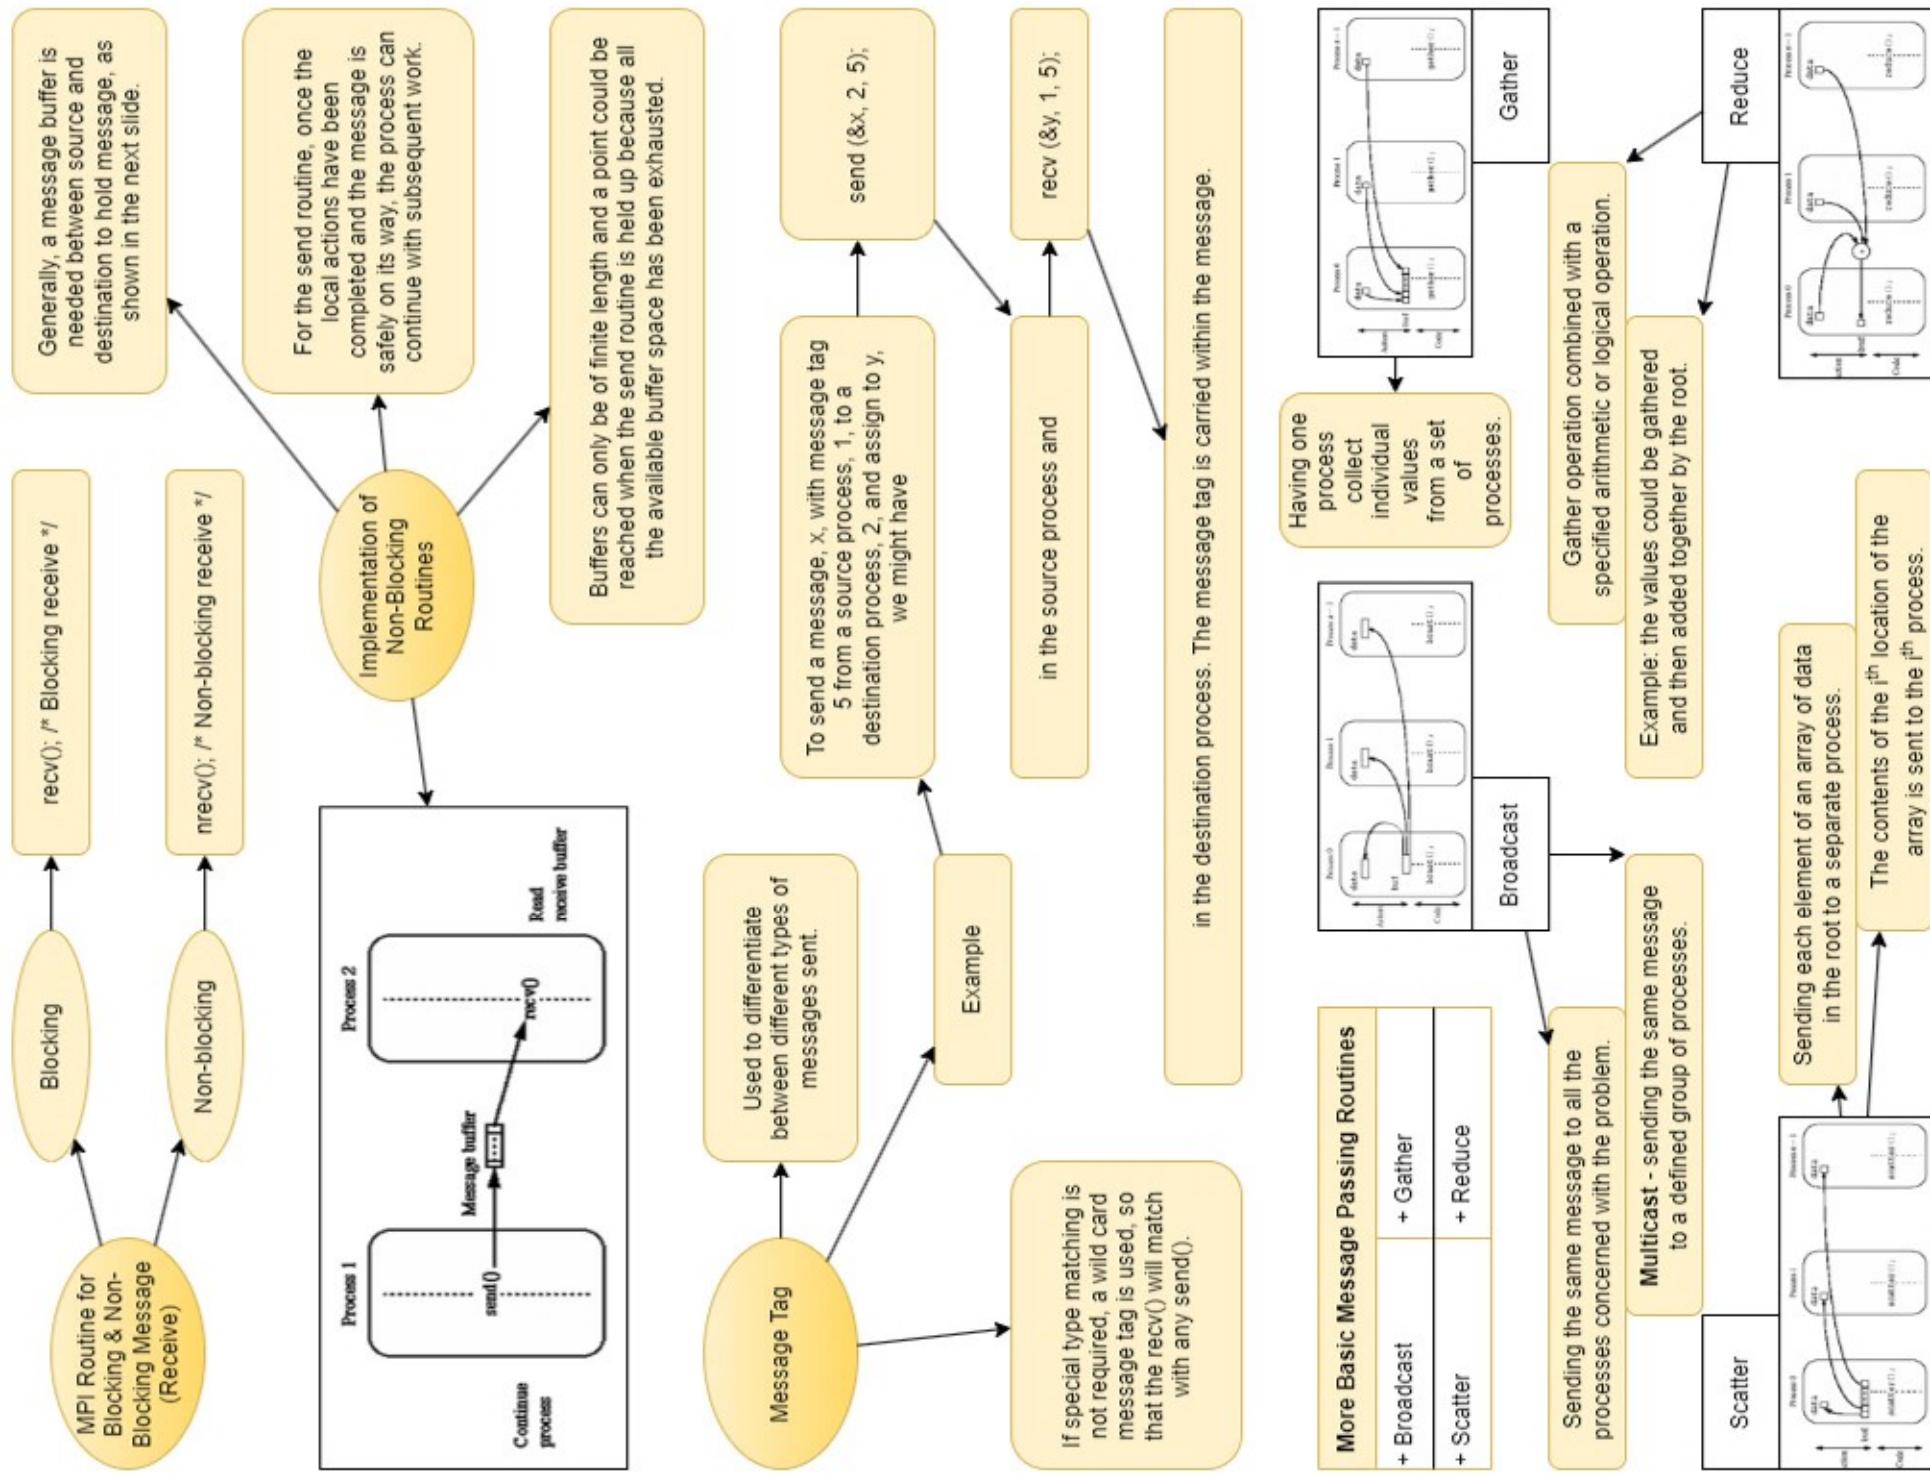


A synchronous receive routine will wait until the message it is expecting arrives.

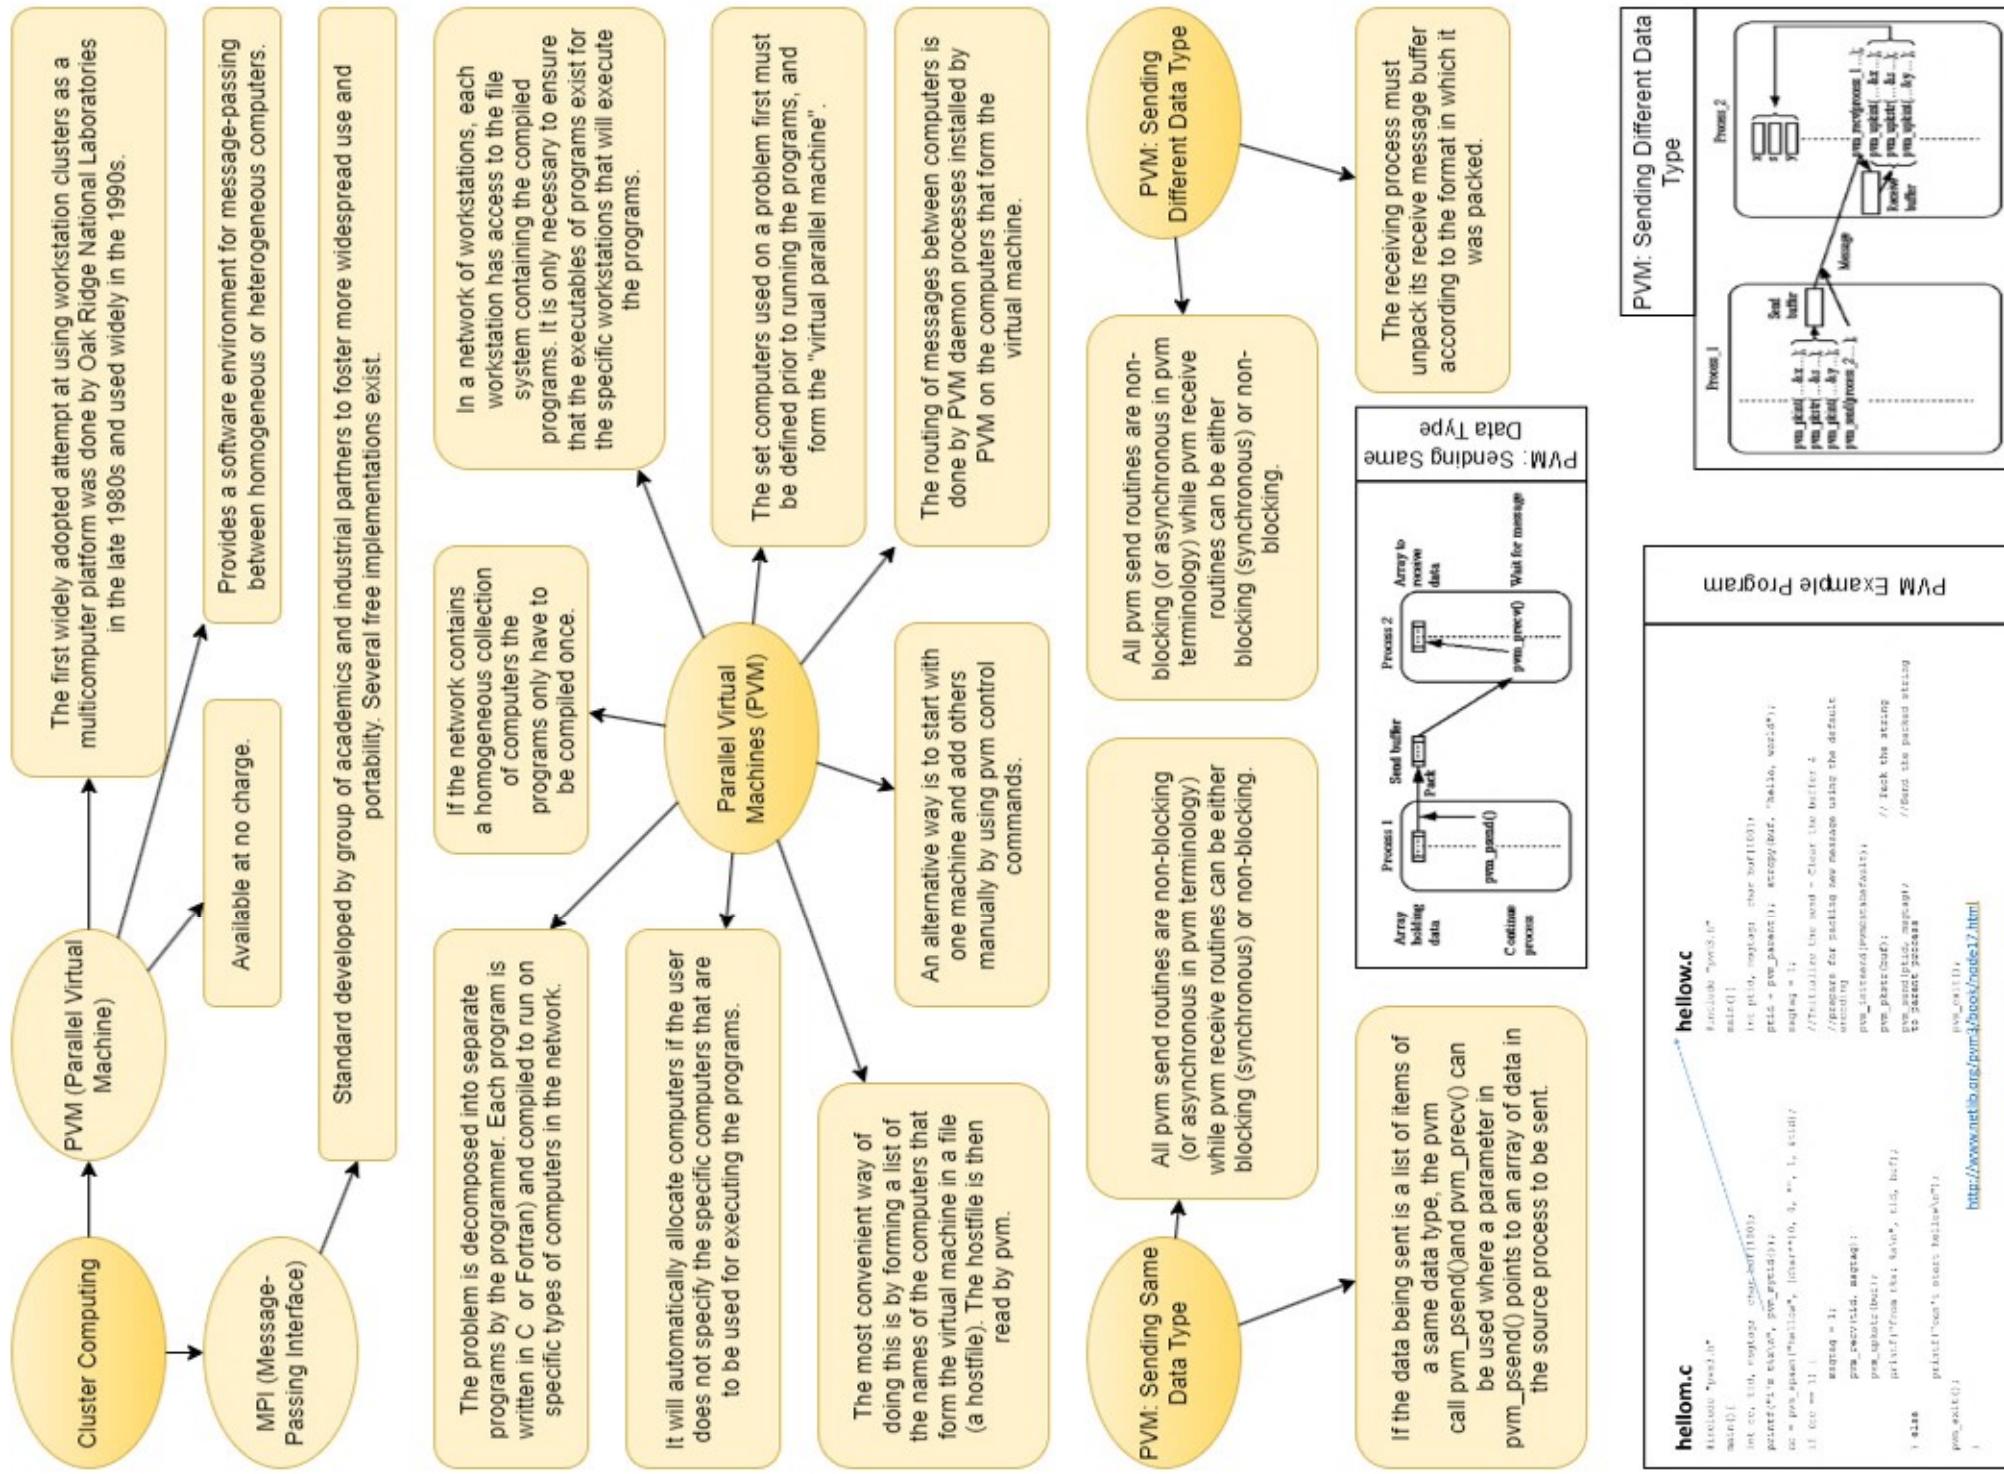
Synchronous routines intrinsically perform two actions:
+ They transfer data and they synchronize processes.
+ Suggests some form of signalling, such as a three-way protocol as shown in next slide.



Lecture 3: Analytical Models for Parallel Processing



Lecture 4: PVM



Lecture 4: MPI

History of MPI

- + It was developed in 1993-94 by a group called the MPI Forum, composed of major IT corporations and research centres to standardize the way that parallel processing is done.
- + Still actively being enhanced, version 4.0 is expected to be released soon.

MPI is an "industrial strength" library with literally hundreds of different routines.

```
Routine          MPI_Bcast (message, count, datatype, root, communicator) {
    MPI_Finalize();
    MPI_Finalize();
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
    MPI_Send(count, type, dest, tag, MPI_COMM_WORLD, &status);
    MPI_Recv(message, count, type, source, tag, MPI_COMM_WORLD, &status);
    MPI_TypeDestroy();
    MPI_TypeDestroy();
}
```

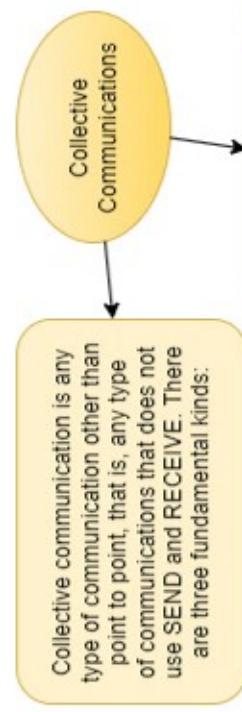
Lecture 4: Further Reference

Goals of MPI

- + Provide source code portability,
- + Allow an efficient range of implementations across many architectures.

Features

- + Provide a great deal of functionality, and
- + Support heterogeneous parallelism, that is, the ability to have different types of processors communicate with each other.



These routines allow you to implement all of the collective communications you will need for virtually any parallel algorithm.

Operation Name	Meaning
HPI_BCAST	Message broadcast from root to all others
HPI_ALLREDUCE	Allreduce operation
HPI_ALLGATHER	Allgather operation
HPI_ALLGATHERV	Allgather with vector operation
HPI_ALLTOALL	Alltoall operation
HPI_ALLTOALLV	Alltoall with vector operation
HPI_GATHER	Gather operation
HPI_GATHERV	Gather with vector operation
HPI_SCATTER	Scatter operation
HPI_SCATTERV	Scatter with vector operation
HPI_WORLDD	World communicator
HPI_COMM_WORLD	World communicator

MPI_Bcast()	<ul style="list-style-type: none"> • MPI_Bcast is message, count, datatype, root, communicator; • One process, called the root, sends the same message to every other process in the specified communicator. So, a broadcast is a 1-to-many transmission. • Every process must execute the same command, but the root is the sender, and everyone else is the receiver. • The indicated message, containing count value, each of type datatype is then sent from the root to every other process in the communicator. • If one of the processes in the communicator does not execute the broadcast, it will result in a deadlock condition. • For example, to broadcast the floating point value x from process 0 to every other process, you would execute the following command:
	<pre>MPIT2_Bcast(&x, 1, MPI_FLOAT, 0, MPI_COMM_WORLD);</pre> <ul style="list-style-type: none"> • When every process has executed this command, the broadcast will take place and every process will have the value of the variable x stored in its local memory

MPI_Reduce()	<ul style="list-style-type: none"> • MPI_Reduce (data_value, &global_value, count, datatype, reduction operator, root, Communicator); • Reduction is the reduction of many values to a single value using a single binary operator. • Now the data value, containing count values of type datatype, is sent by every node in the communicator to the indicated root processor and combined there into a single global value using the specified reduction operator. • For example, if every processor had a floating point value called x, and we wanted to sum them and store them in the variable called xsum on processor number 4, we would do the following:
	<pre>MPIT2_Reduce(&x, &xsum, 1, MPI_FLOAT, 4, MPI_COMM_WORLD);</pre>

MPI_Allreduce()	<ul style="list-style-type: none"> • There is a variant of Reduce that will iterate on the reduction operation, using each node as the root. • That is, we will reduce X to Xsum at node 0. Then we will reduce Xsum at node 1. Then we will reduce X to Xsum at node 2. This variant is called All_reduce. • The calling sequence is the same except that we eliminate the root field, since every node will function as the root. • All reduce is identical to doing P simultaneous reductions. • It is an example of a many → many collective communication.

MPI_Gather()	<ul style="list-style-type: none"> • MPI_Gather (&send_data, send_count, send_type, receive_type, root, communicator) • This is an operation that takes an N element data structure evenly distributed among P processors, and gathers them together on a single processor, called the root. • This operation assumes that every node contains send count (= N/P) data elements, each of type send type. • These elements then are sent to the processor called root where they are all stored in a data structure, called receive_data, of the same type. • Receive_count is identical to send count and represents the number of elements received from each individual node.

Lecture 4: Further Reference

- There is also a version of the Gather operation in which the elements are repeatedly gathered at every node.
 - That is, we do a gather with root = 0, then another gather with root = 1, etc.
 - This is called the AllGather operation.
 - Its parameter sequence is the same as in Gather except that the root field is not present

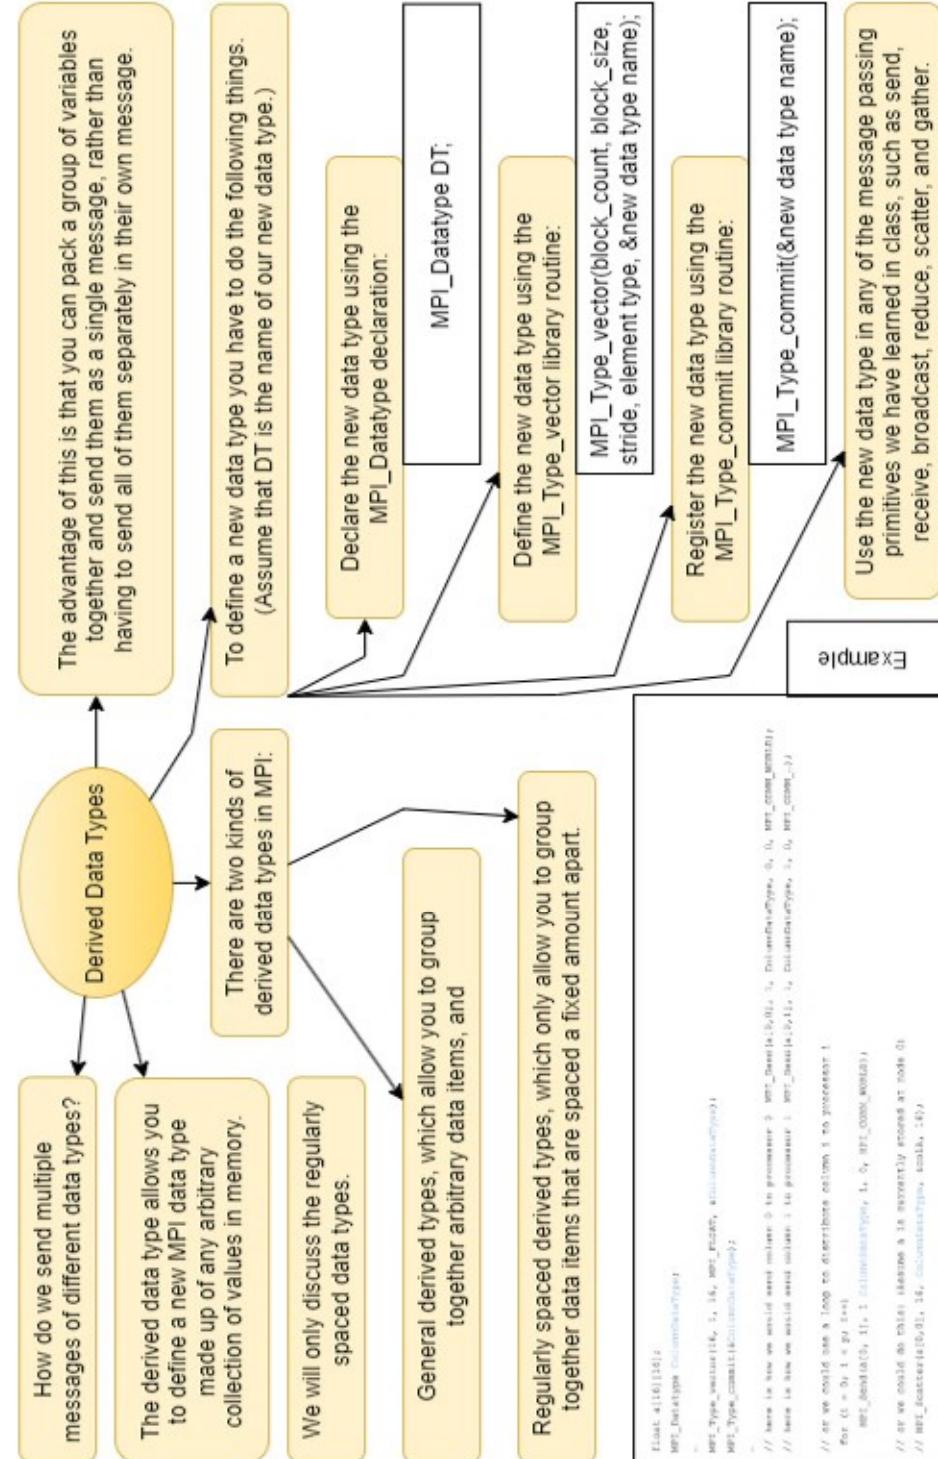
```
MPI_Scatter(&send_data, send_count, send_type, &receive_data,
receive_count, receive_type, root, communicator)
```

- This function is the exact opposite of Gather.
 - Now we take a data structure called send data containing N data elements of type send type and we distribute it evenly to P processors, giving each processor send count ($= N/P$) data elements.
 - The node that does the distribution is called the root.
 - So, lets assume that a floating point array Y located at processor 2 contained 64 data elements, and we wanted to distribute it evenly to four processors, each one getting 16 elements and each processor storing it as a local variable called `localY`.

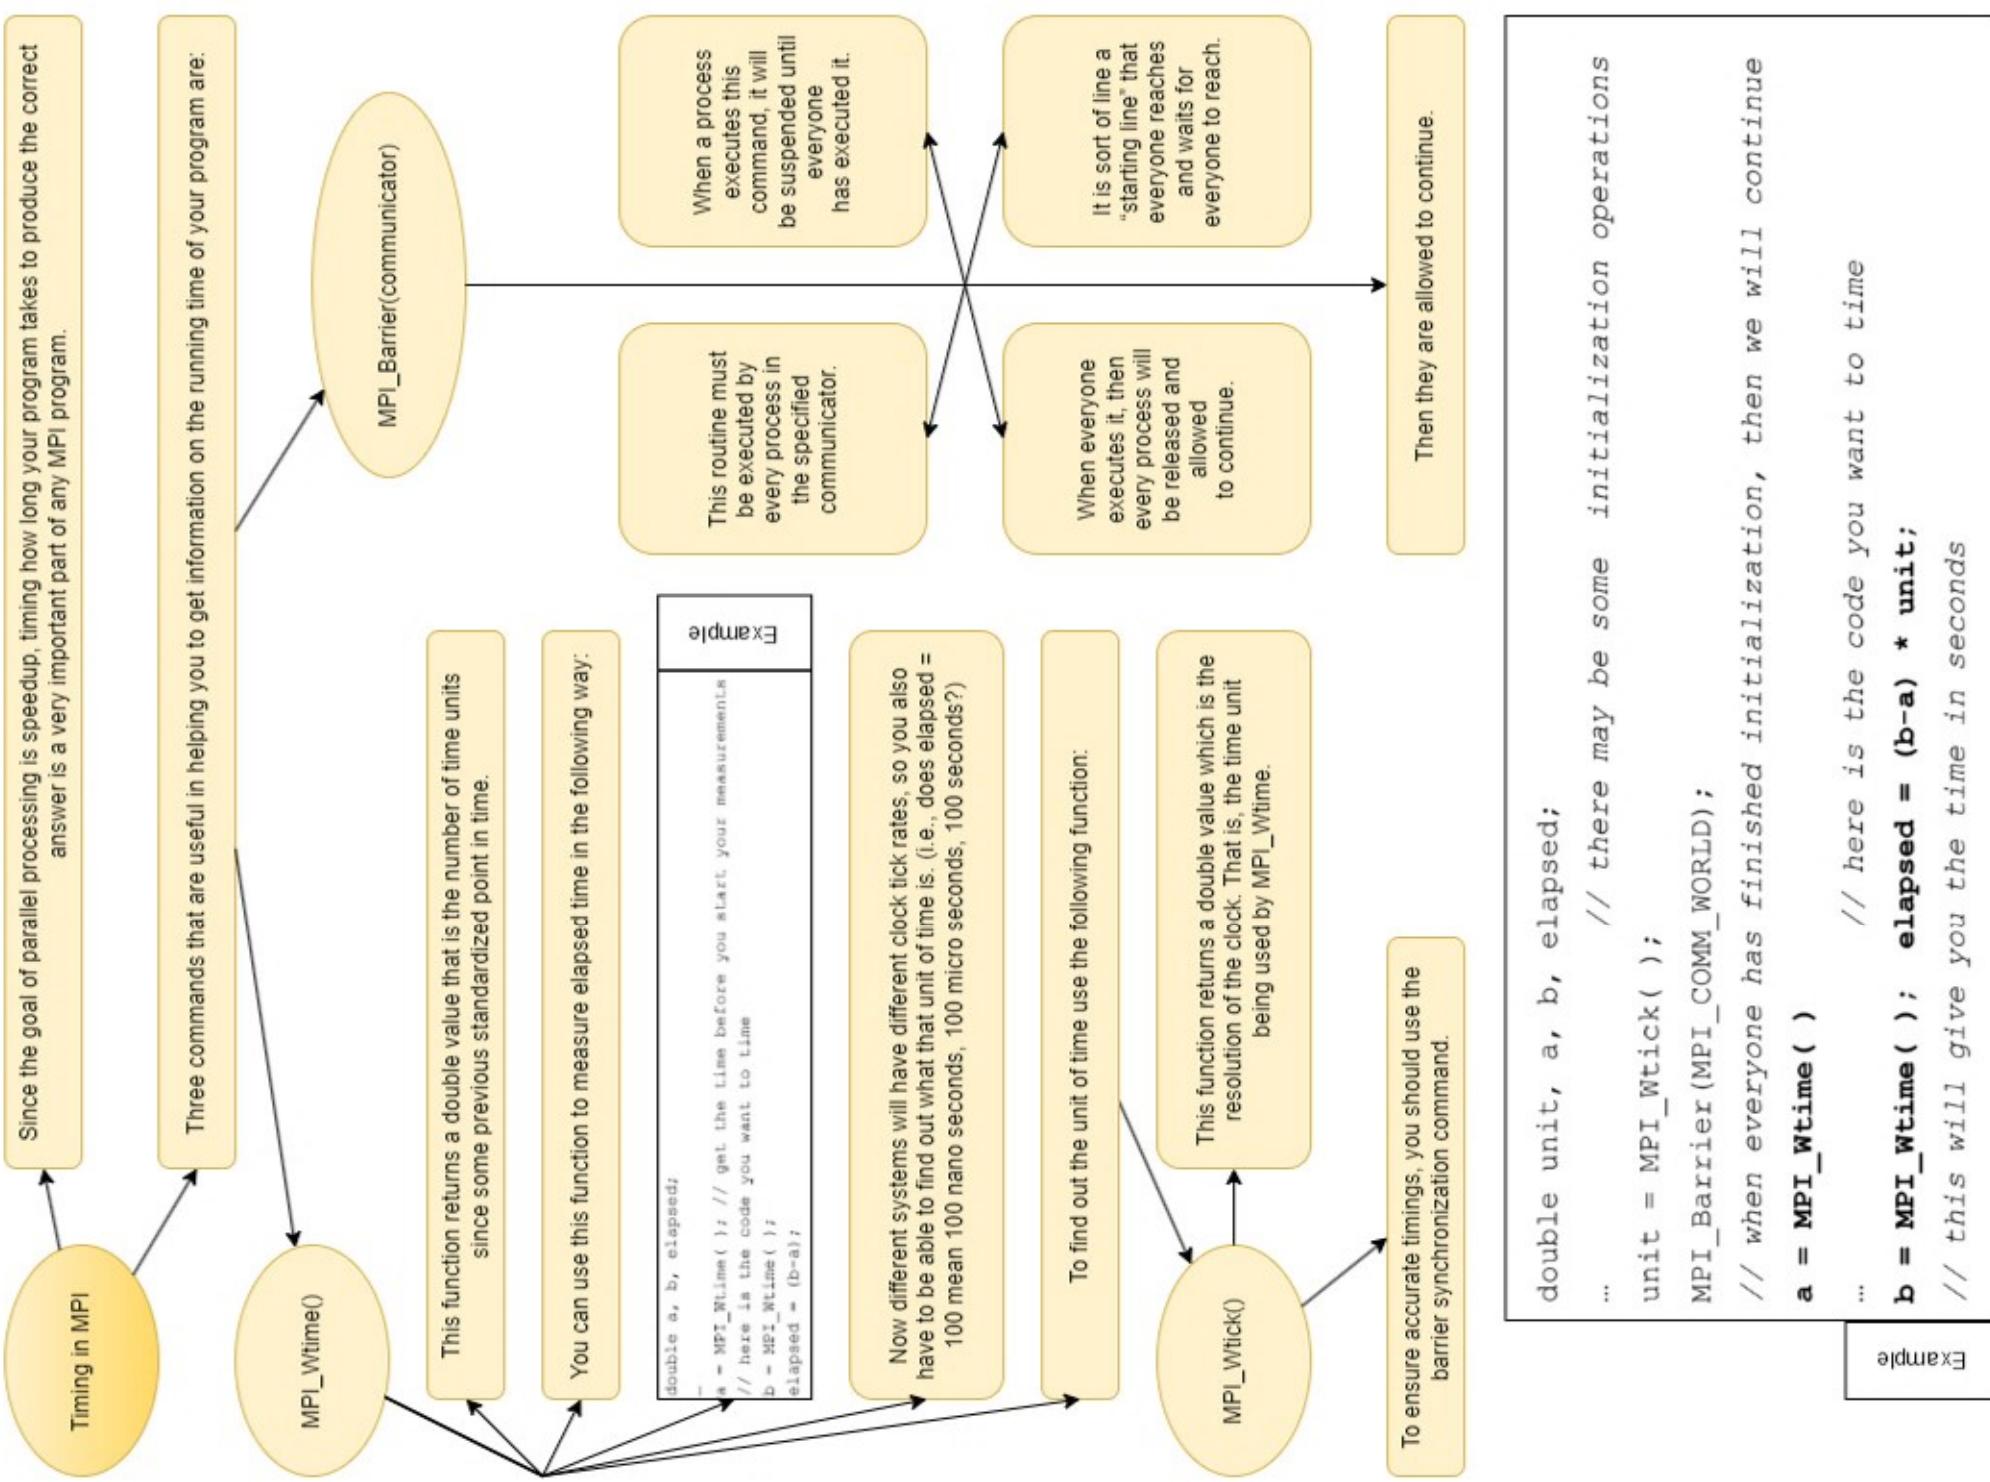
卷之三

- When we have finished executing this operation, the variable myY on processor 0 would contain the values of $y[0]$ to $y[15]$, the variable myY on processor 1 would contain $y[16]$ to $y[31]$, the variable myY on processor 2 would contain $y[32]$ to $y[47]$, etc.

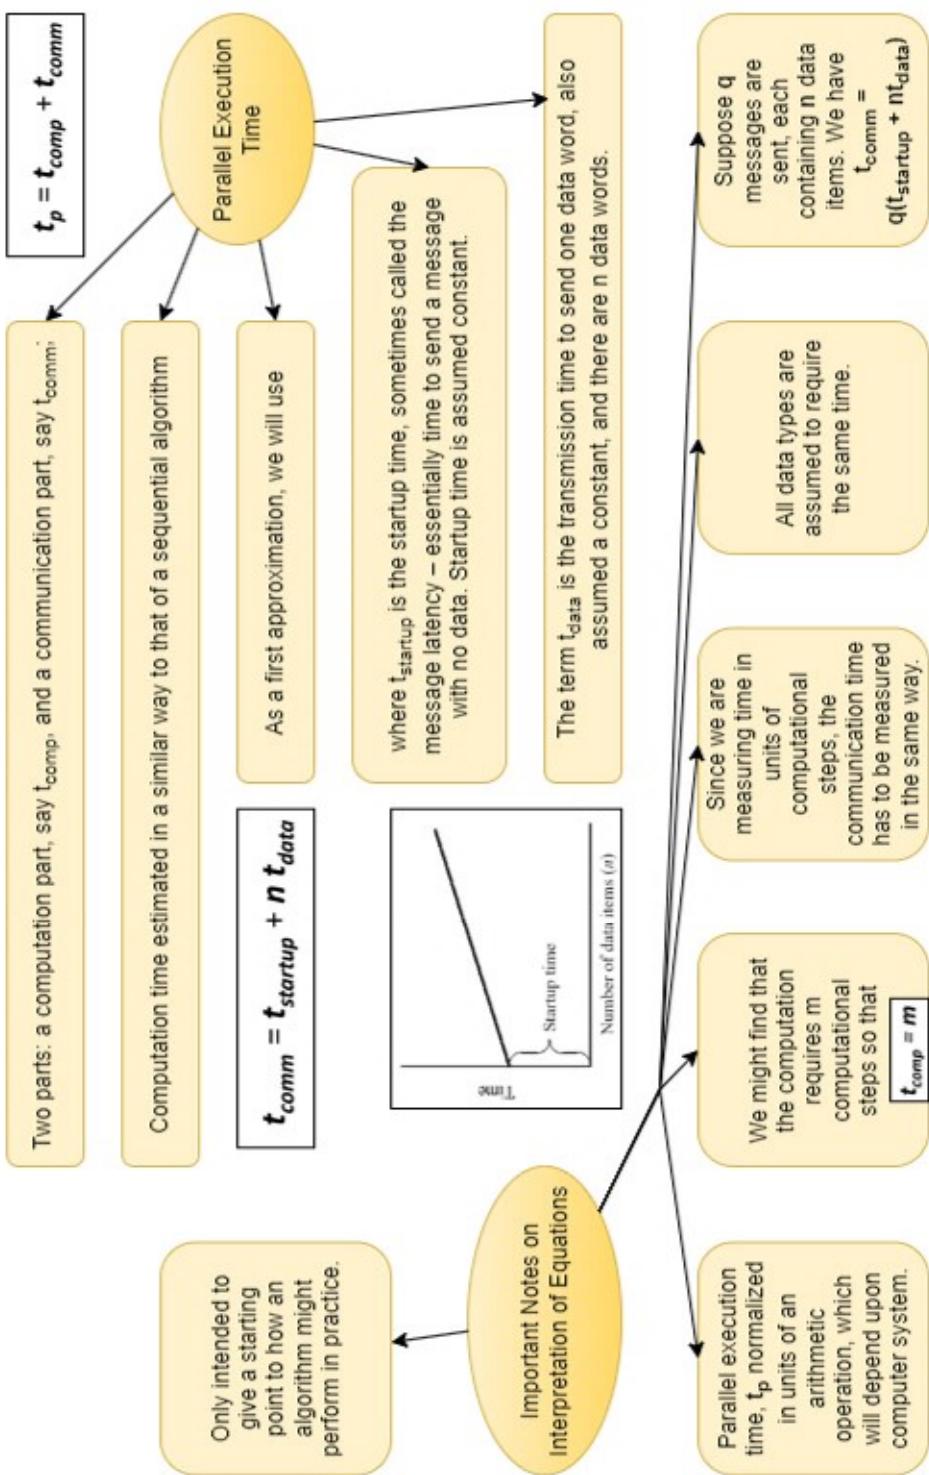
Lecture 4: Derived Data Types



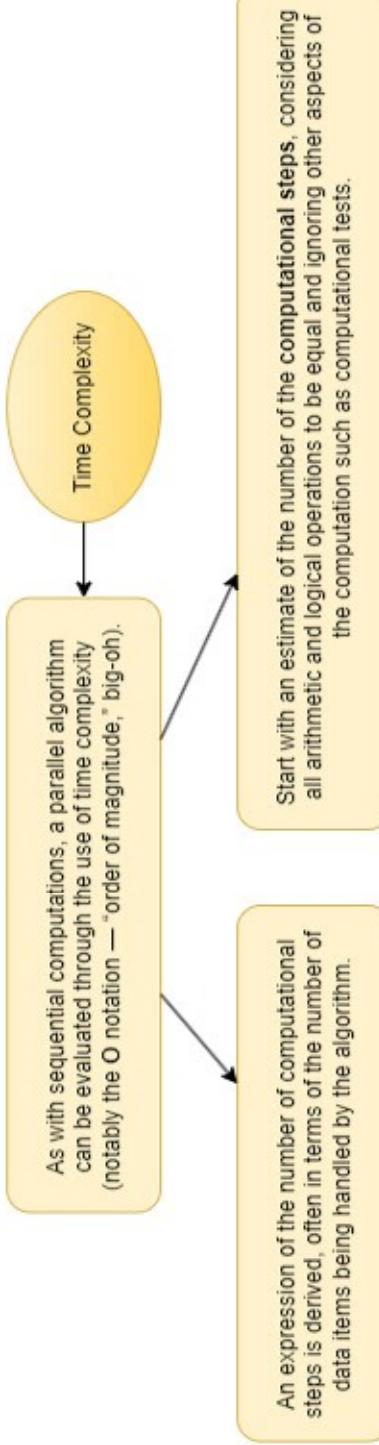
Lecture 4: Timing



Lecture 5: Parallel Execution Time



Lecture 5: Time Complexity

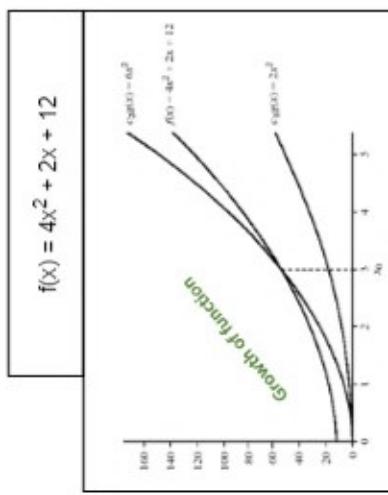


Lecture 5: Time Complexity

<ul style="list-style-type: none">Formal Definition<ul style="list-style-type: none">$f(x) = O(g(x))$ if and only if there exists positive constants c and x_0 such that $0 \leq f(x) \leq cg(x)$ for all $x \geq x_0$.where $f(x)$ and $g(x)$ are functions of x.For example, if $f(x) = 4x^2 + 2x + 12$, the constant $c = 6$ would work with the formal definition to establish that $f(x) = O(x^2)$, since $0 < 4x^2 + 2x + 12 \leq 6x^2$ for $x \geq 3$.	The big-O notation
<ul style="list-style-type: none">Unfortunately, the formal definition also leads to alternative functions for $g(x)$ that will satisfy the definition. For example,<ul style="list-style-type: none">$g(x)=x^3$ also satisfies the definition $4x^2+2x+12 \leq 2x^3$ for $x \geq 3$.Normally, we would use the function that grows the least for $g(x)$.	

<ul style="list-style-type: none">Many cases we have a “tight bound” in that function $f(x)$ equals $g(x)$ to within a constant factor. This can be captured in a Θ notation.	Theta notation - upper bound
<ul style="list-style-type: none">$f(x) = \Theta(g(x))$<ul style="list-style-type: none">if and only if there exists positive constants c_1, c_2, and x_0 such that $0 \leq c_1 g(x) \leq f(x) \leq c_2 g(x)$ for all $x \geq x_0$.If $f(x) = \Theta(g(x))$, it is clear that $f(x) = O(g(x))$ is also true.We can say that $f(x) = 4x^2 + 2x + 12 = \Theta(x^2)$, which is more precise than using $O(x^2)$.We should really only use the big O notation if and only if the upper bound on growth can be satisfied. However it is common practice to use big O in any event.	

<ul style="list-style-type: none">The lower bound on growth can be described by the Ω notation which is defined as	Omega notation - upper bound
<ul style="list-style-type: none">$f(x) = \Omega(g(x))$ if and only if there exists positive constants c and x_0 such that $0 \leq cg(x) \leq f(x)$ for all $x \geq x_0$.It follows from this definition that $f(x) = 4x^2 + 2x + 12 = \Omega(x^2)$.We can read $\Omega()$ as “grows at most as fast as” and $\Omega()$ as “grows at least as fast as.”The function $f(x) = \Theta(g(x))$ is true if and only if $f(x) = \Omega(g(x))$ and $f(x) = O(g(x))$.The Ω notation can be used to indicate the best case situation.	



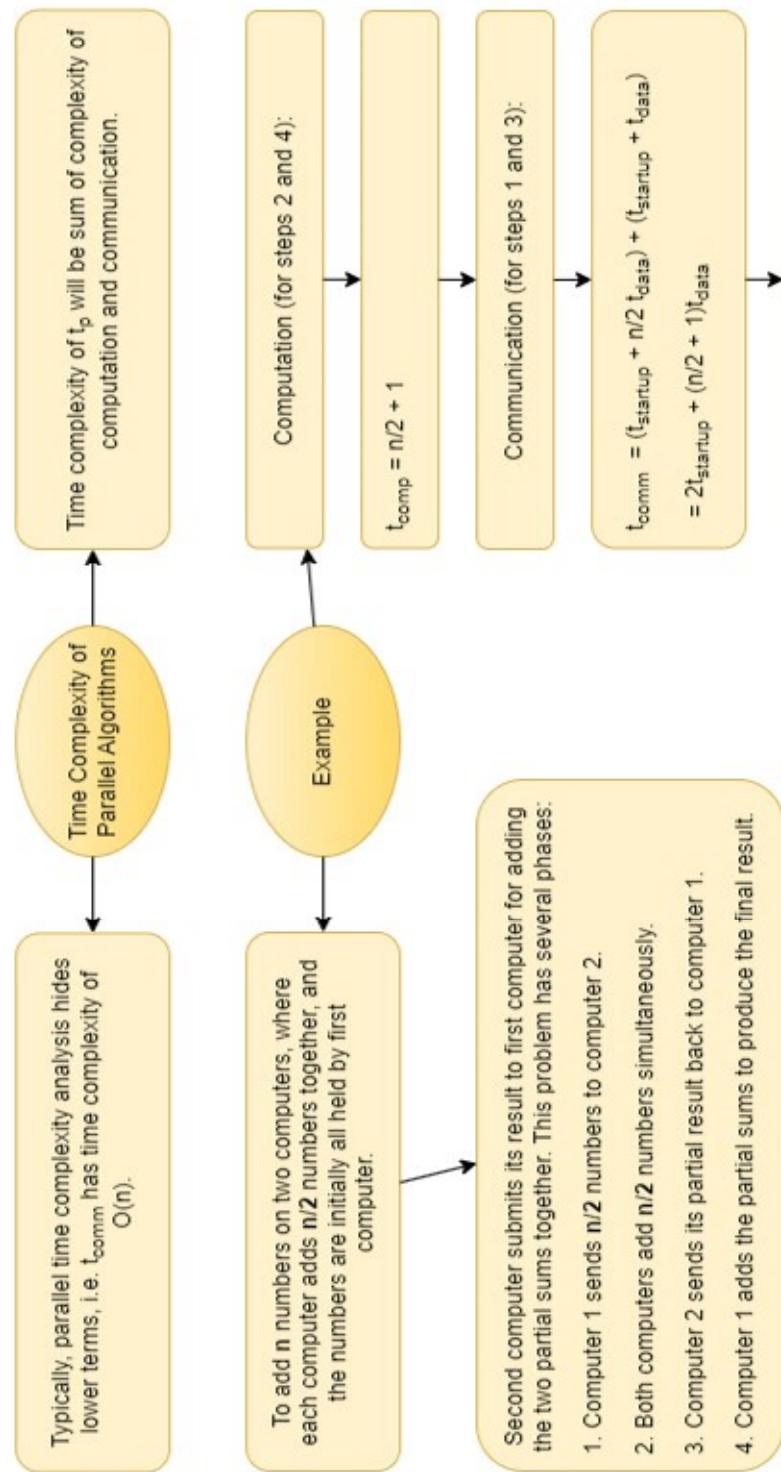
This would be indicated by a time complexity of Theta($n \log n$) and Omega(n^2).

Example

It may be that it requires at least $n \log n$ steps, but could require n^2 steps for n numbers depending upon the order of the numbers.

The execution time of a sorting algorithm often depends upon the original order of the numbers to be sorted.

Lecture 5: Time Complexity



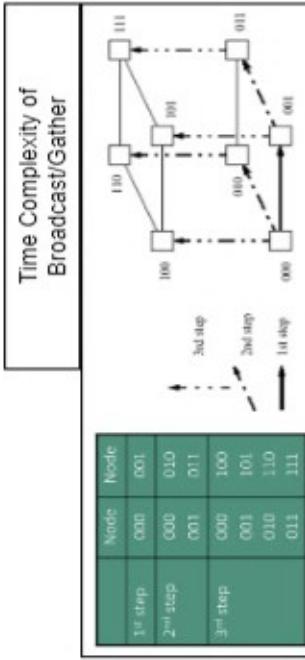
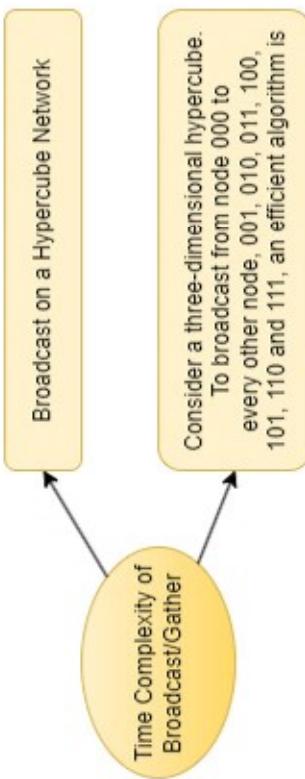
Lecture 5: Cost-Optimal Algorithms

A cost-optimal (or work-efficient or processor-time optimality) algorithm is one in which the cost to solve a problem is proportional to the execution time on a single processor system (using the fastest known sequential algorithm). i.e. $O(n \log n)$.

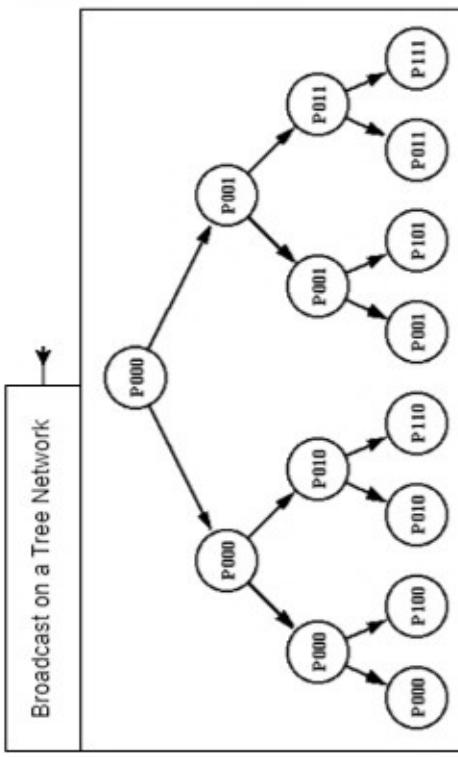
Suppose the best known sequential algorithm for a problem has time complexity of $O(n \log n)$.
Example

(Parallel time complexity) \times (number of processors) = sequential time complexity
A parallel algorithm for the same problem that uses n processes and has a time complexity of $O(\log n)$ is cost optimal, whereas a parallel algorithm that uses n^2 processors and has time complexity of $O(1)$ is not cost optimal.

Lecture 5: Communication routines & time complexity

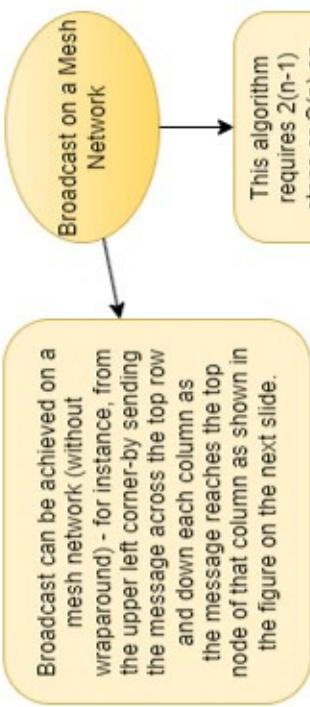


Lecture 5: Communication routines & time complexity

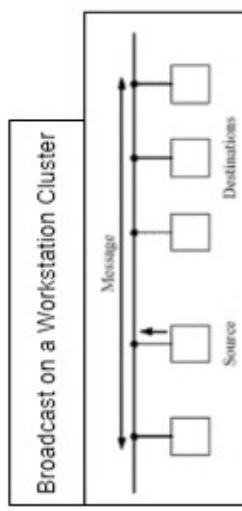
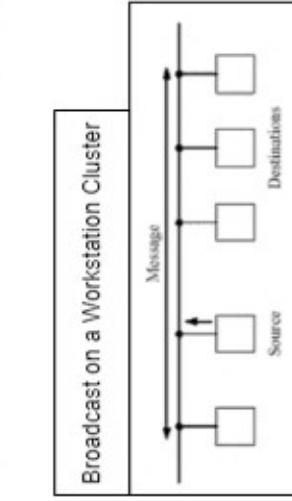
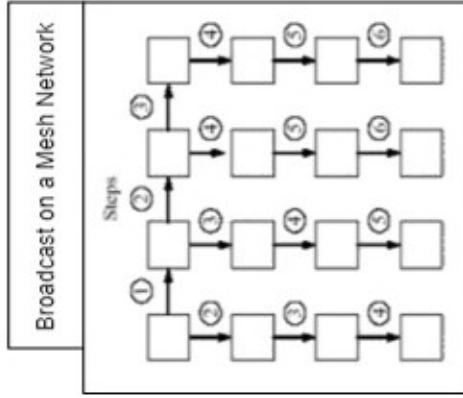


Broadcast can be done using a single connection that is read by all the destinations on the network simultaneously. Hence broadcast could be very efficient with an $O(1)$ time complexity for one data item; for n data item it is $O(n)$.

Most networked computers use a variety of network structures, and the convenient broadcast medium available in a single Ethernet will not generally be applicable.



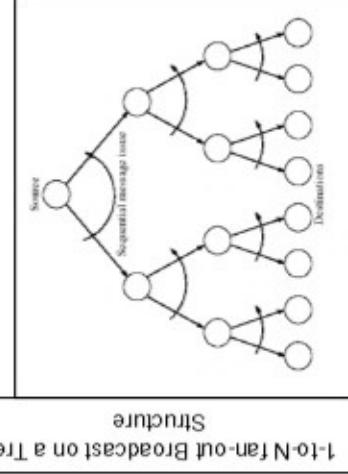
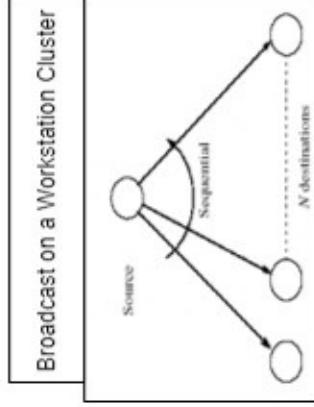
This algorithm requires $2(n-1)$ steps or $O(n)$ on the $n \times n$ mesh, again an optimal algorithm in terms of the number of steps because the diameter of a mesh without wraparound is given by $2(n-1)$.



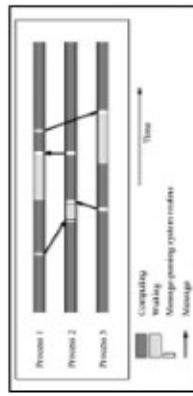
The 1-to-N fan-out broadcast applied to a tree structure has time complexity of $O(M+N/M)$ for one data item broadcast to N destinations, where there are M daemons.

Each of these daemons must receive the data in turn, and N/M is the number of destinations under each daemon.

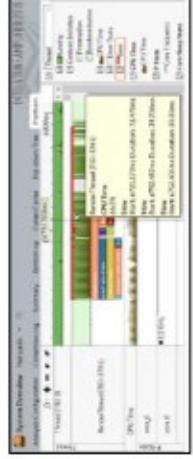
In addition a binary tree would be difficult to implement in a library call.



Lecture 5: Parallel program debugging and evaluation



Programs can be watched as they are executed in a space-time diagram (or process-time diagram).



A space-time diagram is useful to spot erroneous actions.

The events that created the space-time diagram can be captured and saved so that the presentation can be replayed without having to re-execute the program.

An example is the Intel VTune program visualization system.

Also shows the amount of time spent by each process on communication, waiting, and on message passing library routines.

If possible, run the program as a single process and debug as a normal sequential program.

Debugging Strategies

Geist et al. (1994) suggest a three-step approach to debugging message-passing programs:

Execute the program using two to four multitasked processes on a single computer. Now examine actions such as checking that messages are indeed being sent to the correct places. It is very common to make mistakes with message tags and have messages sent to the wrong places.

Measuring Execution Time

Elapsed time include waiting time for messages and it is assumed that the processor is not executing other programs at the same time.

MPI provides the routine `MPI_Wtime()` for returning time (in seconds).

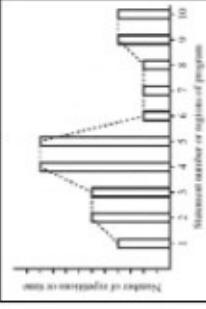
Timing

```
L1: time(&t1); /* start timer */  
...  
L2: time(&t2); /* stop timer */  
elapsed_time = difftime(t2, t1); /* elapsed_time = t2 - t1 */  
printf("Elapsed time = %5.2f seconds", elapsed_time);
```

To measure the execution time between point L1 and point L2 in the code, we might have a construction such as

One process, say P_0 , is made to send a message to another process, say P_1 . Immediately upon receiving the message, P_1 sends the message back to P_0 .

Profiling

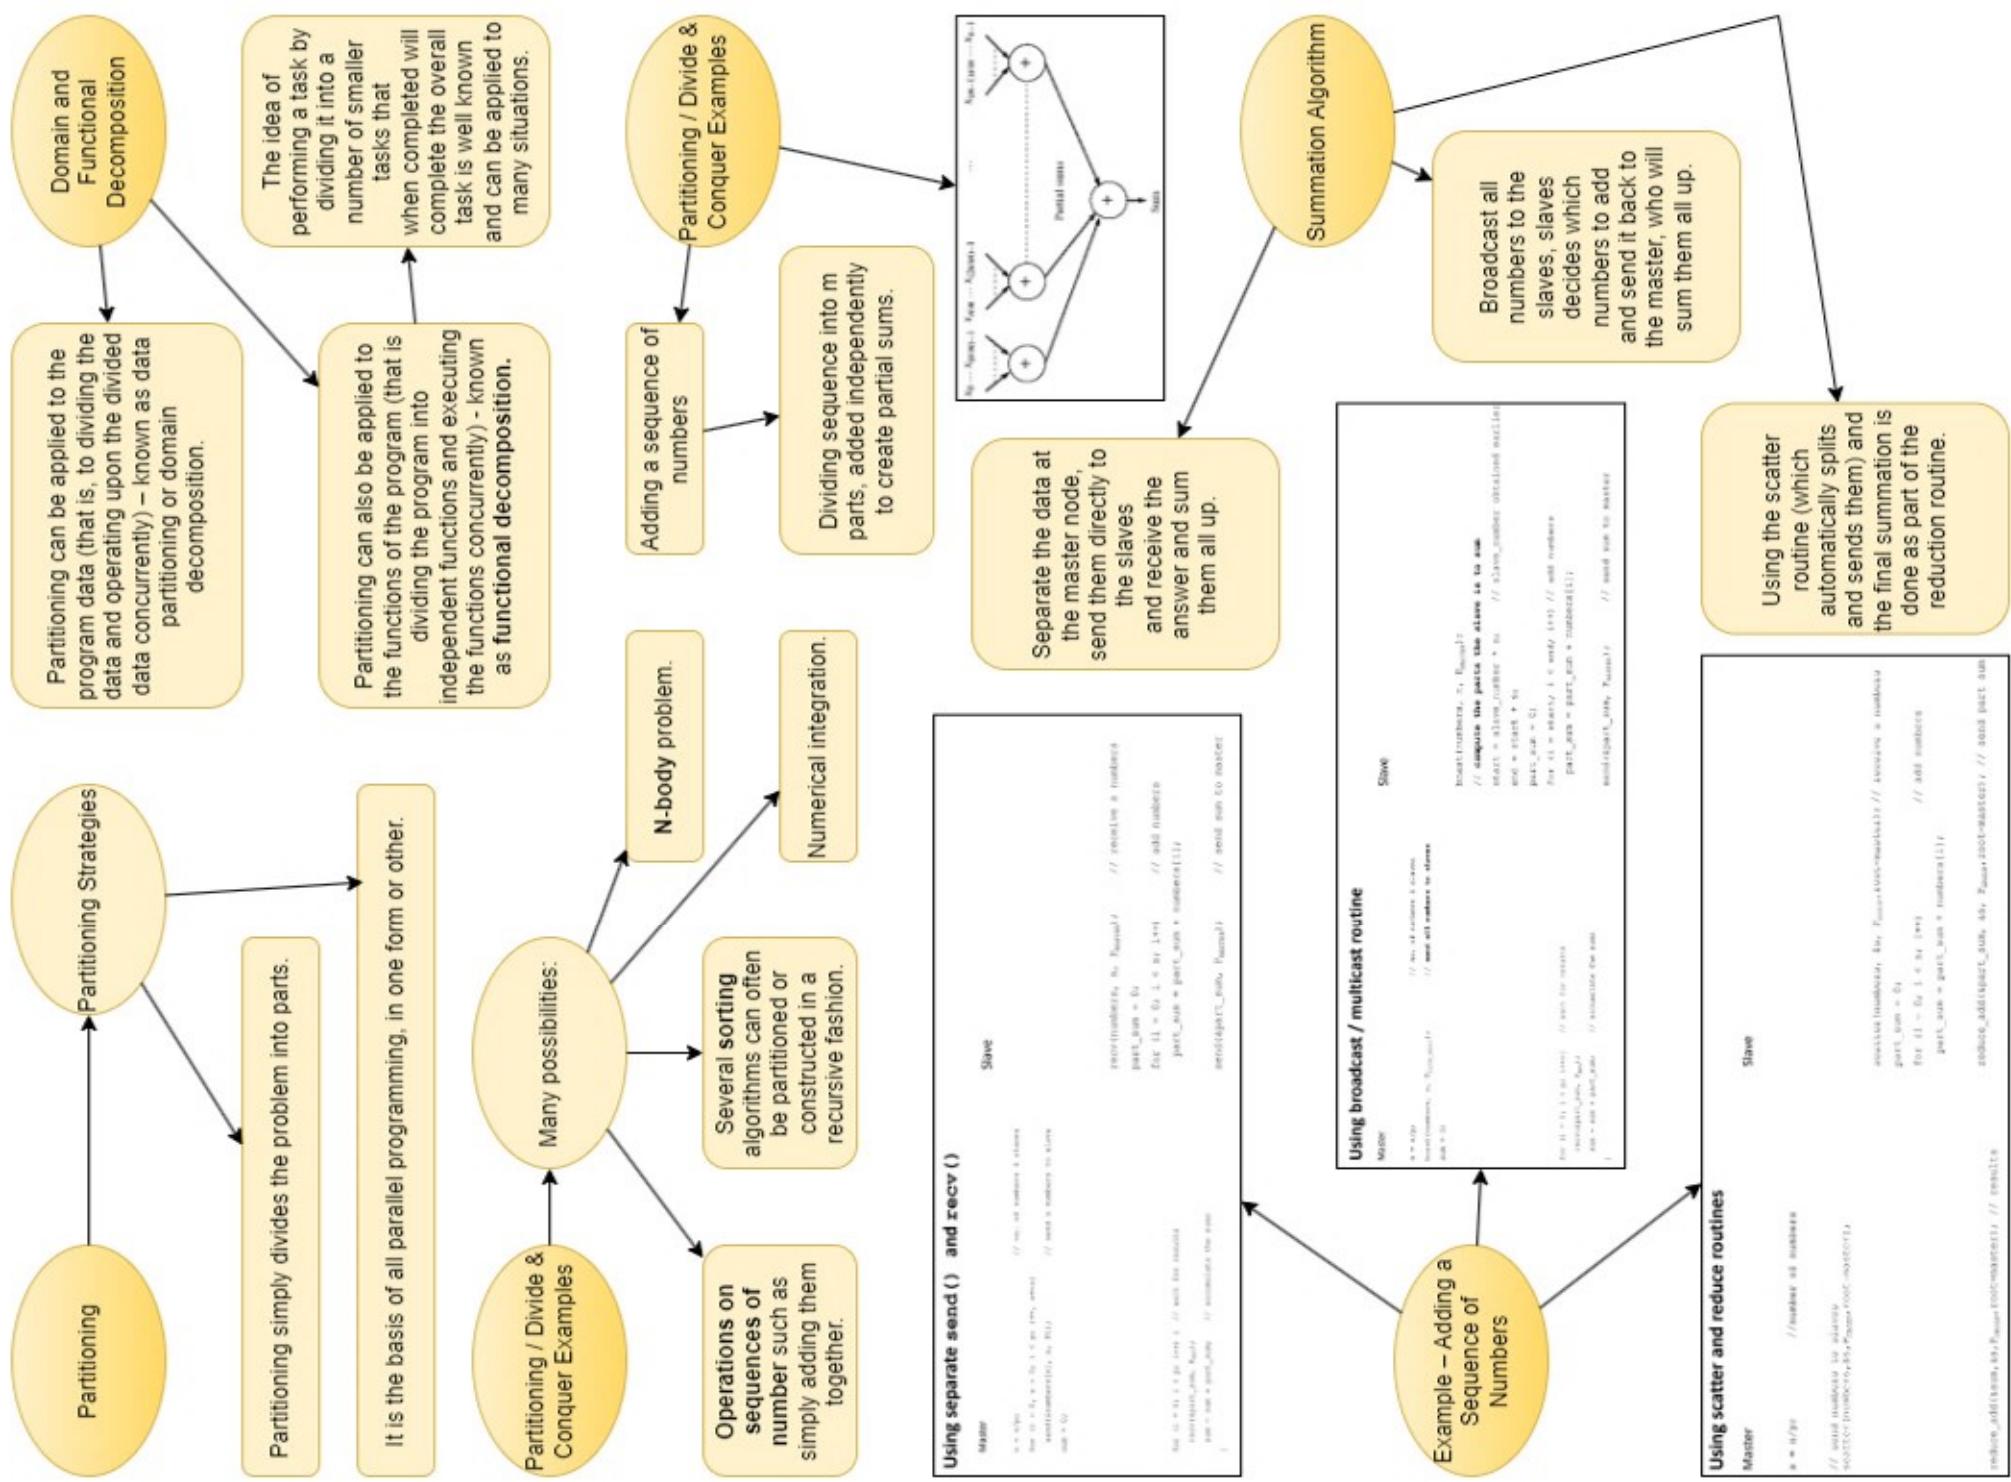


A profile of a program is a histogram or graph showing the time spent on different parts of the program.

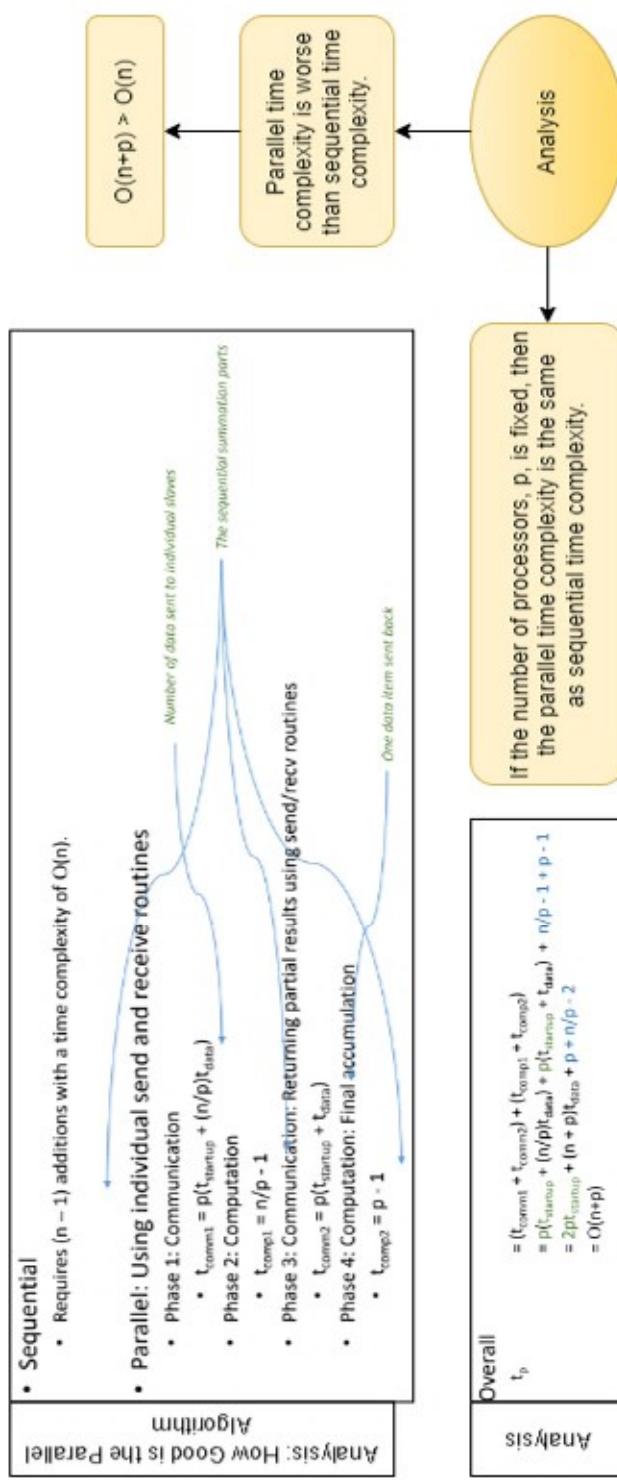
```
P0  
L1: time(&t1);  
send(&x, P1);  
recv(&x, P1);  
L2: time(&t2);  
elapsed_time = 0.5 * difftime(t2, t1);  
printf("Elapsed time = %5.2f seconds", elapsed_time);  
  
P1  
recv(&x, P0);  
send(&x, P0);
```

A profile can show the number of times particular source statements are executed.

Lecture 6: Partitioning



Lecture 6: Partitioning



Lecture 6: Divide and Conquer

Divide and Conquer

- Characterized by dividing a problem into sub-problems that are of the same form as the larger problem.
- Further divisions into smaller sub-problems are usually done by recursion. The recursive method will continually divide the problem until the tasks cannot be broken down into smaller parts.
- Then the very simple tasks are performed and results are combined, with the combining continued with larger and larger tasks.

Example

```

/* This is a pseudo code, actual implementation has been omitted */
int addlist(int *a) /* * a: add list of numbers, 0 */ {
    if (intradiv(a) == 1) /* if a is a single element */
        return (a[0]); /* return a[0] */
    else {
        /* divide a into two parts, a1 and a2 */
        /* recursive calls to add sub-lists */
        part_a1 = addlist(a1);
        part_a2 = addlist(a2);
        return (part_a1 + part_a2);
    }
}

```

Explanation for Code

As in all recursive definitions, a method must be present to terminate the recursion when the division can go no further.

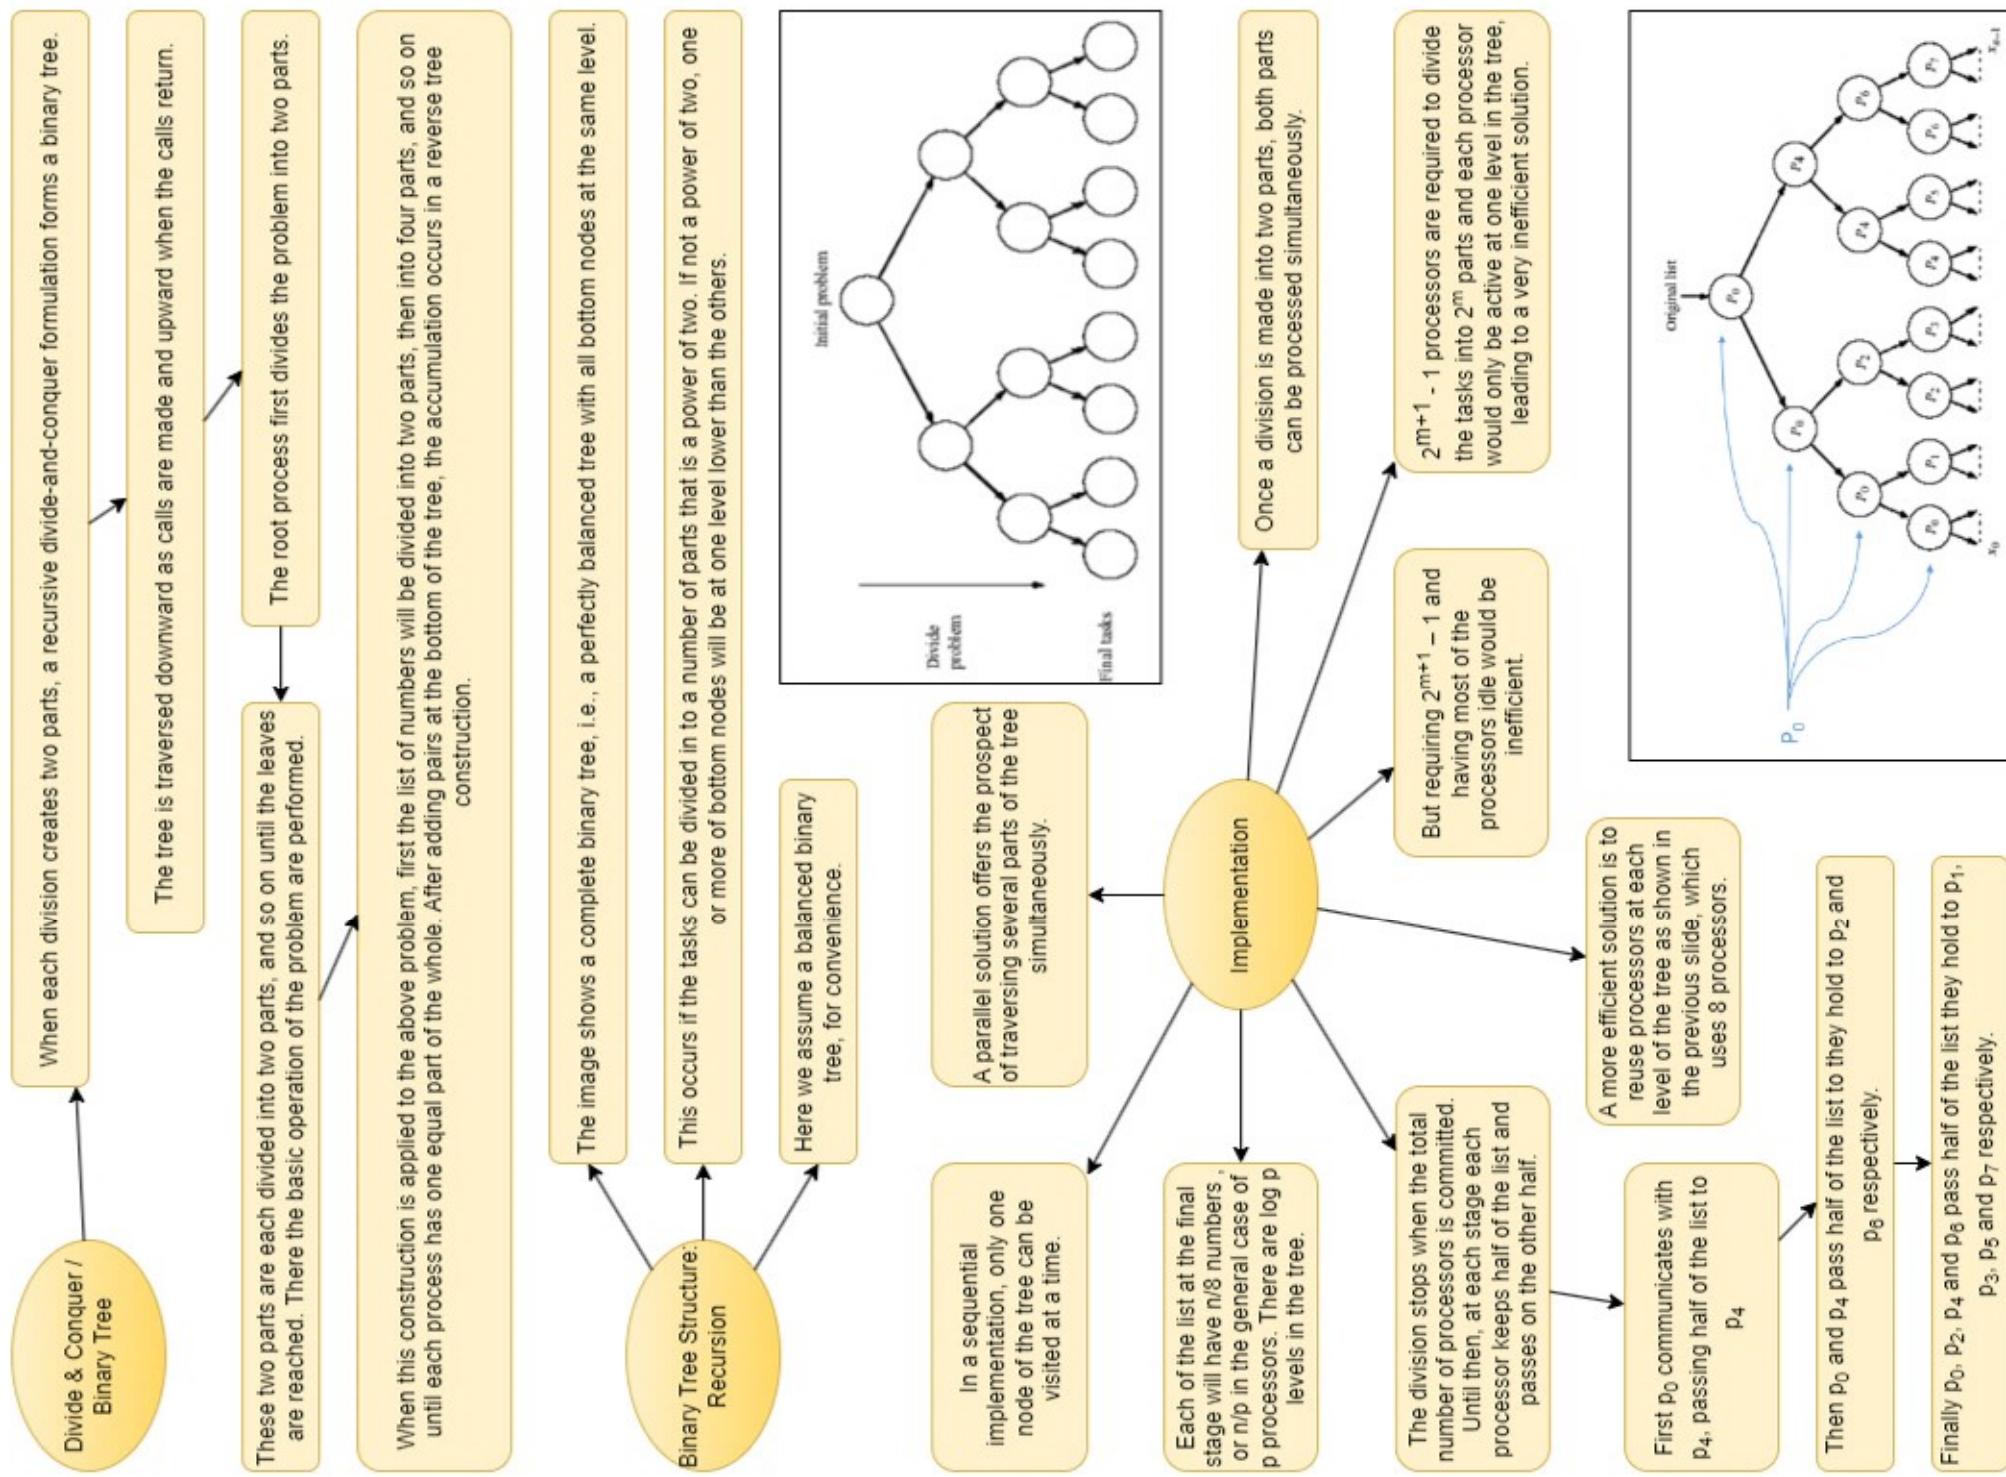
number(s) returns the number of numbers in the list pointed by s. If there are two numbers in the list, they are called n₁ and n₂. If there is one number in the list, it is called n₁ and n₂. If there are no numbers, both n₁ and n₂ are zero.

Separate if statements could be used for each of the cases: 0, 1, or 2 numbers in the list. Each could cause termination of the recursive call.

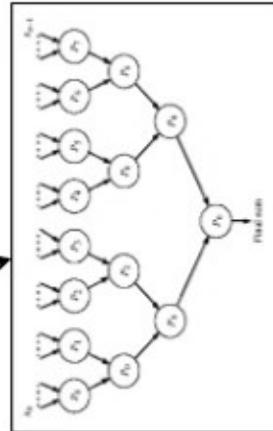
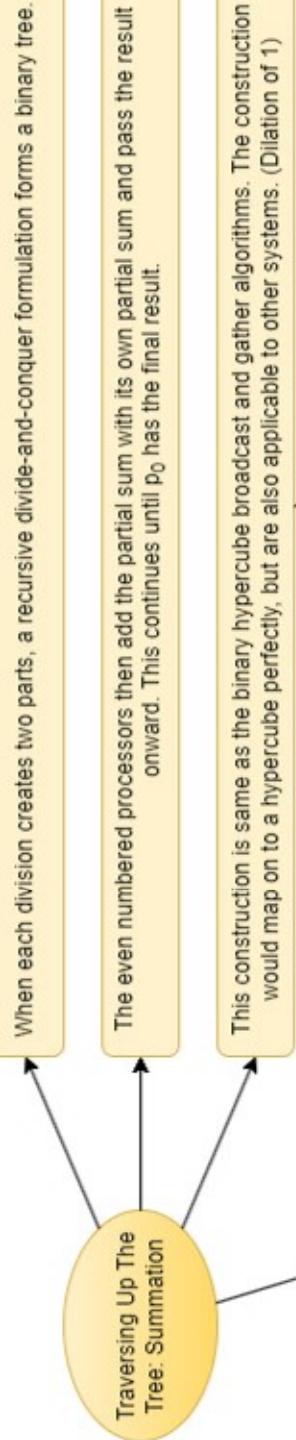
This method can also be used for other global operations on a list. E.g.:

Sorting a list by dividing the list in to smaller and smaller lists to sort. Mergesort and quicksort-sorting algorithms are usually described by such recursive definitions.

Lecture 6: Divide and Conquer



Lecture 6: Divide and Conquer



m-ary Tree

Divide and conquer can also be applied where a task is divided into more than two parts at each stage.

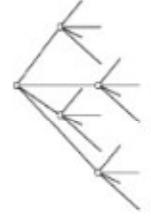
Example: Task broken into four parts. The sequential recursive definition would be

```

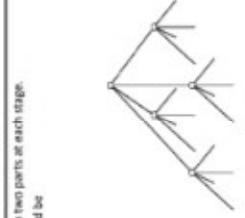
task calculate(s1) {
    if (node_id == 1) {
        send(0,1, P1);
        calculate(s1);
        send(0,2, P2);
        calculate(s1);
        send(0,3, P3);
        calculate(s1);
        part_sum = *s1;
        send(part_sum, P1);
        part_sum = part_sum + part_sum;
        send(part_sum, P2);
        part_sum = part_sum + part_sum;
        send(part_sum, P3);
        part_sum = part_sum + part_sum;
    }
}

```

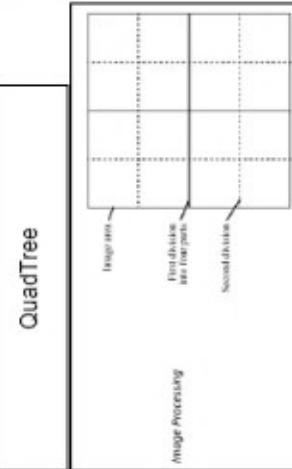
A tree in which each node has four children is called a quadtree.



quadtree

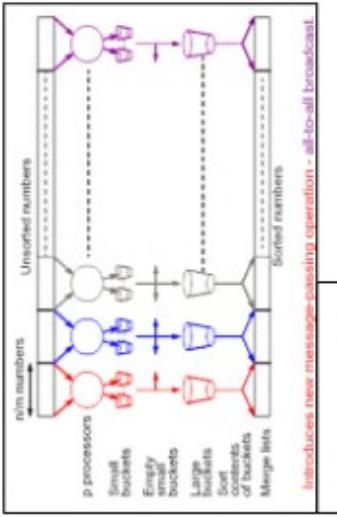
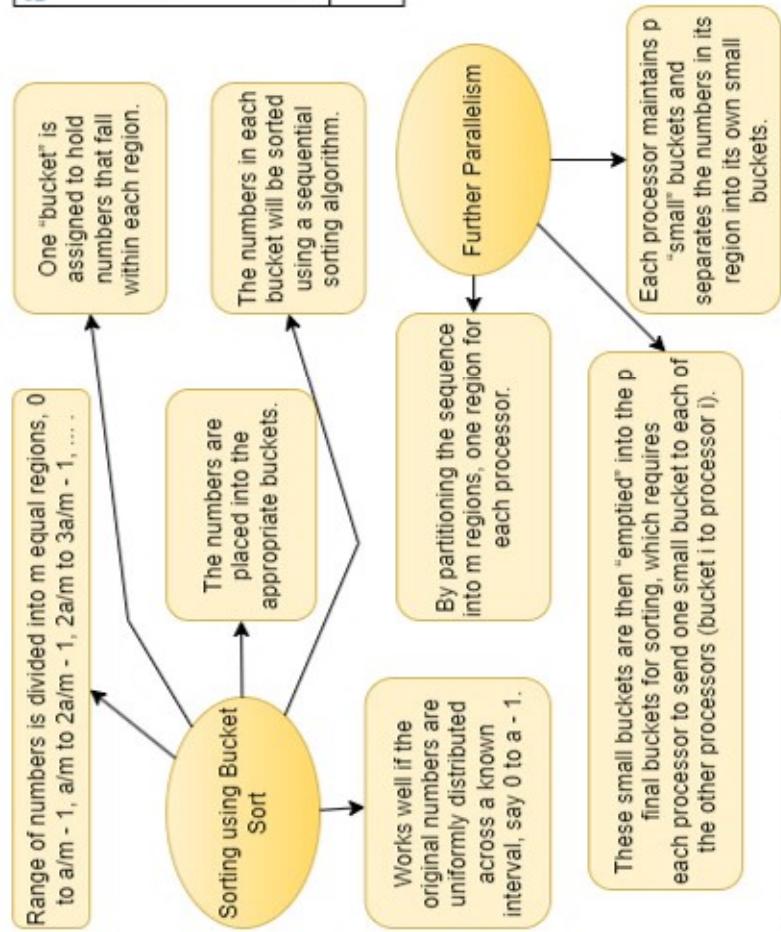


Parallel Code (P4)



- Similar sequences are required for the other processes

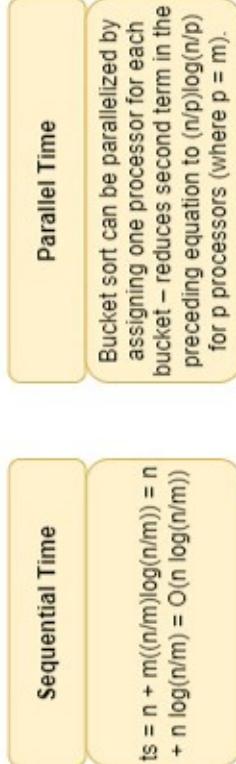
Lecture 6: Partitioning & Divide and Conquer – Example I



More Parallelism

An all-to-all routine sends data from each process to every other process and is illustrated on the next slide. This type of routine is available in MPI (MPI_Alltoall()), which is more efficient than using individual send()s and recv()s.

The all-to-all routine will actually transfer the rows of an array to column (and hence transpose a matrix).



Parallel Time

Bucket sort can be parallelized by assigning one processor for each bucket – reduces second term in the preceding equation to $(n/p)\log(n/p)$ for p processors (where $p = m$).

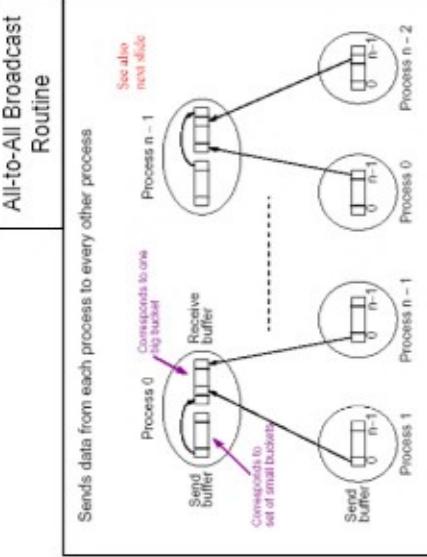
- Phase 1 — Computation and Communication (partition numbers)

$$t_{\text{unif}} = t_{\text{setup}} + t_{\text{gap}}$$

$$t_{\text{comp4}} = n/p$$
- Phase 2 — Computation (Sort into Small Buckets)

$$t_{\text{comp2}} = n/p$$
- Phase 3 — Communication (Send to Large Buckets)

$$t_{\text{send3}} = (p-1)t_{\text{gap}} + (n/p^2)t_{\text{gap}}$$
- Analysis



- Phase 4 — Computation (Sort Large Buckets)

$$t_{\text{comp4}} = (n/p)\log(n/p)$$
- Overall
 - It is assumed that the numbers are uniformly distributed to obtain these formula.
 - Worst-case scenario would occur when all the numbers fall into one bucket!

Lecture 7: Numerical Integration

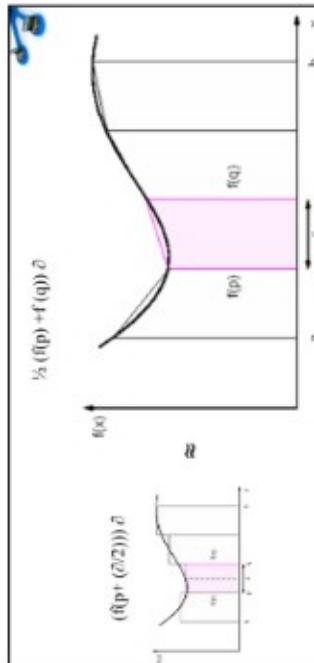
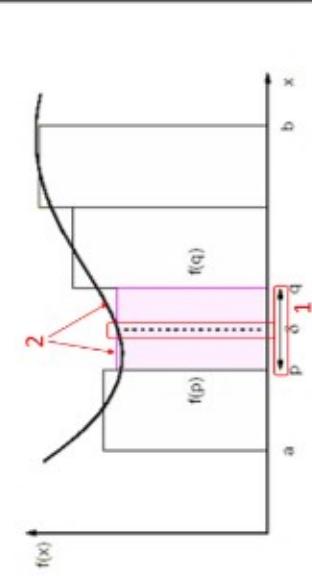
Previous example: Divide the problem and solve each subproblem. The problem was assumed to be divided into equal parts and simple partitioning was employed.

Numerical Integration

Some times simple partitioning will not give the optimum solution, especially if the amount of work in each part is difficult to estimate.

A general divide-and-conquer technique divides the region continually into parts and lets some optimization function decide when certain regions are sufficiently divided. Let us take a different example, numerical integration

Each region calculated using an approximation given by rectangles:
Aligning the rectangles:



Suppose we were to sum the area from $x=a$ to $x=b$ using p processes, numbered 0 to $p-1$. The size of the region for each process is $(b-a)/p$. To calculate the area in the described manner, a section of SPMD pseudo-code could be as follows for Process Pi

```

int i;
printf("Enter number of intervals ");
scanf("%d", &n);
// read number of intervals required

broadcast(n, P_parallel);
region = (b - a)/p;
start = i * region + i;
end = start + region;
d = (b - a)/n;
area = 0.0;
for (x = start; x < end; x = x + d)
    area = area + f(x) * f(x+d);
area = 0.5 * area * d;
reduce_and_intergral(area, P_parallel);
// form sum of areas
}

```

$$I = \int_a^b f(x) dx$$

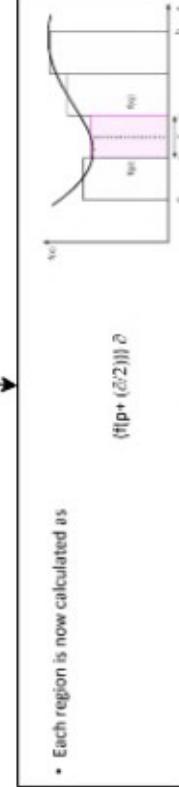
To integrate the function (i.e. to compute the 'area under the curve'), we can divide the area into separate parts, each of which can be calculated by a separate process.

Each region could be calculated using an approximation given by rectangle, where $f(p)$ and $f(q)$ are the heights of the two edges of the rectangular region and sigma is the width (interval).

Bucket sort, for example, is only effective when each region has approximately the same number of numbers.

The complete integration can be approximated by the summation of rectangular regions from a to b by aligning the rectangles so that the upper midpoint each rectangle intersect with the function. This has the advantage that errors on each side of midpoint end tend to cancel.

Here we take the actual intersections of the lines with the function to create rectangles as shown in the fig.



* Each region is now calculated as

Such approximate numerical methods for computing a definite integral using linear combination of values are called quadrature methods

Prior to start of the computation, one process is statically assigned to be responsible for computing each region. By making the interval smaller, we come closer to attaining the exact solution.

Let us consider a trapezoidal method.

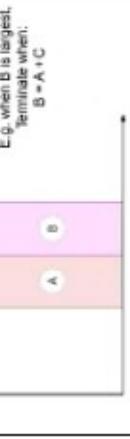
Since each calculation is of the same form, the SPMD model is appropriate.

Lecture 7: Adaptive Quadrature

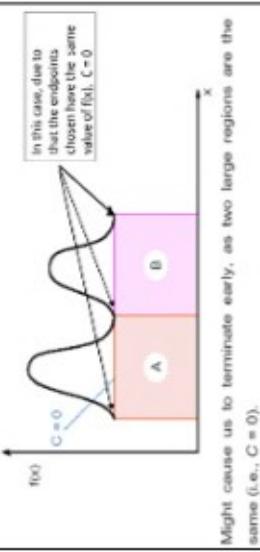
Method whereby the solution adapts to the shape of the curve. Example

Use three areas, A, B, and C. The computation is terminated when the area computed for the largest of the A and B regions is sufficiently close to the sum of the areas computed for the other two regions.

Adaptive Quadrature



Some care might be needed in choosing when to terminate



Gravitational N-Body Problem

Objective is to find the positions and movements of bodies in space (say planets) that are subject to gravitational forces from other bodies using Newtonian laws of physics.

Computations of areas under slowly varying parts of the curve stop earlier than computations of areas under more rapidly varying parts.

The spacing sigma will vary across the interval. The consequent of this is that a fixed process assignment will not lead to the most efficient use of processors, for which some care might be needed in choosing when to terminate (i.e. $C=0$).

Note that once the bodies move into a new position, the force (F) changes. The computation has to be repeated.

Over the time Δt interval, the position changes

$$x^{t+1} = x^t + v \Delta t$$

where x^t is its position at time t

The new velocity will be

$$v^{t+1} = v^t + \frac{F \Delta t}{m}$$

where v^{t+1} will be the velocity at time $t+1$, and v^t is the velocity at time t .

Adaptive Quadrature with False Termination

For example, the function such as shown below in the fig might cause us to terminate early, as the two large regions are the same (i.e. $C=0$).

The gravitational force between two bodies of masses m_a and m_b is given by

$$F = \frac{G m_a m_b}{r^2}$$

where G is the gravitational constant and r is the distance between the bodies.

Subject to forces, a body will accelerate according to Newton's second law:

$$F = m a$$

Let the time interval be Δt . Hence, for a body with mass m , the resulting F applied will be

$$F = \frac{m(v^{t+1} - v^t)}{\Delta t}$$

Most of the problems that we want to solve, such as weather forecasting, involves such computation and they are in a 3 dimensional space.

Sequential Code

- The overall gravitational N-body computation can be described by the algorithm.


```

for (t = 0; t < maxTime; t++) {
    for each time period
        for each body (particle)
            g = Force_Routine(i);
            v_new[i] = v[i] + g * dt; // compute new velocity
            x_new[i] = x[i] + v_new[i] * dt; // new position
        for each body (particle)
            x[i] = x[i] + x_new[i]; // update the position
            v[i] = v[i] + v_new[i]; // and the velocity
    repeat for next time period
}
      
```

$$F = \frac{G m_a m_b}{r^2}$$

$$\begin{aligned}
 F_x &= \frac{G m_a m_b (x_b - x_a)}{r^2} \\
 F_y &= \frac{G m_a m_b (y_b - y_a)}{r^2} \\
 F_z &= \frac{G m_a m_b (z_b - z_a)}{r^2}
 \end{aligned}$$

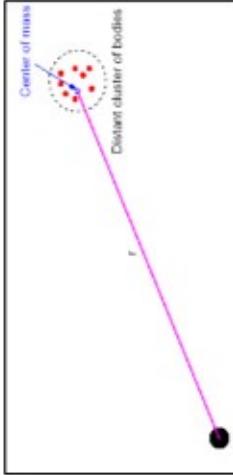
where the particles (bodies) of mass m_a and m_b and have the coordinates (x_a, y_a, z_a) and (x_b, y_b, z_b)

Lecture 7: Adaptive Quadrature

The algorithm is an $O(N^2)$ algorithm (for one iteration) as each of the N bodies is influenced by each of the other $N - 1$ bodies.

Not feasible to use this direct algorithm for most interesting N-body problems where N is very large.

Time complexity can be reduced using the observation that a cluster of distant bodies can be approximated as a single distant body of the total mass of the cluster situated at the centre of mass of the cluster.



After the tree has been constructed, the total mass and centre of mass of the sub-cube is stored at each node.

Start with whole space (3-Dimensional) in which one cube contains the bodies (or particles).

The leaves represent cells each containing one body.

First, this cube is divided into eight sub-cubes.

This process creates an oct-tree; that is, a tree with up to eight edges from each node.

If a sub-cube contains no particles, the sub-cube is deleted from further consideration.

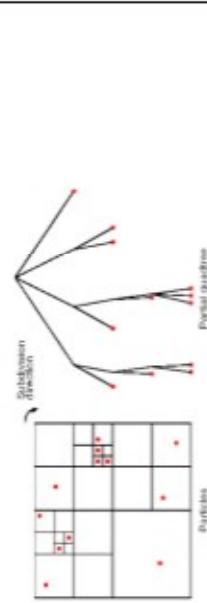
Barnes-Hut Algorithm (2)

- For a 2 dimensional problem, each recursive sub division will create four sub-areas and a quad-tree.

- In general, the tree will be very unbalanced.

- Figure below illustrates the decomposition for a two dimensional space and resultant quadtree.

- The three dimensional case follows the same construction except with up to eight edges from each node



If a sub-cube contains more than one body, it is recursively divided until every sub-cube contains one body. (If a sub-cube contains one body, stop recursion).

$$r > \frac{d}{\theta}$$

The Barnes-Hut Algorithm creates a very unbalanced tree.

where θ is a constant typically 1.0 or less (θ is called the opening angle).

Constructing the tree requires a time of $O(n \log n)$, and so does computing all the forces, so that the overall time complexity of the method is $O(n \log n)$. (Remember earlier, the complexity of a tree structure is $O(\log n)$ and at each level we need to "decompose" it into n bodies.

Orthogonal Recursive Bisection

Repeated until there are as many areas as processors, and then one processor is assigned to each area.

For 12 processors, i.e. $n = 12$

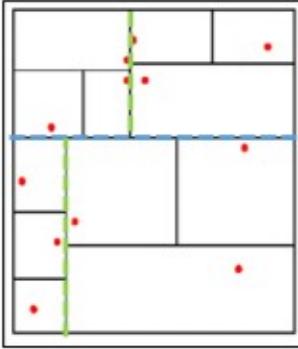
Consider a two-dimensional square area

A better way of dividing the groups would be Orthogonal Recursive Bisection

Repeated until there are as many areas as processors, and then one processor is assigned to each area.

First, a vertical line is found that divides the area into two areas each with an equal number of bodies.

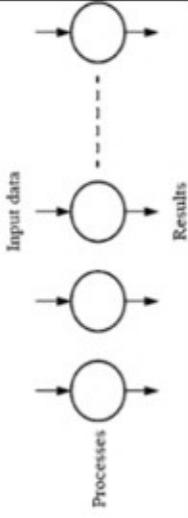
For each area, a horizontal line is found that divides it into two areas each with an equal number of bodies.



Lecture 7: Embarrassingly Parallel

A computation that can be divided into a number of completely independent parts, each of which can be executed by a separate processor.

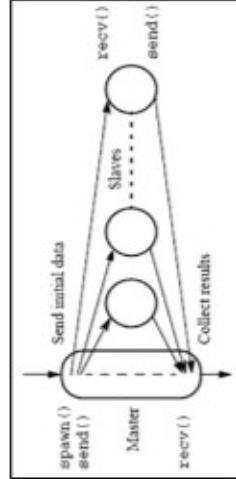
Naturally parallel



Each process requires different or the same data and produces results from its input data without any need for results from other processes.

Embarrassingly Parallel

Master-slave approach for data transmission and processing in embarrassingly parallel computation is shown below:



Can achieve Maximum possible speedup